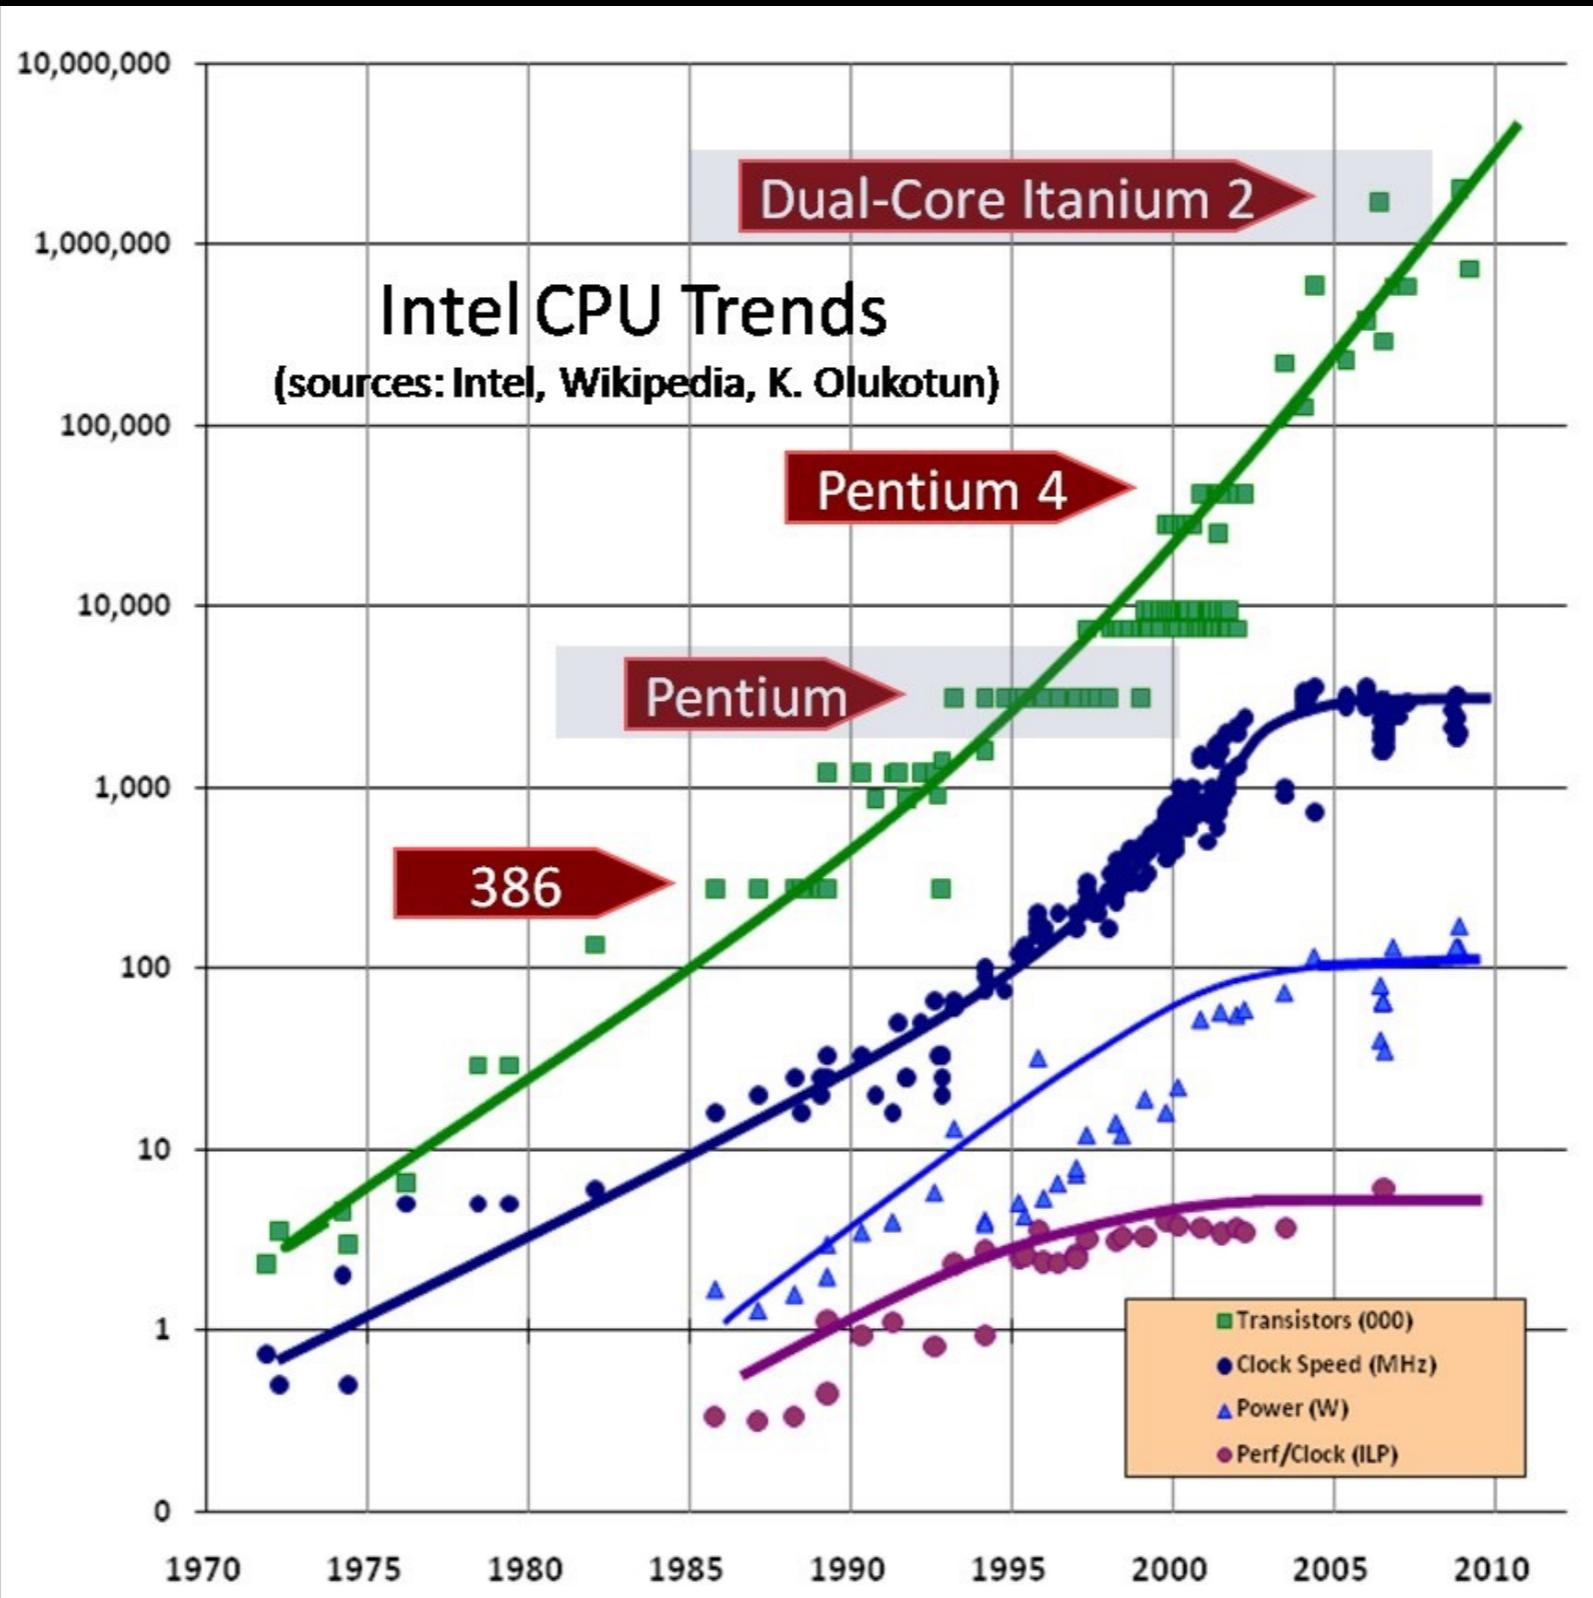


Utilización de esqueletos algorítmicos en EcmaScript para paralelizar el computo en navegadores Web a través de Web Workers

Schenkelman, Damian
Servetto, Matias

Tutora: Rosa Wachenchauzer



<http://www.extremetech.com/wp-content/uploads/2013/08/CPU-Scaling.jpg>

Paralelización

Comunicación de datos

Comunicación de código

Sincronización

Consolidación de resultados

Esqueletos algorítmicos

Modelo de alto nivel para programación paralela

Operaciones trivialmente
parallelizables

```
dobles = map(e => e * 2, [1,2,3,4]);
```

```
// dobles [2,4,6,8]
```

map

```
map :: (a -> b) -> [a] -> [b]
```

```
dobles = map(e => e * 2, [1,2,3,4]);  
// dobles [2,4,6,8]
```

map

map :: (a -> b) -> [a] -> [b]

```
pares = filter(e => e % 2 == 0, [1,2,3,4]);  
// pares [2,4]
```

filter

```
filter :: (a -> Bool) -> [a] -> [a]
```

```
pares = filter(e => e % 2 == 0, [1,2,3,4]);  
// pares [2,4]
```

filter

filter :: (a -> Bool) -> [a] -> [a]

```
suma = reduce((ac, cur) => ac + cur,  
0,  
[1,2,3,4]);  
  
// suma 10
```

reduce

reduce :: (a -> b -> a) -> a -> [b] -> a

```
suma = reduce((ac, cur) => ac + cur,  
0,  
[1,2,3,4]);  
  
// suma 10
```

reduce

reduce :: (a -> b -> a) -> a -> [b] -> a

```
pipe(filter(e => e % 2 == 0),  
  map(e => e * 2), [1,2,3,4]);
```

```
// [4, 8]
```

Encadenamiento

```
pipe :: ([a] -> [b]) -> ([b] -> [c]) -> [a] -> [c]
```

```
pipe(filter(e => e % 2 == 0),  
  map(e => e * 2), [1,2,3,4]);  
  
// [4, 8]
```

Encadenamiento

pipe :: ([a] -> [b]) -> ([b] -> [c]) -> [a] -> [c]

Caso de estudio

UI

[1,2,3,4]

$e \Rightarrow e * 2$

T1

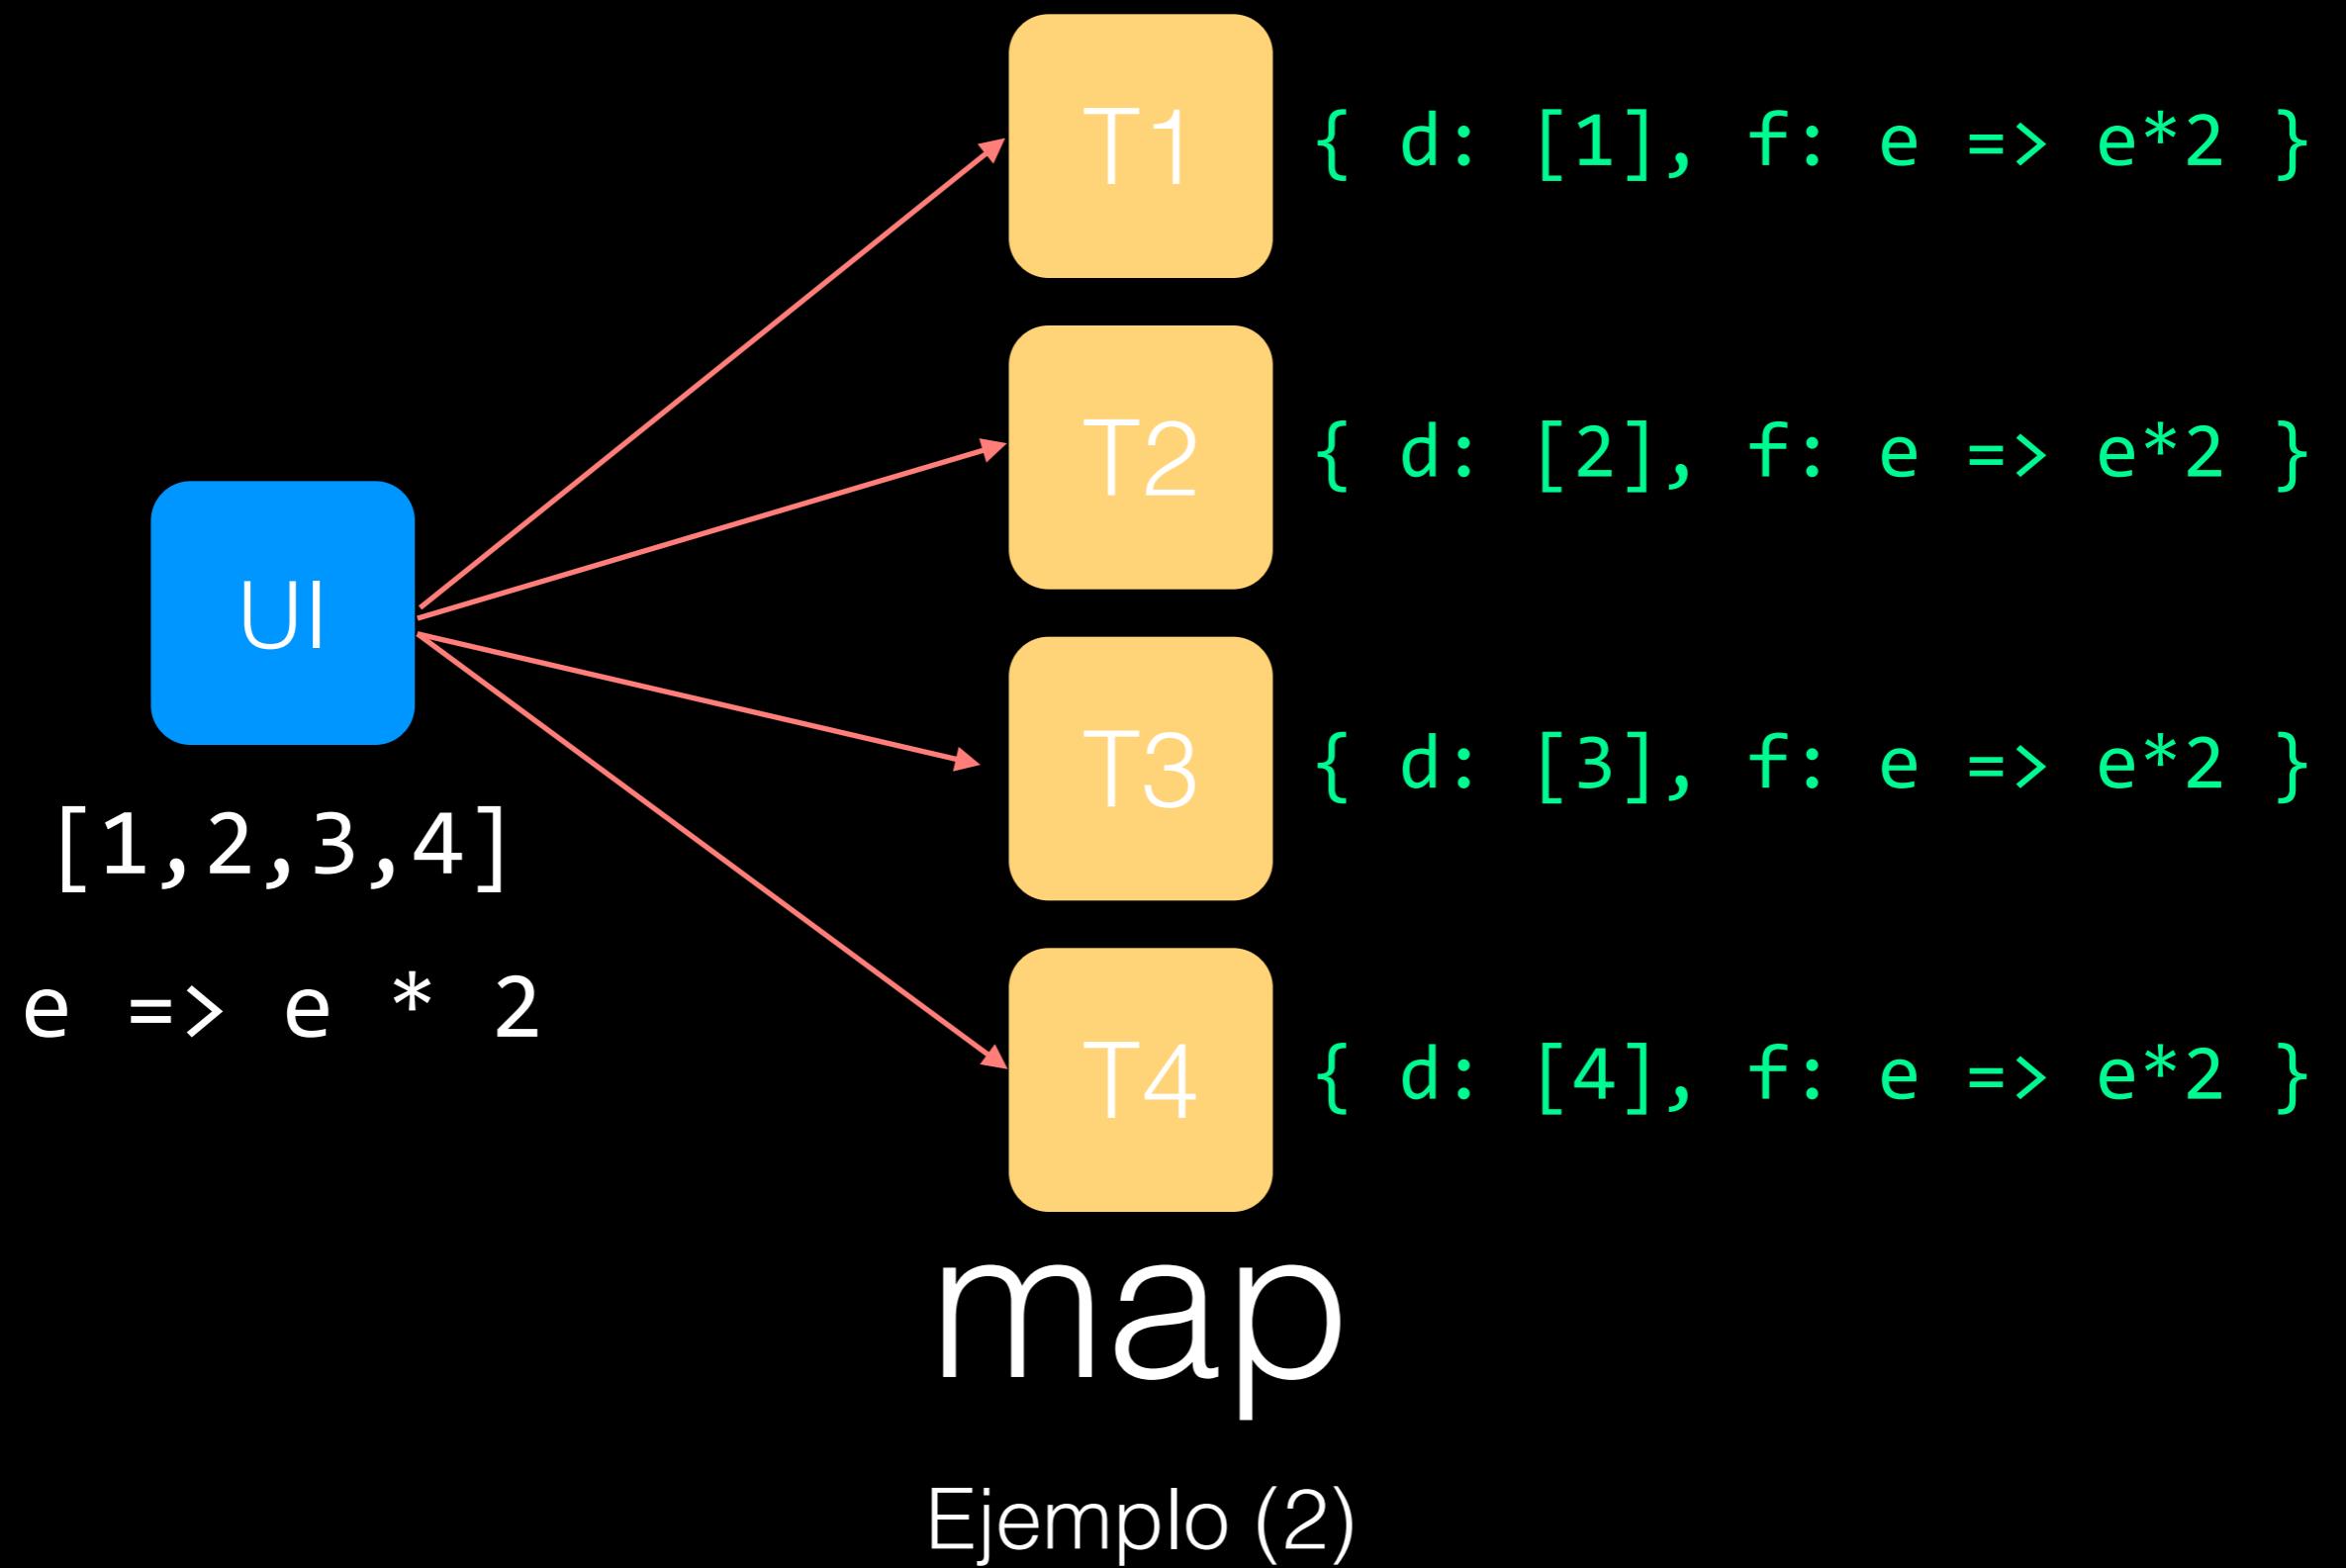
T2

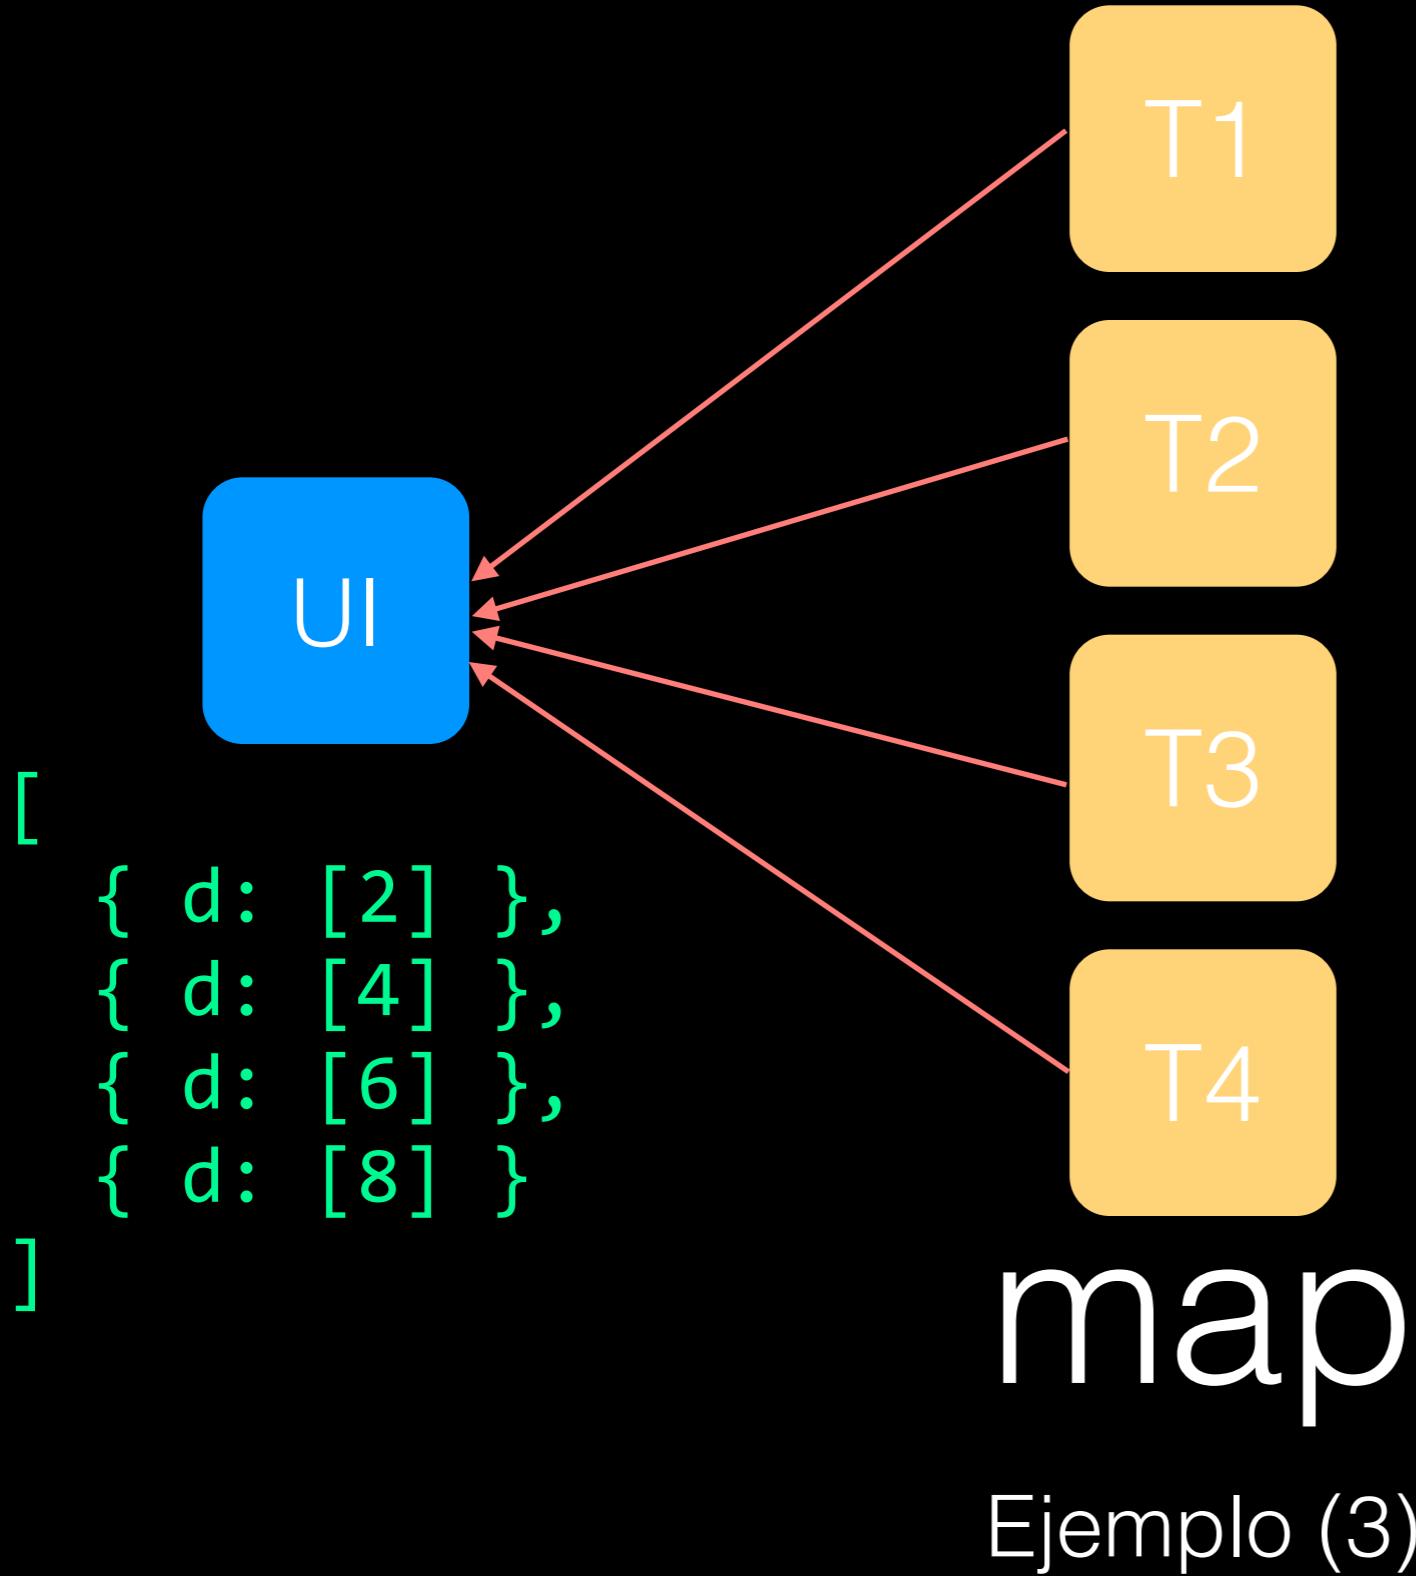
T3

T4

map

Ejemplo





UI

[2,4,6,8]

T1

T2

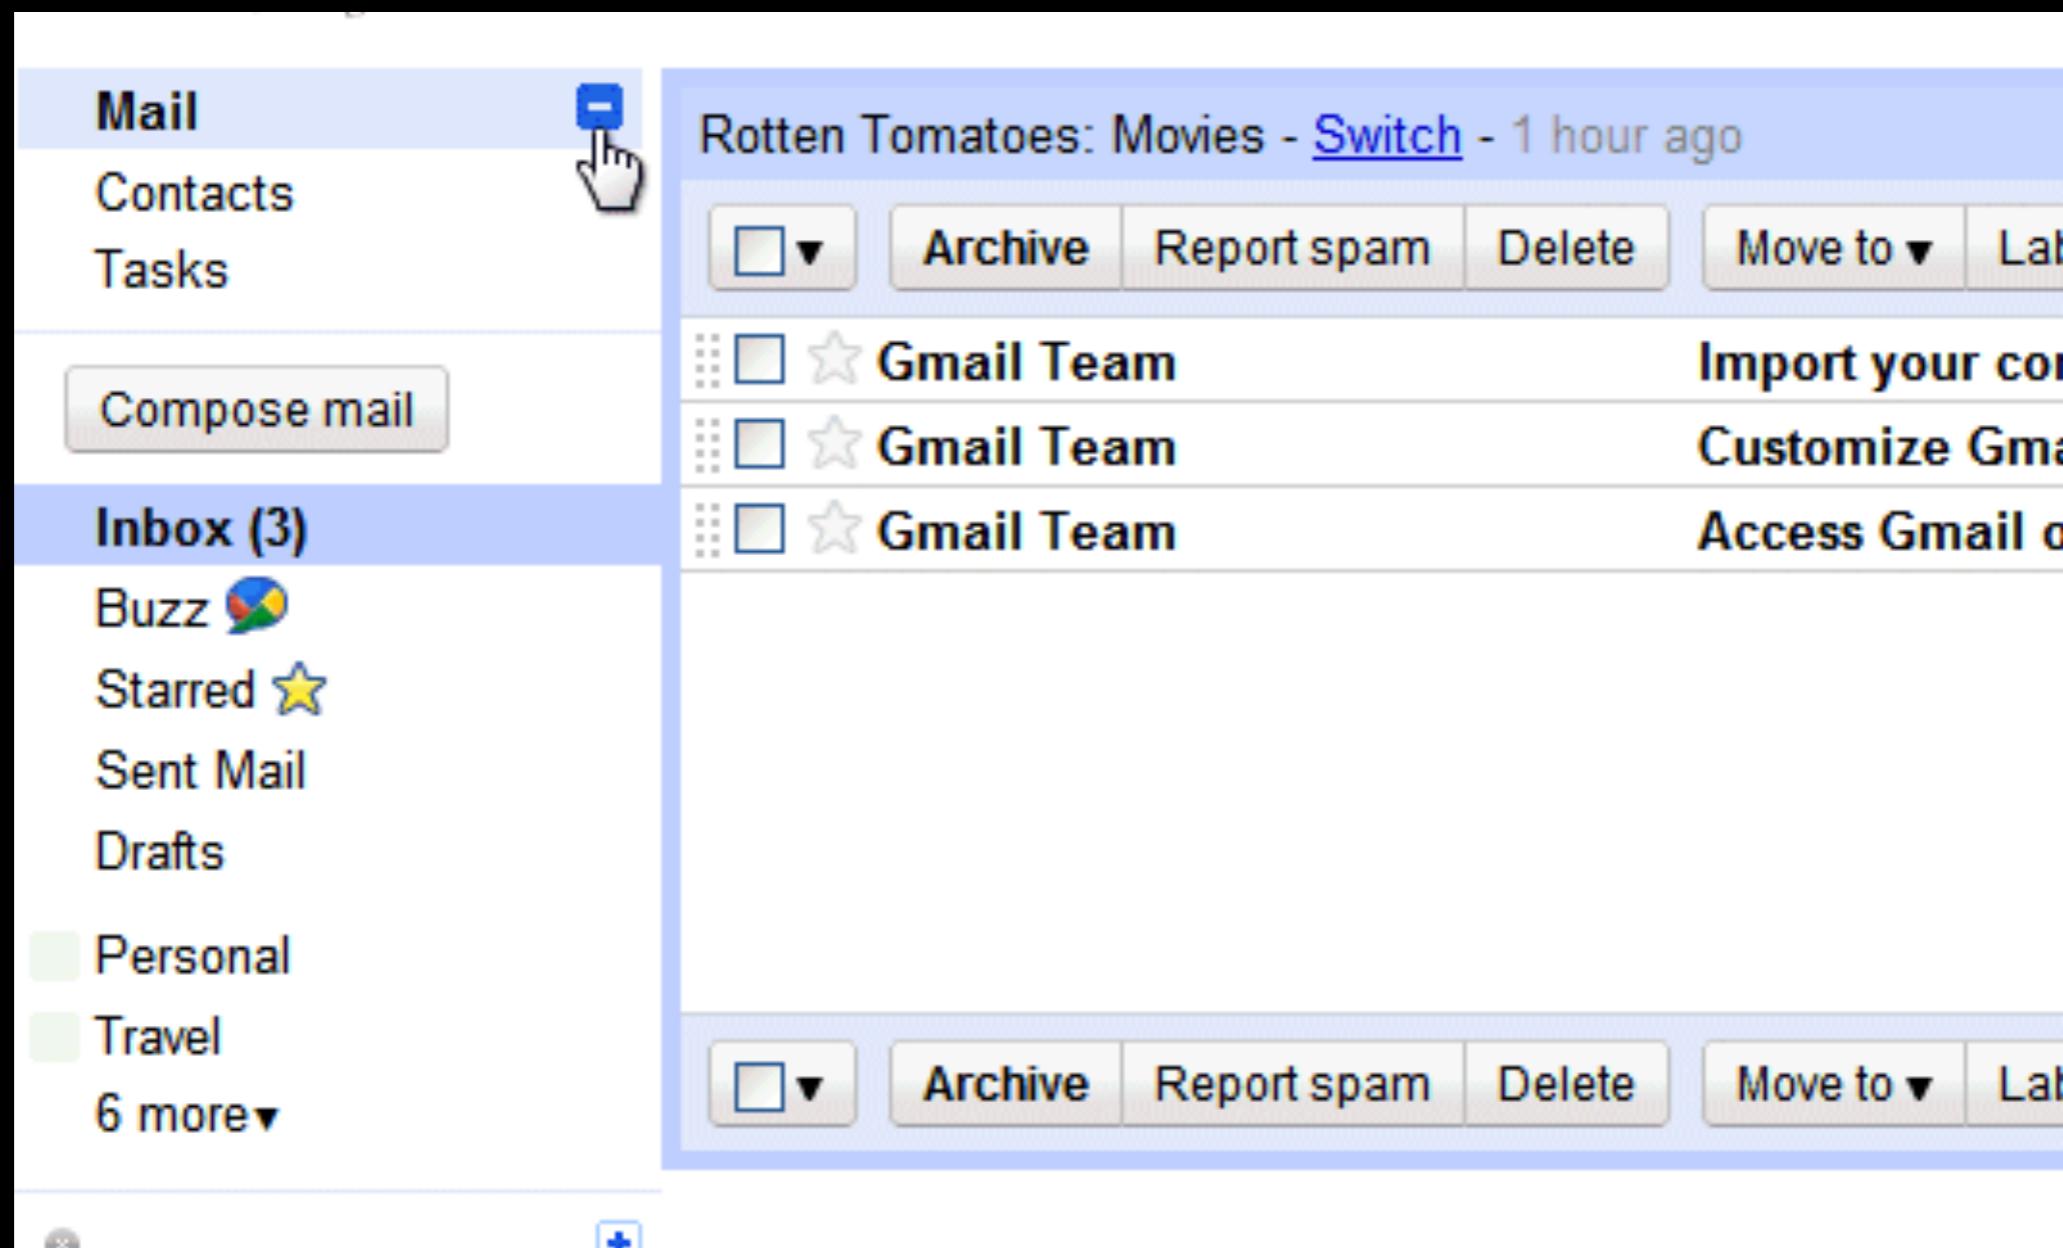
T3

T4

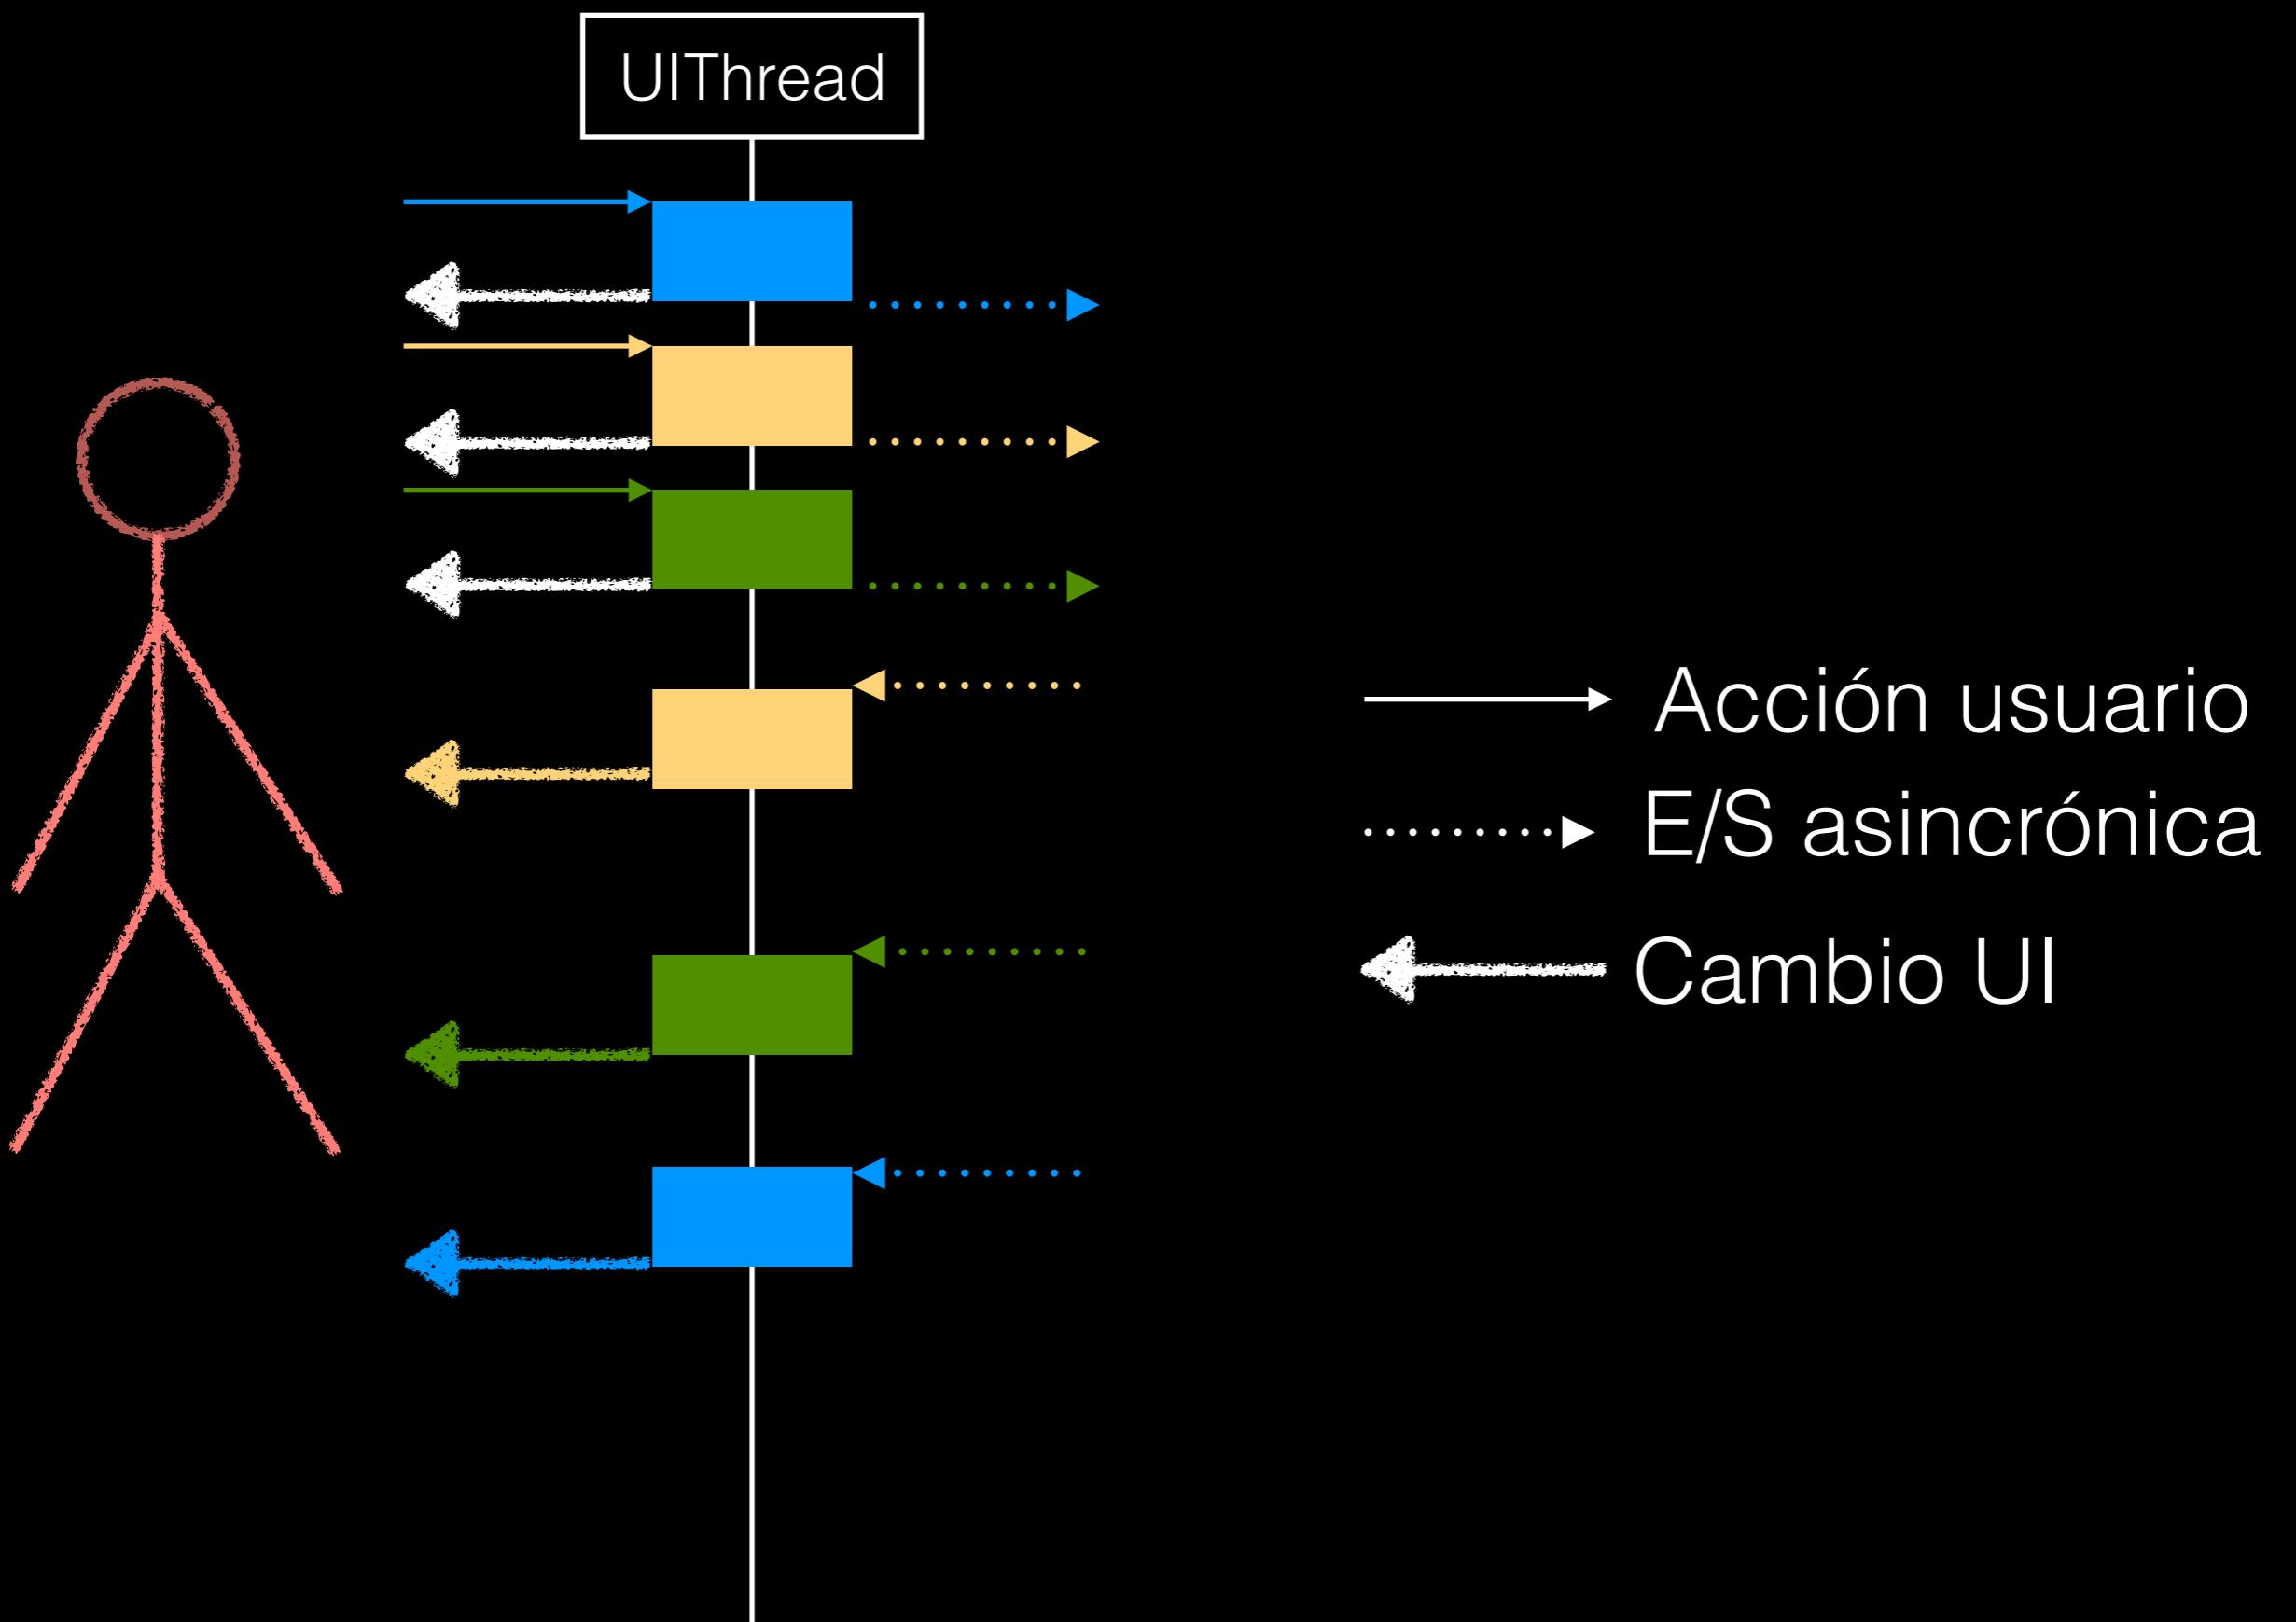
map

Ejemplo (4)

JavaScript

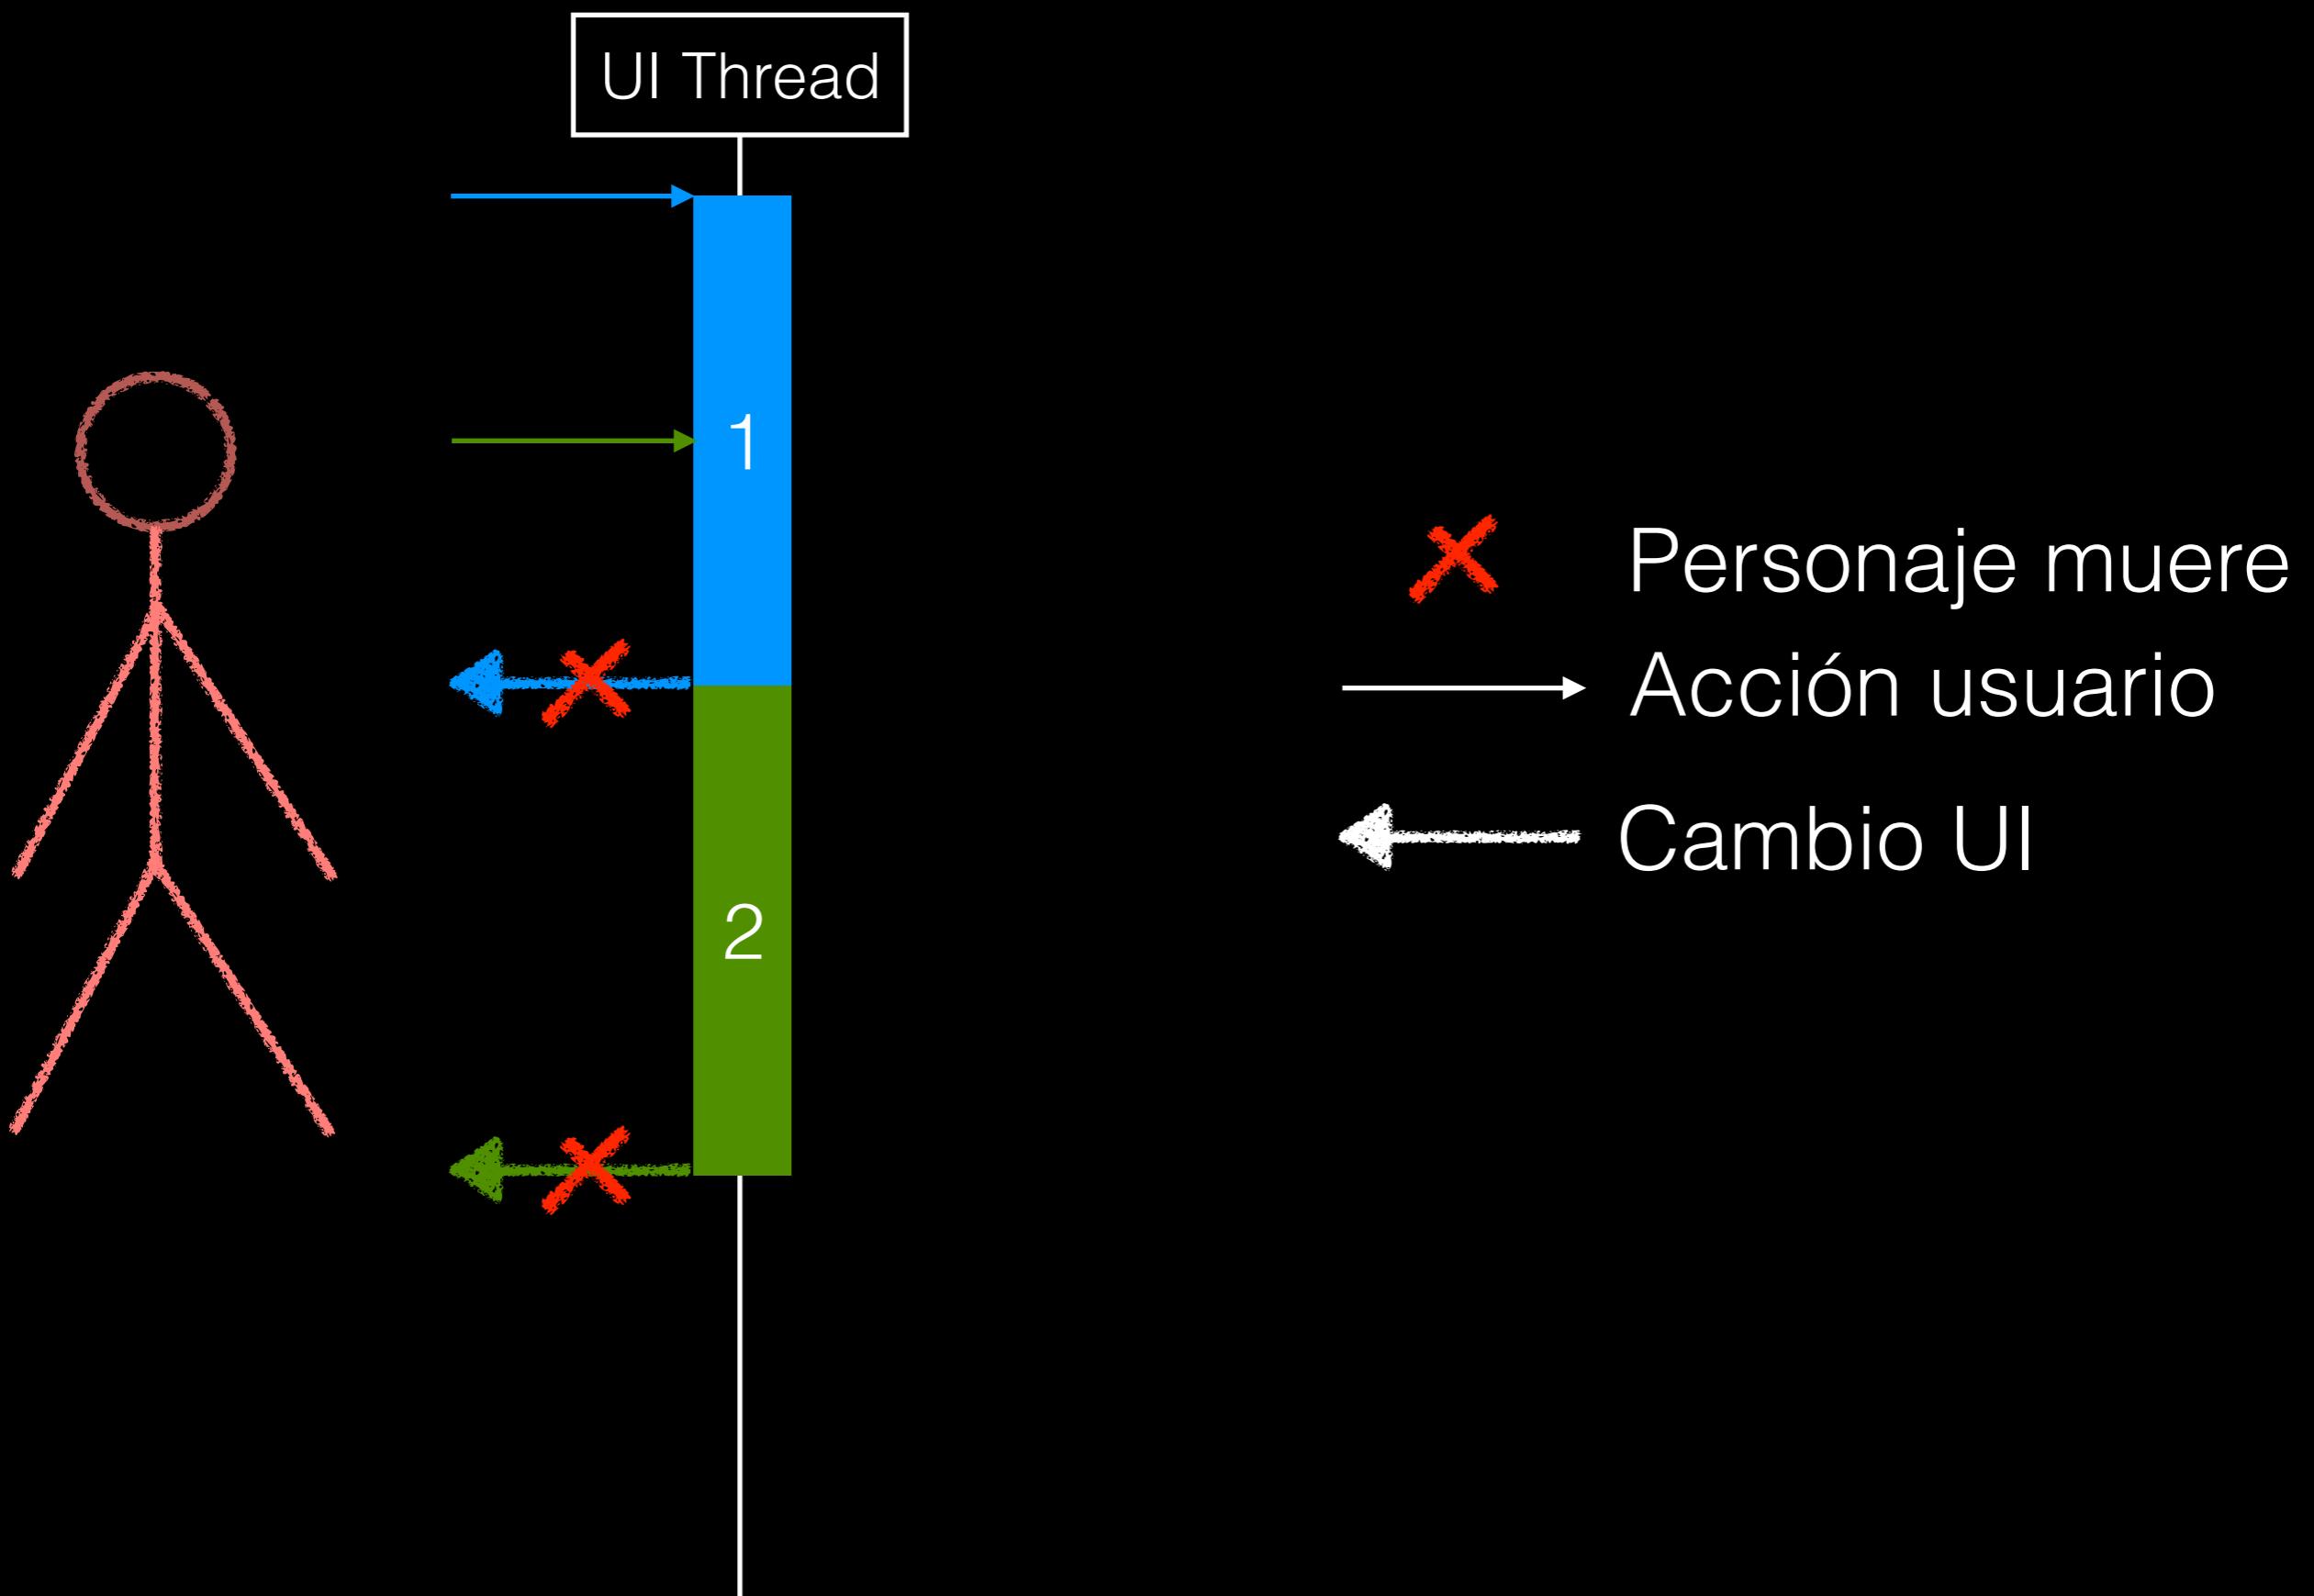


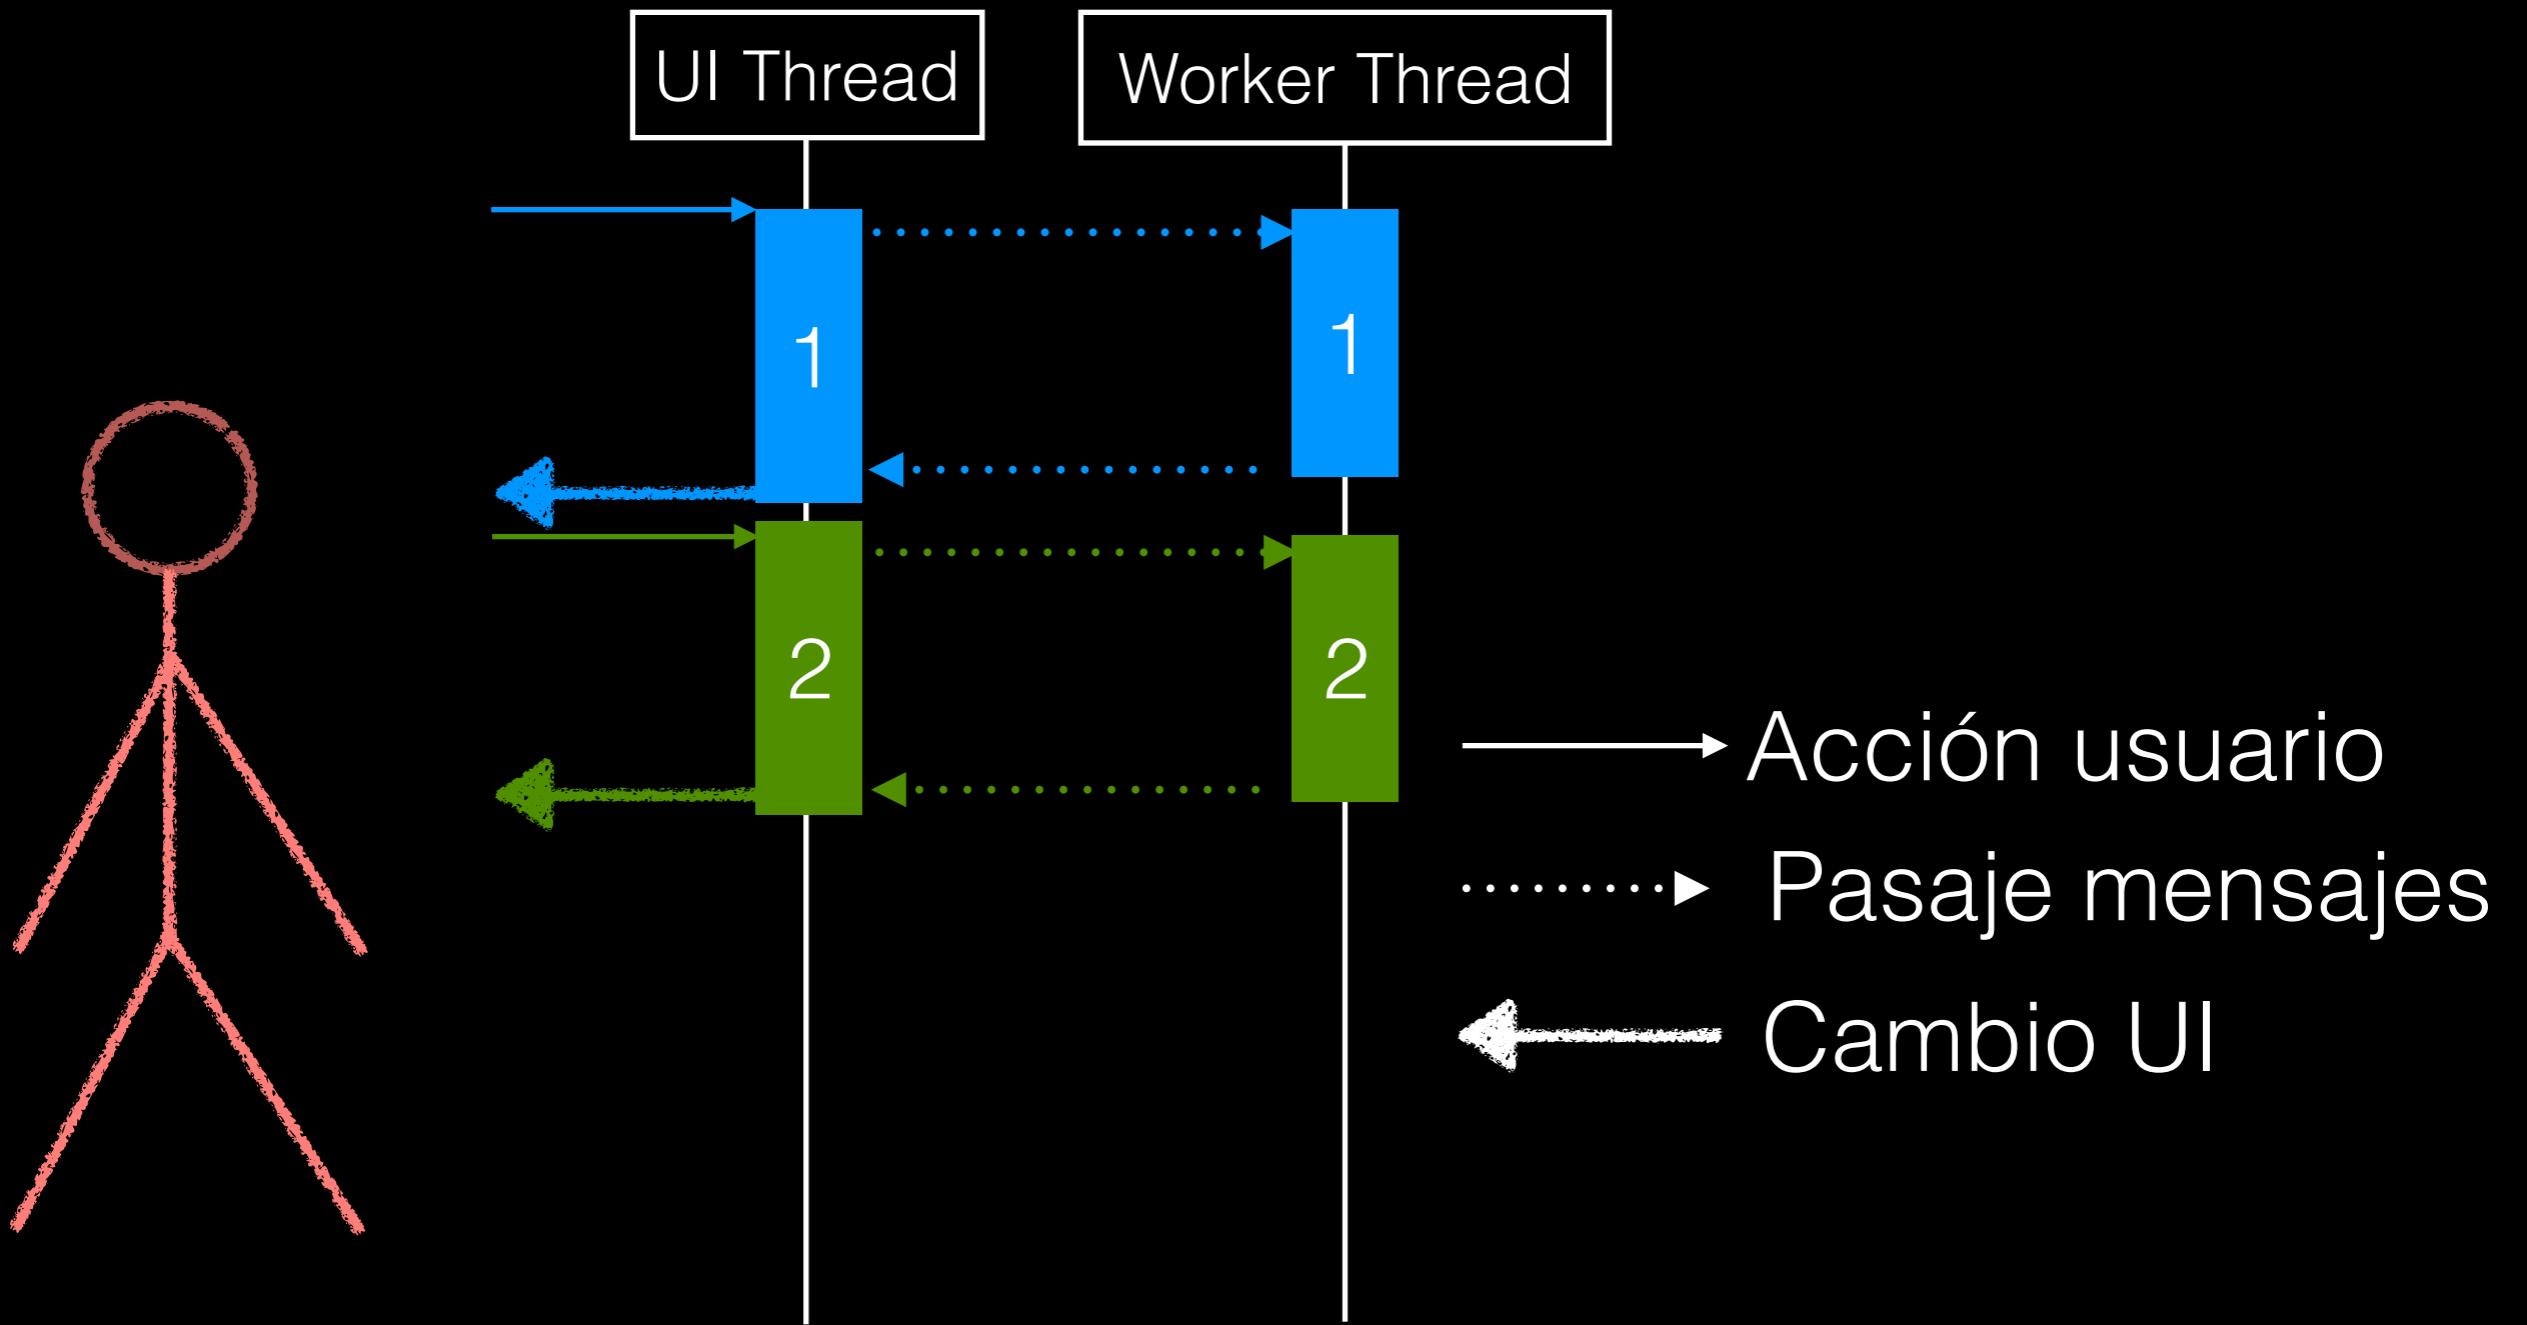
GMail (2010)





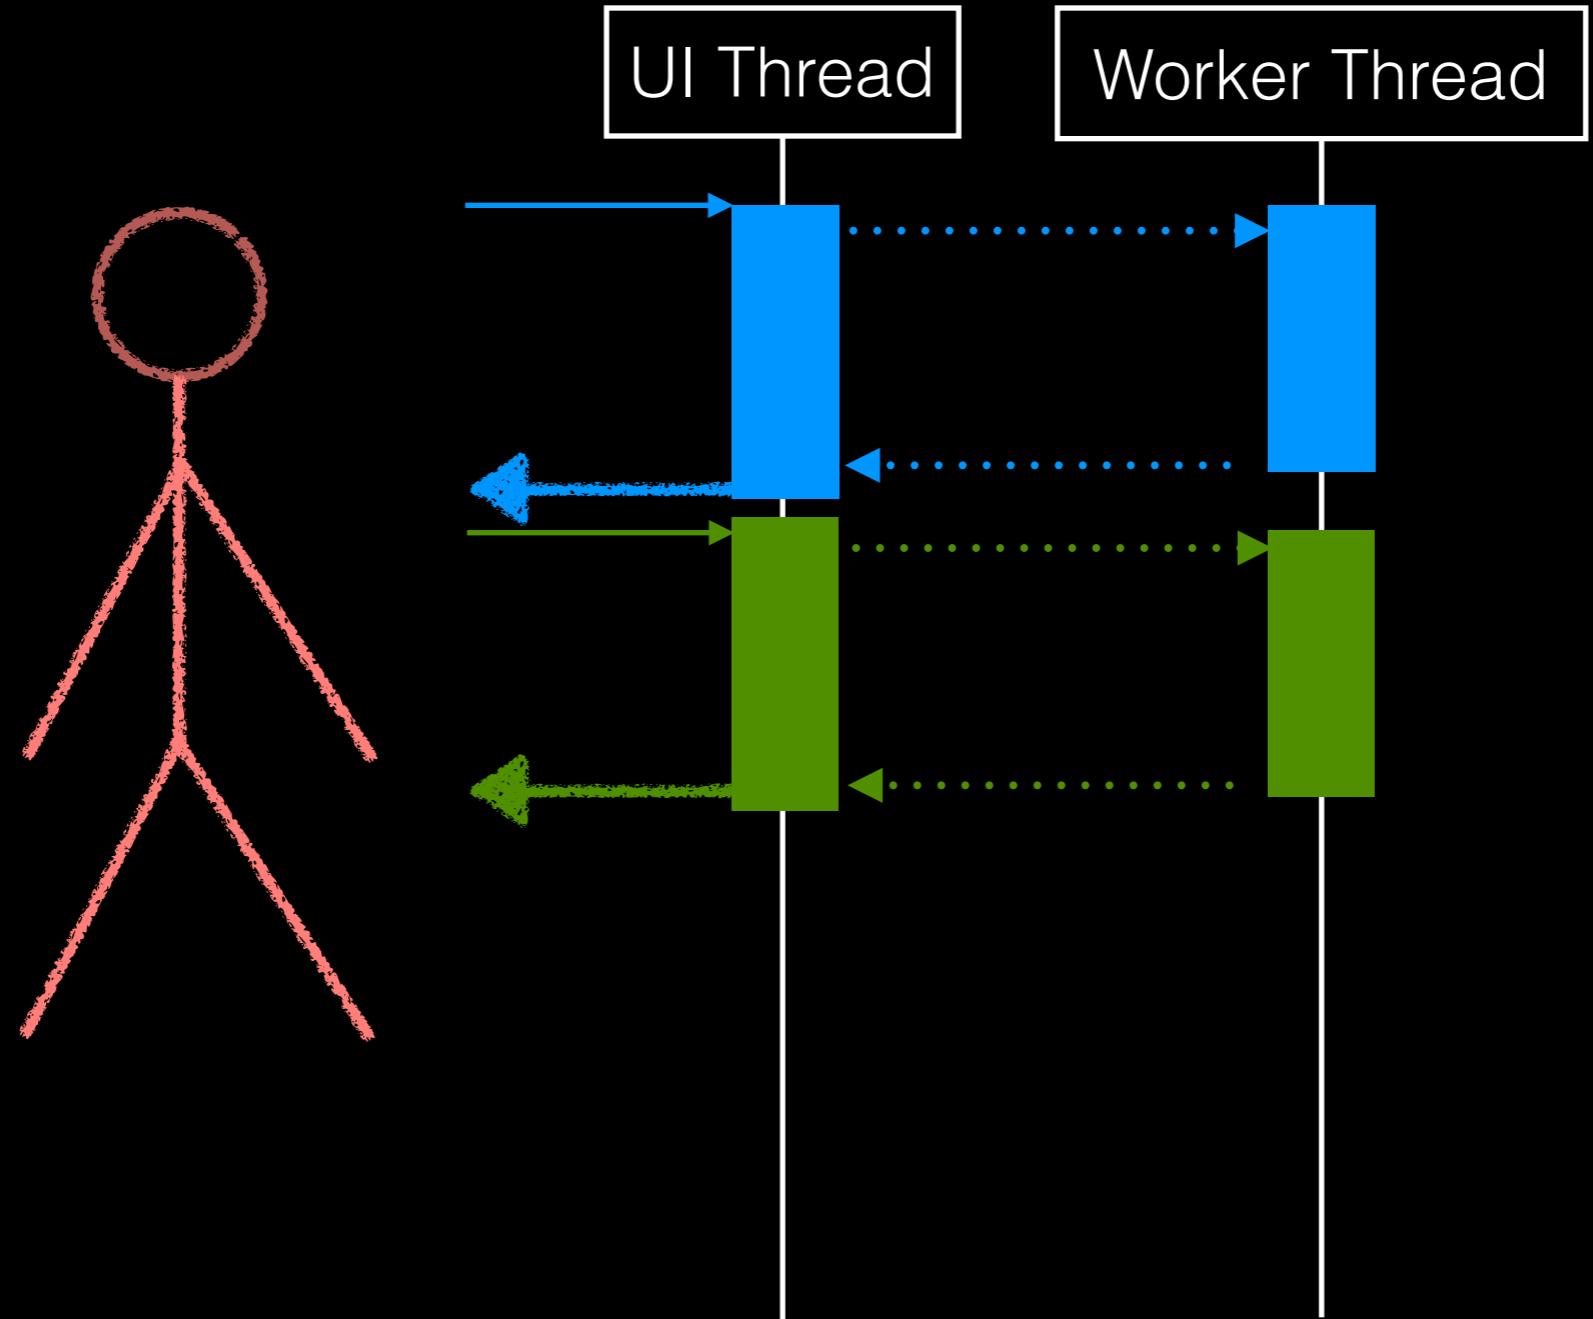
Unreal Engine en Firefox con asm.js





Dividir el trabajo

Paralelización



Web Workers

index.html

```
const w = new Worker('worker.js');
w.addEventListener('message', function(e){
  console.log('Recibido del worker', e.data);
  worker.terminate();
})
worker.postMessage('Hola mundo');
```

worker.js

```
self.addEventListener('message', function(e) {
  const data = e.data;
  self.postMessage(data + '!!!!');
});
```

Uso básico

Soportado

- XHR
- importScripts(scriptName)
- Crear otros workers

No soportado

- DOM
- window
- document
- parent

Sintaxis

```
var xs = [1,2,3,4];
```

```
var dobles = xs.map(x => x * 2);
```

```
var pares = xs.filter(x => x % 2 == 0);
```

```
var suma = xs.reduce((ac, el) => ac + el, 0);
```

JavaScript

map, reduce, filter

```
var pjs = require('p-j-s');

// cantidad de workers para el pool
pjs.init({ maxWorkers: 4 });

var xs = new Uint32Array([1,2,3,4]);

pjs(xs).map(e => e * 2).seq()
  // asíncrono
  .then(cuadrados => {
    // res [2,4,6,8]
  });

```

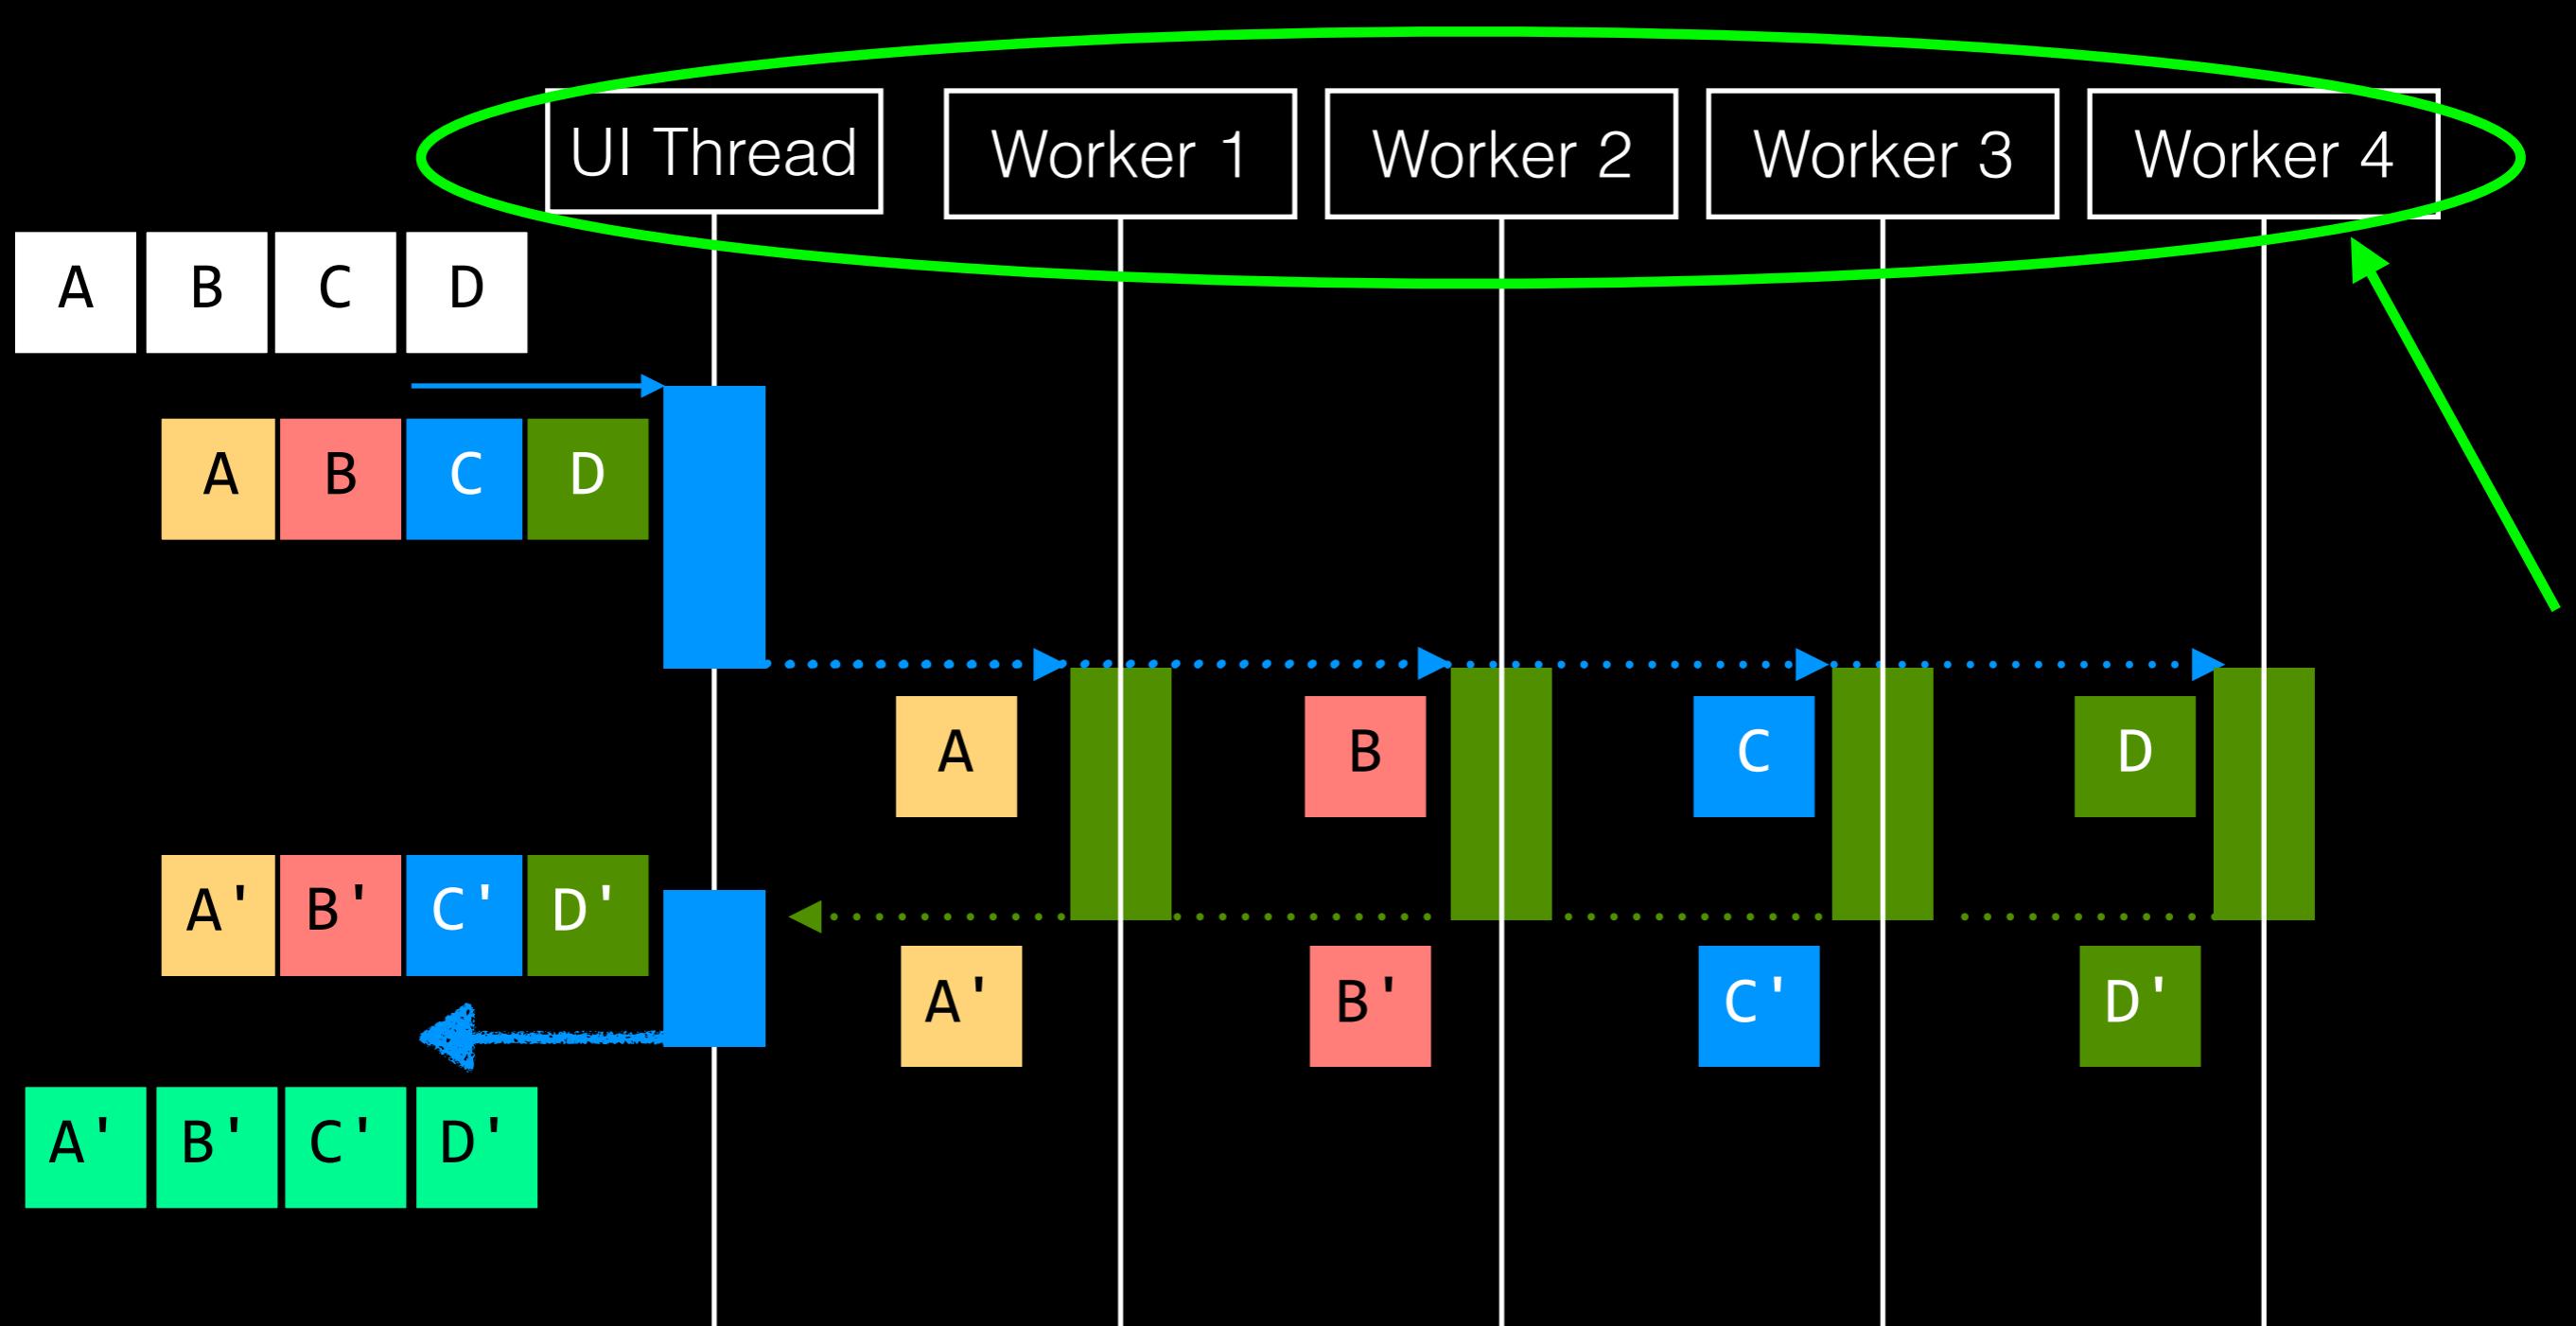
p-j-S

Básico

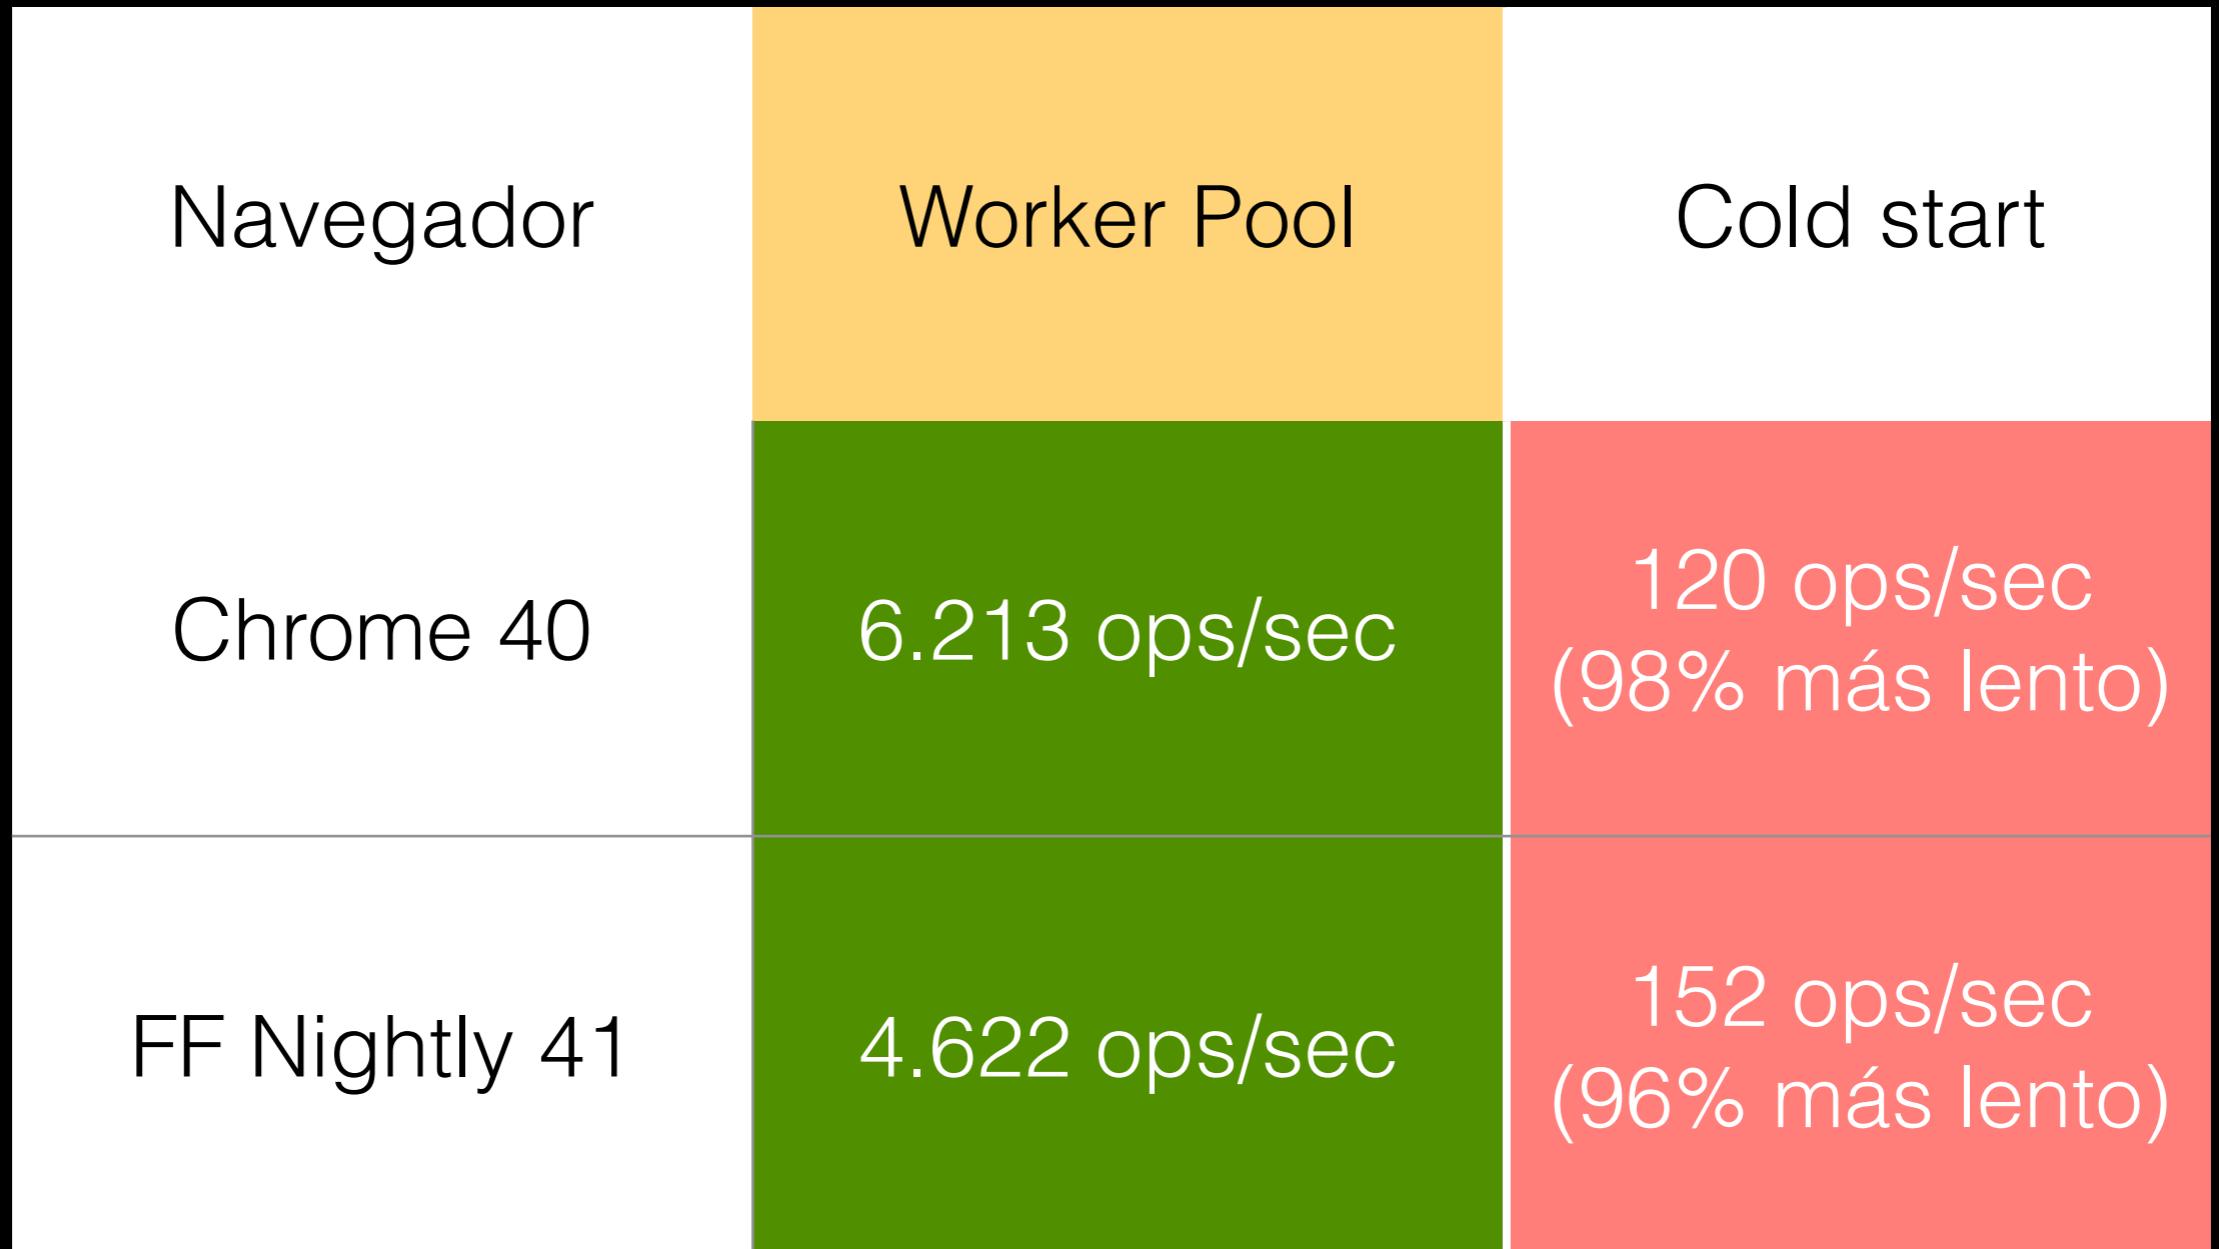
```
var xs = [1,2,3,4];          var xs = new Uint32Array([1,2,3,4]);  
  
var ys = xs                  pjs(xs)  
  // paso 1                  // paso 1  
  .filter(x => x % 2 == 0)  .filter(x => x % 2 == 0)  
  // paso 2                  // paso 2  
  .map(x => x * 2);        .map(x => x * 2)  
                            // terminar cadena  
// ys [4,8]                  .seq().then(ys => {  
                           // ys [4,8]  
                           });
```

Comparación

Inicialización



p-j-S



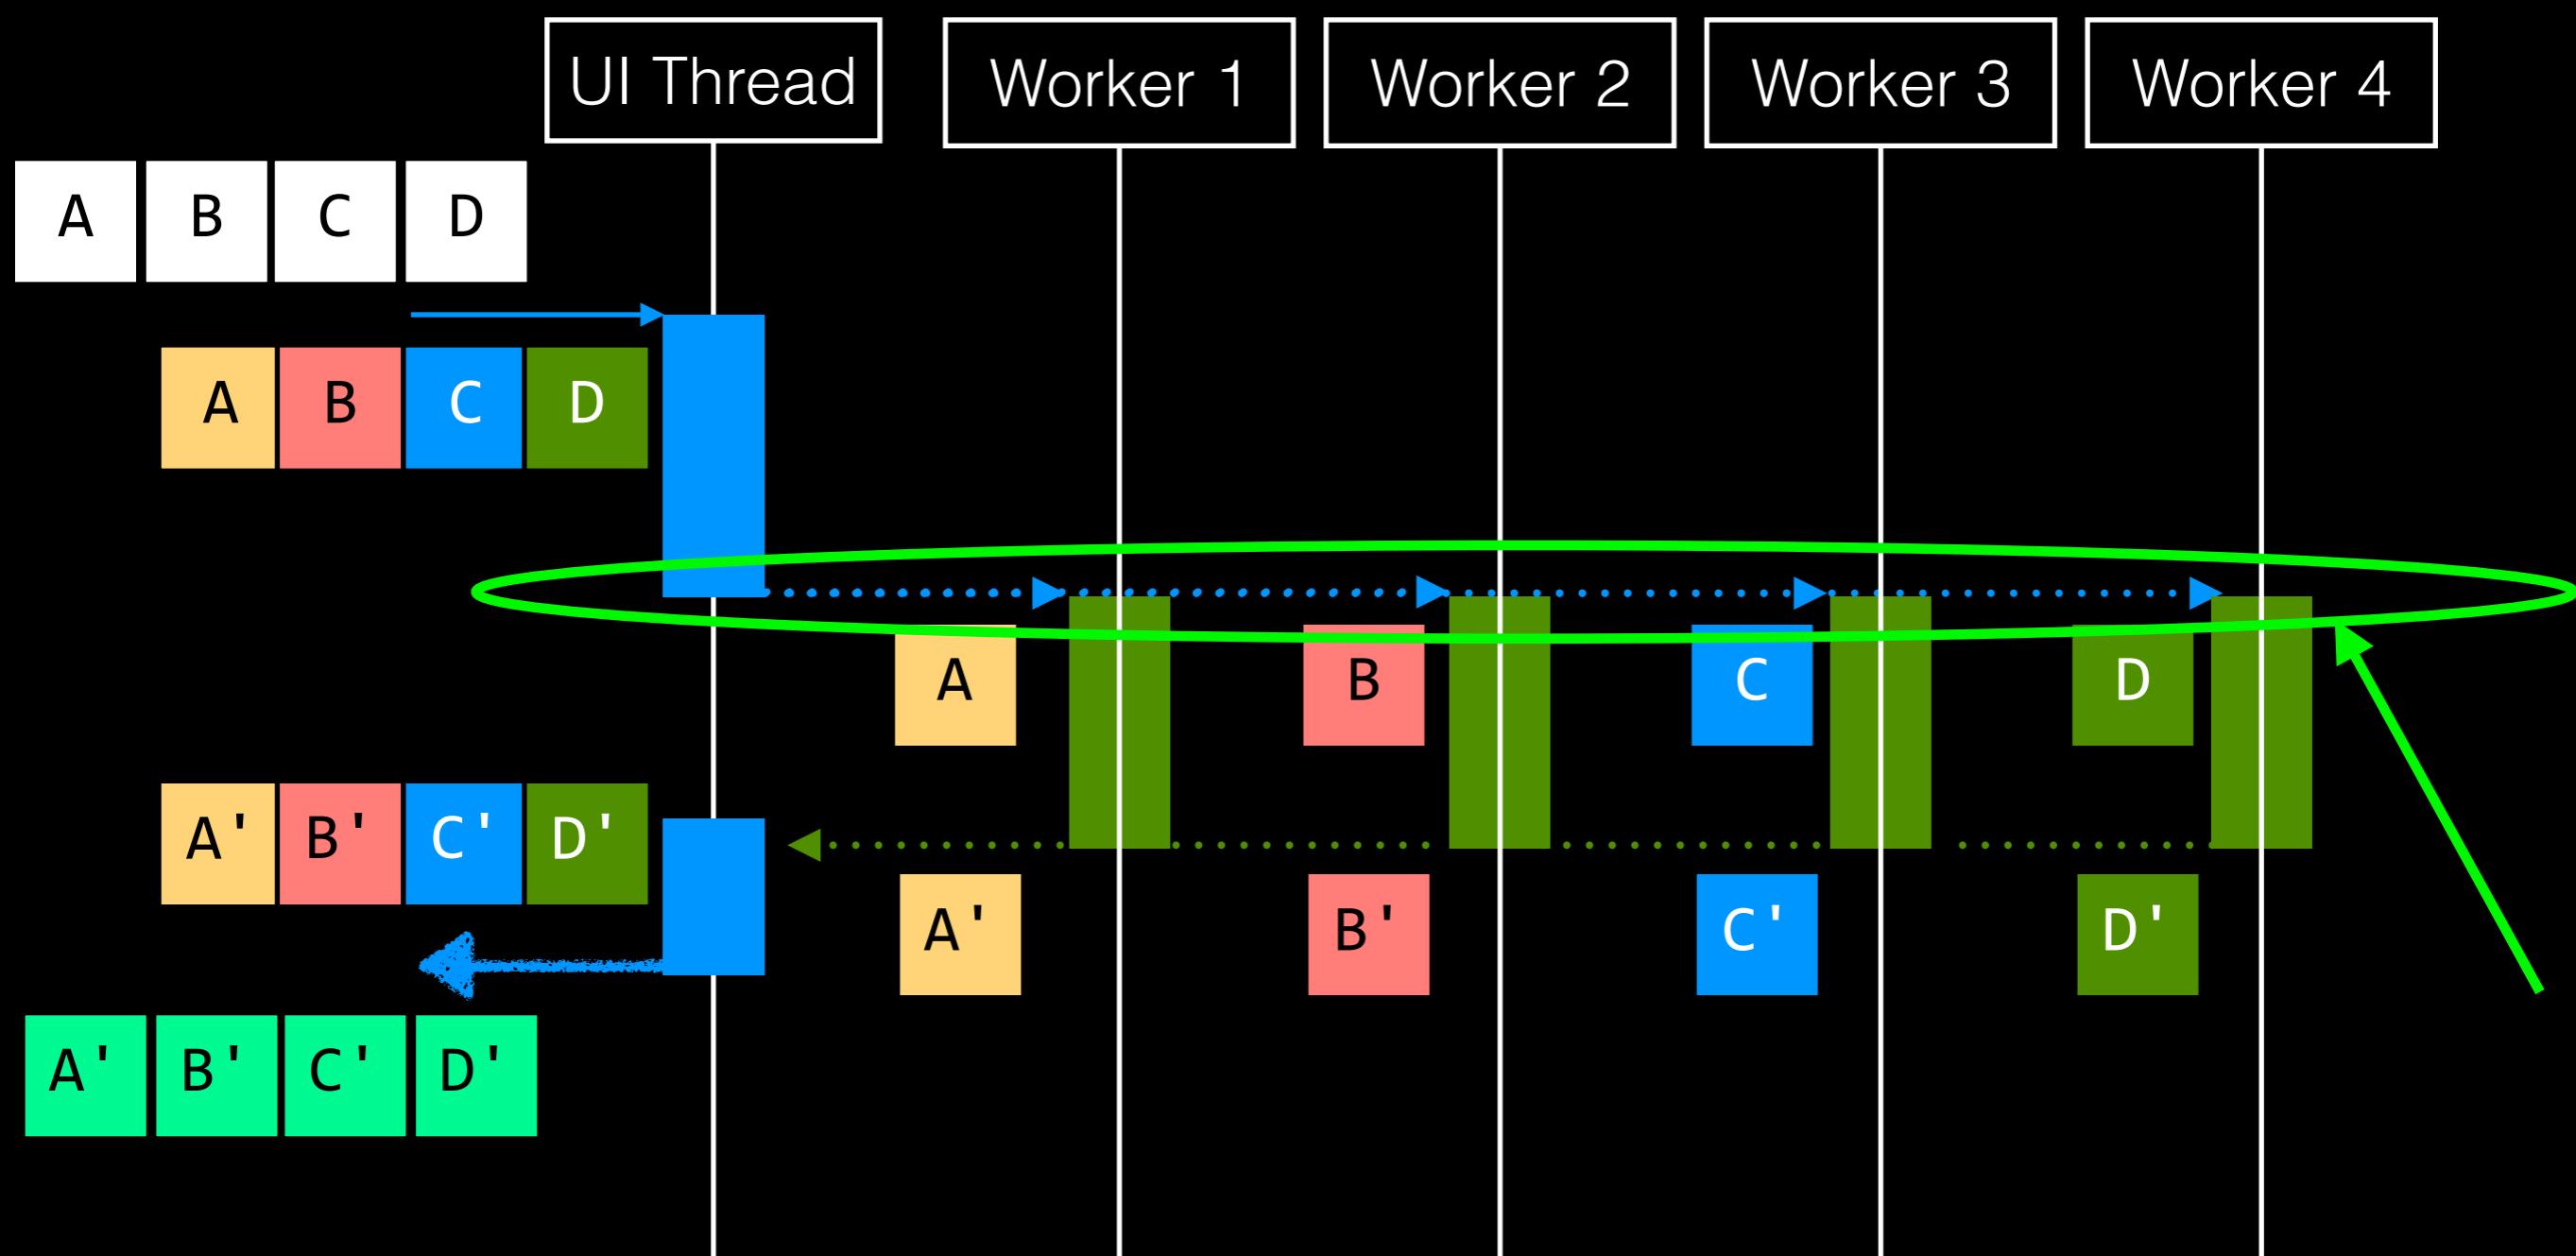
Creación de Workers

<http://jsperf.com/worker-cold-start/6>

Cantidad de Workers

- Memoria consumida por instancia
- "Intercalación" (*interleaving*)

Comunicación de datos



p-j-S

$$T_{ser} = \sum_{i=0}^N t = Nt$$

T_{ser} : Tiempo en serie

N : Cantidad de elementos

t : Tiempo promedio por elemento

Tiempo en serie

$$T_{par} = \frac{T_{ser}}{K}$$

T_{ser} : Tiempo en serie

T_{par} : Tiempo en paralelo

K : Cantidad de Workers

Tiempo en paralelo

Ideal

$$T_{sync} = T_t + T_s + T_p$$

$$T_{par} \approx \frac{T_{ser}}{K} + T_{sync}$$

T_t : Tiempo de transferencia

T_s : Tiempo de serializacion/deserializacion

T_p : Tiempo de particion/union

Tiempo en paralelo

Real

spec

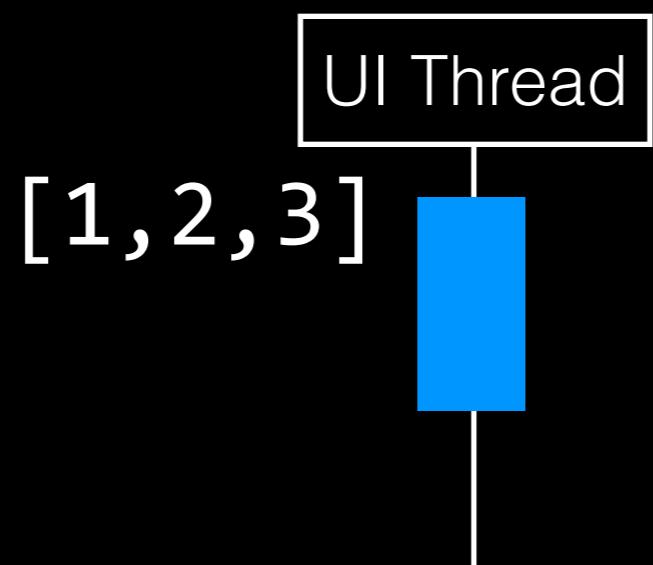
```
void postMessage(any message ,  
optional sequence < Transferable > transfer)
```

ejemplo

```
const worker = new Worker('worker.js');  
...  
const items = new Uint8Array([1,2,3,4]);  
worker.postMessage(items, [ items.buffer ]);
```

postMessage

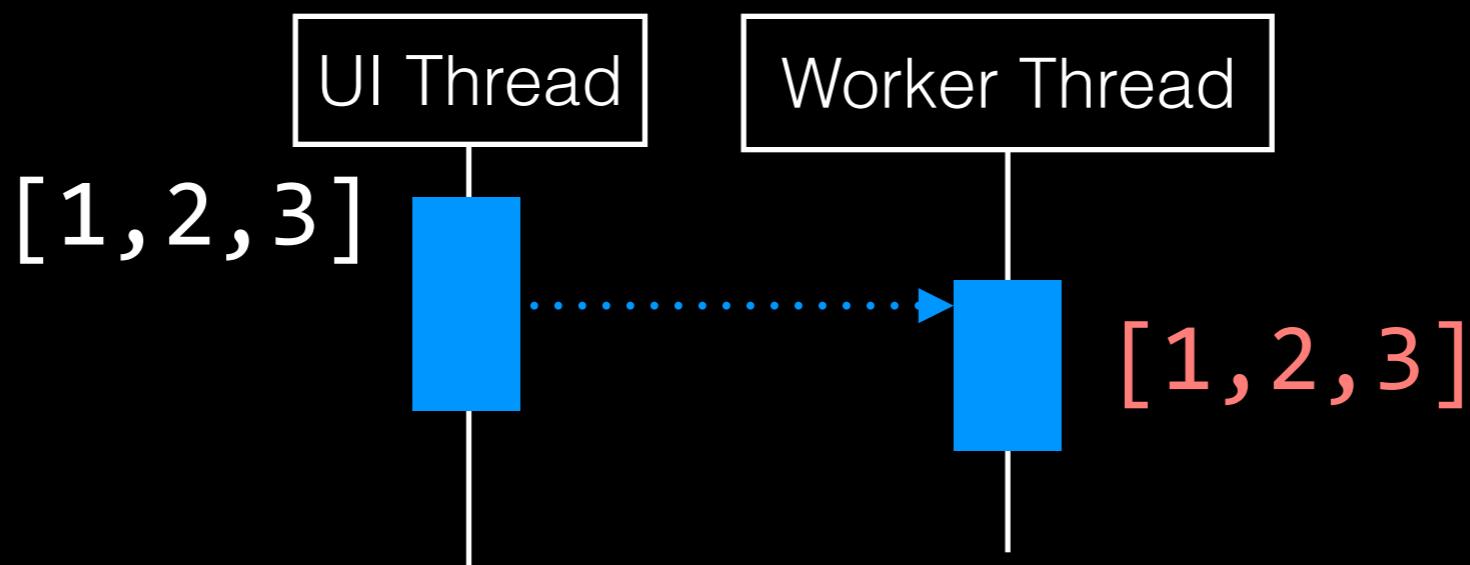
```
const xs = new Uint32Array([1,2,3]);
```



Cloning

No transferable

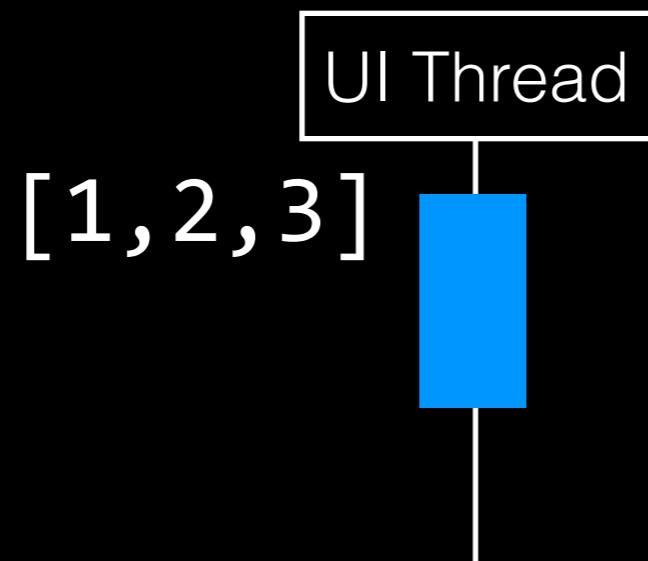
```
const xs = new Uint32Array([1,2,3]);  
worker.postMessage(xs);
```



Cloning (2)

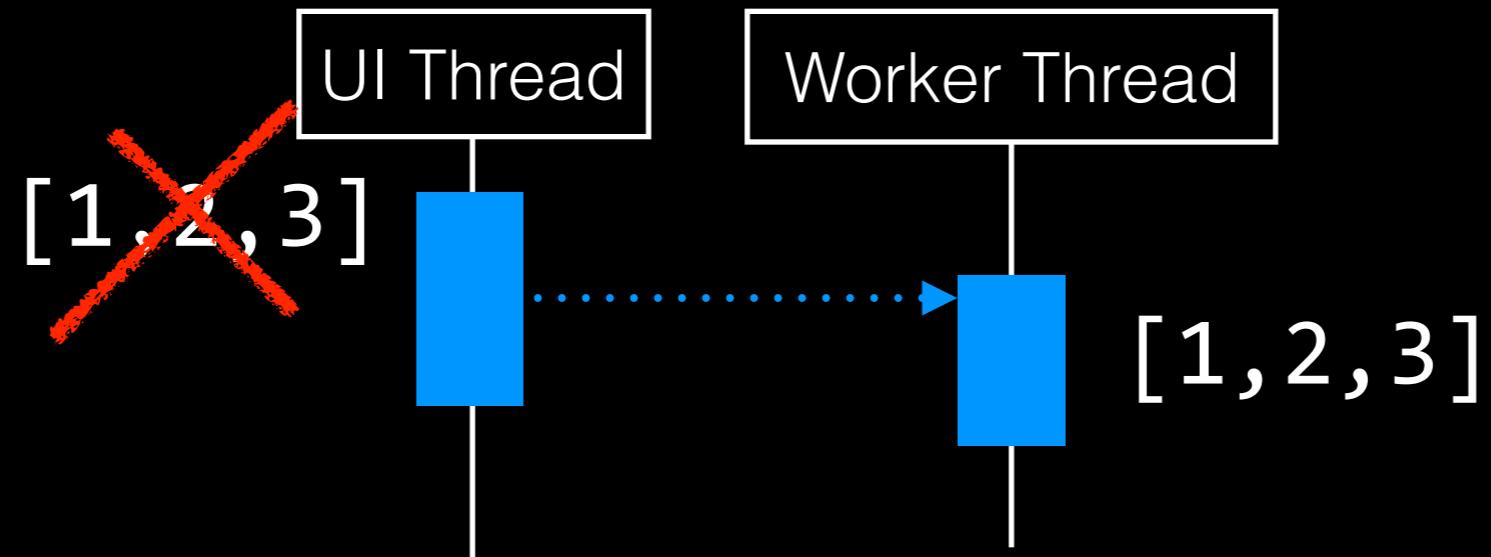
No transferable

```
const xs = new Uint32Array([1,2,3]);
```

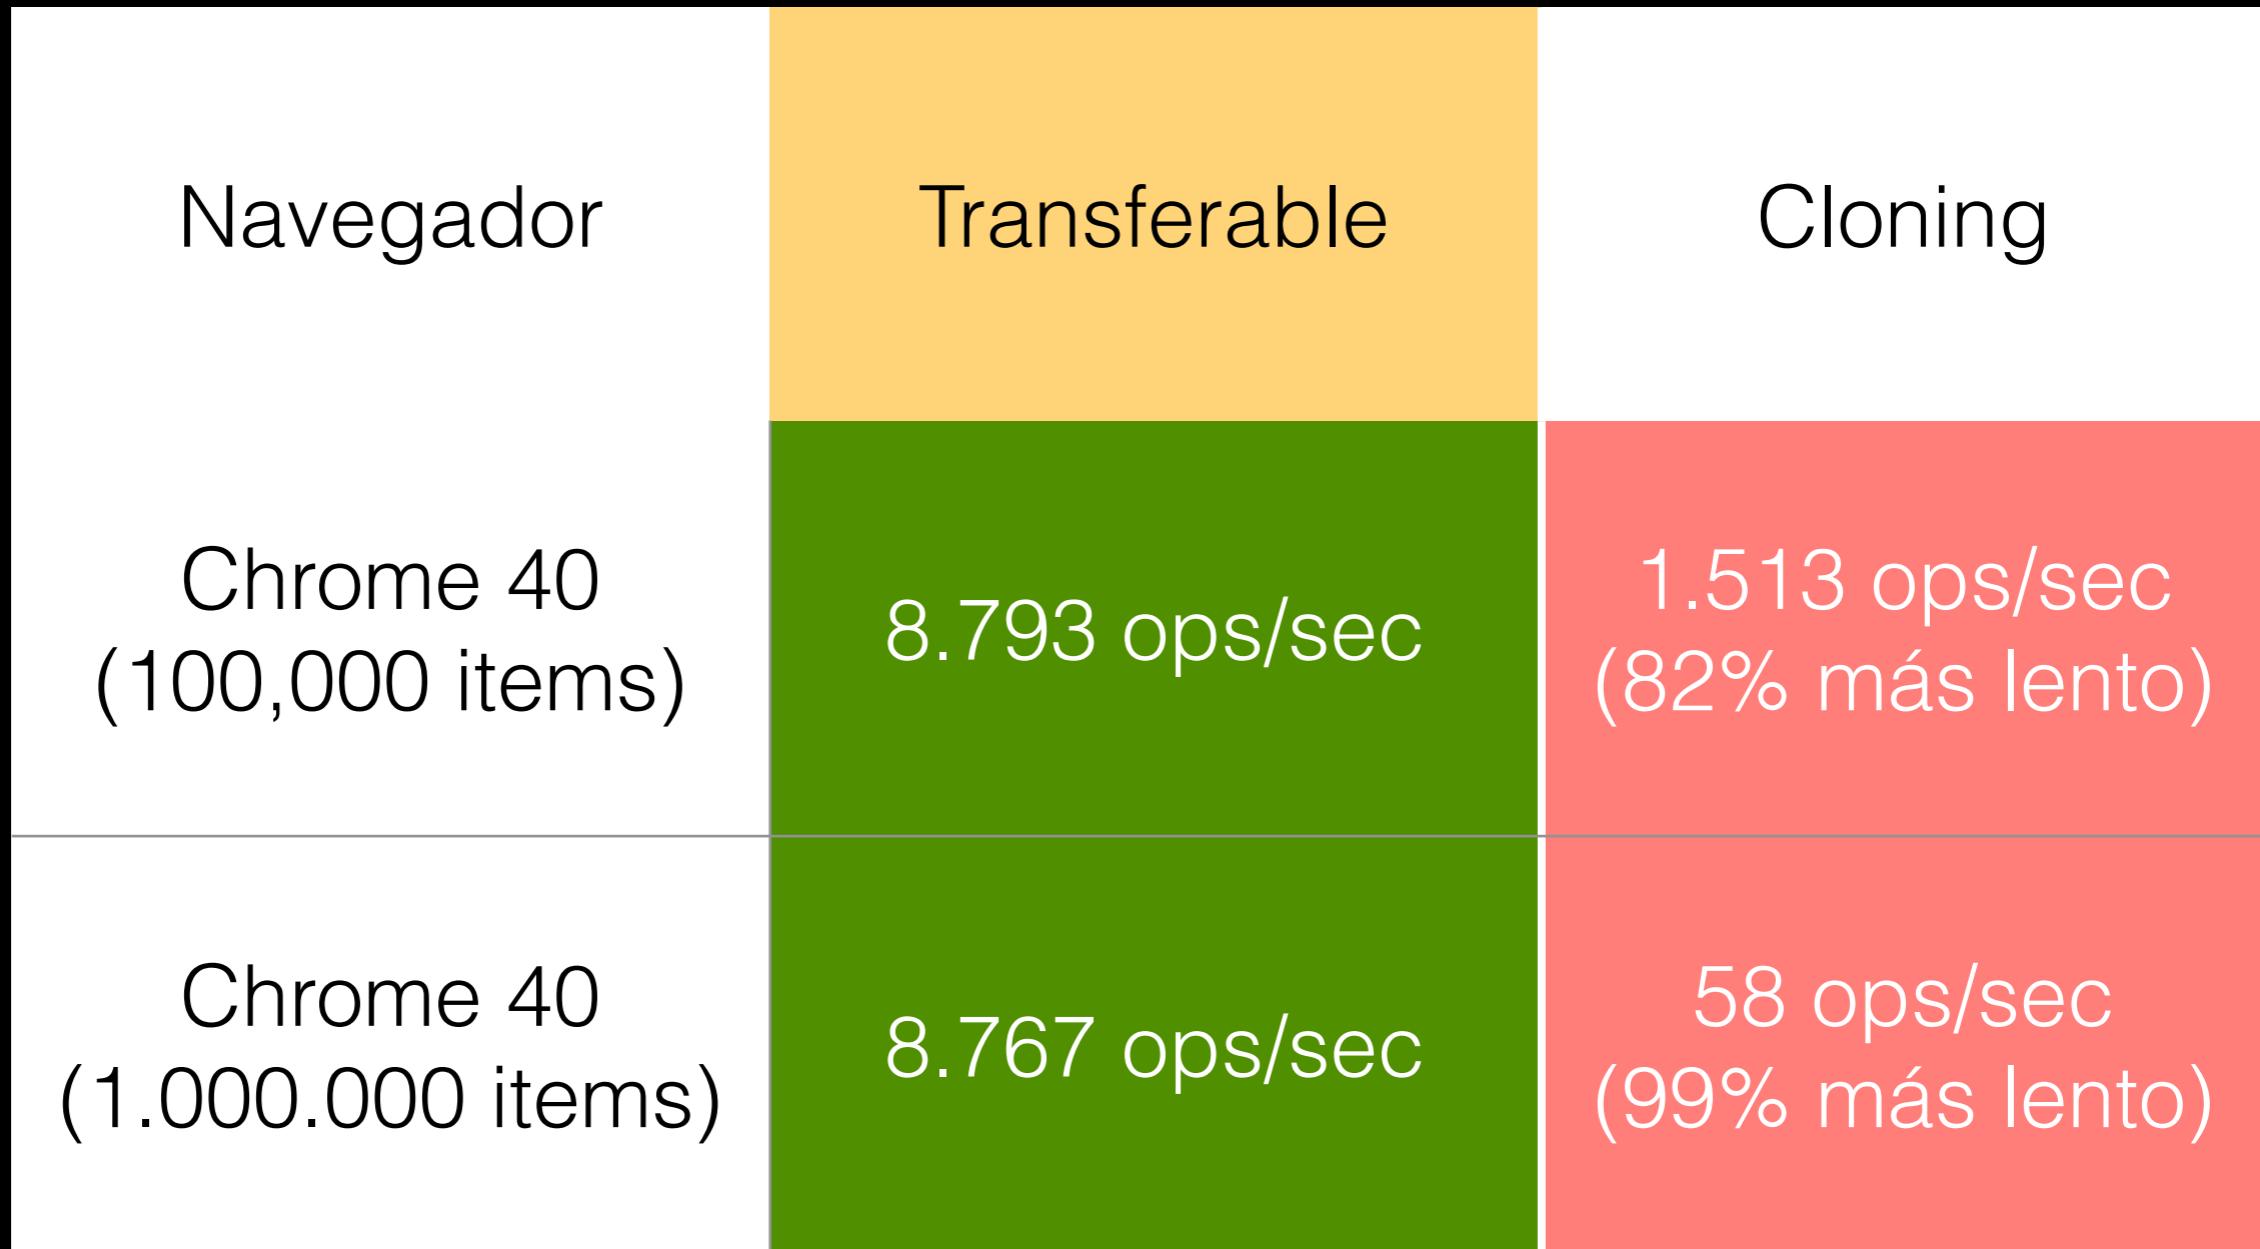


Transferables

```
const xs = new Uint32Array([1,2,3]);  
worker.postMessage(xs, [ xs.buffer ]);
```



Transferables(2)



TypedArrays

<http://jsperf.com/transferrable-vs-cloning/3>

<http://jsperf.com/longer-transferrable-vs-cloning/3>

Navegador

SharedTypedArray

TypedArray

FF Nightly 41
(100.000 items)

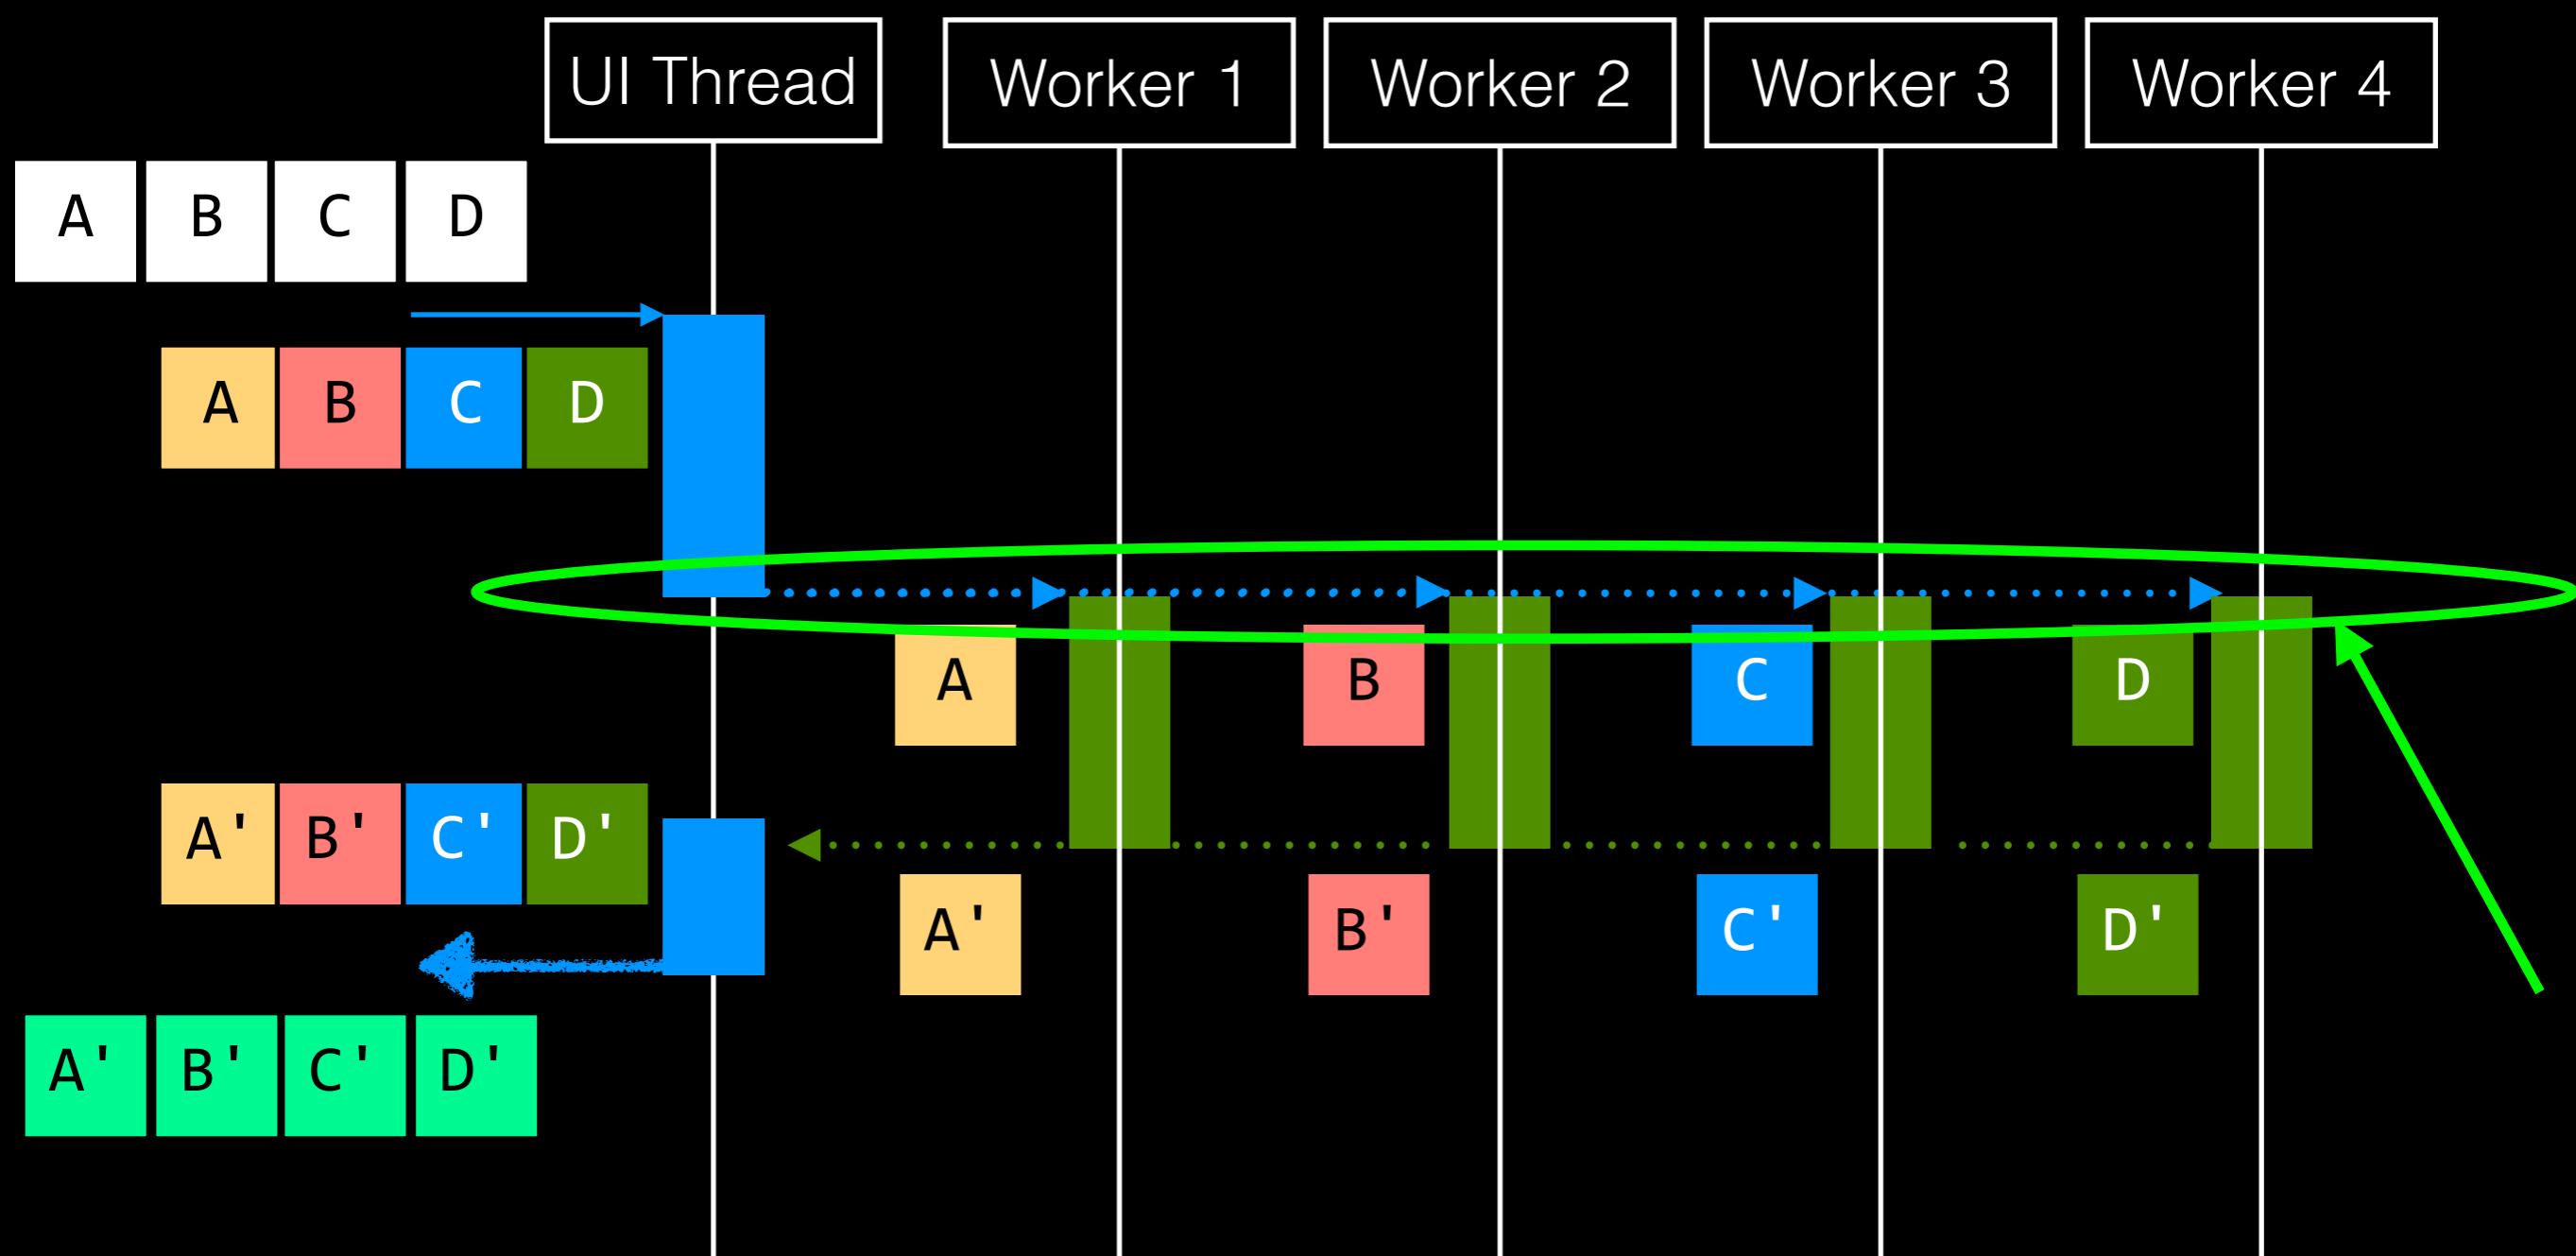
6.223 ops/sec

435 ops/sec
(93% más lento)

SharedTypedArrays

<http://jsperf.com/transfer-shared-vs-not-shared>

Comunicación de código



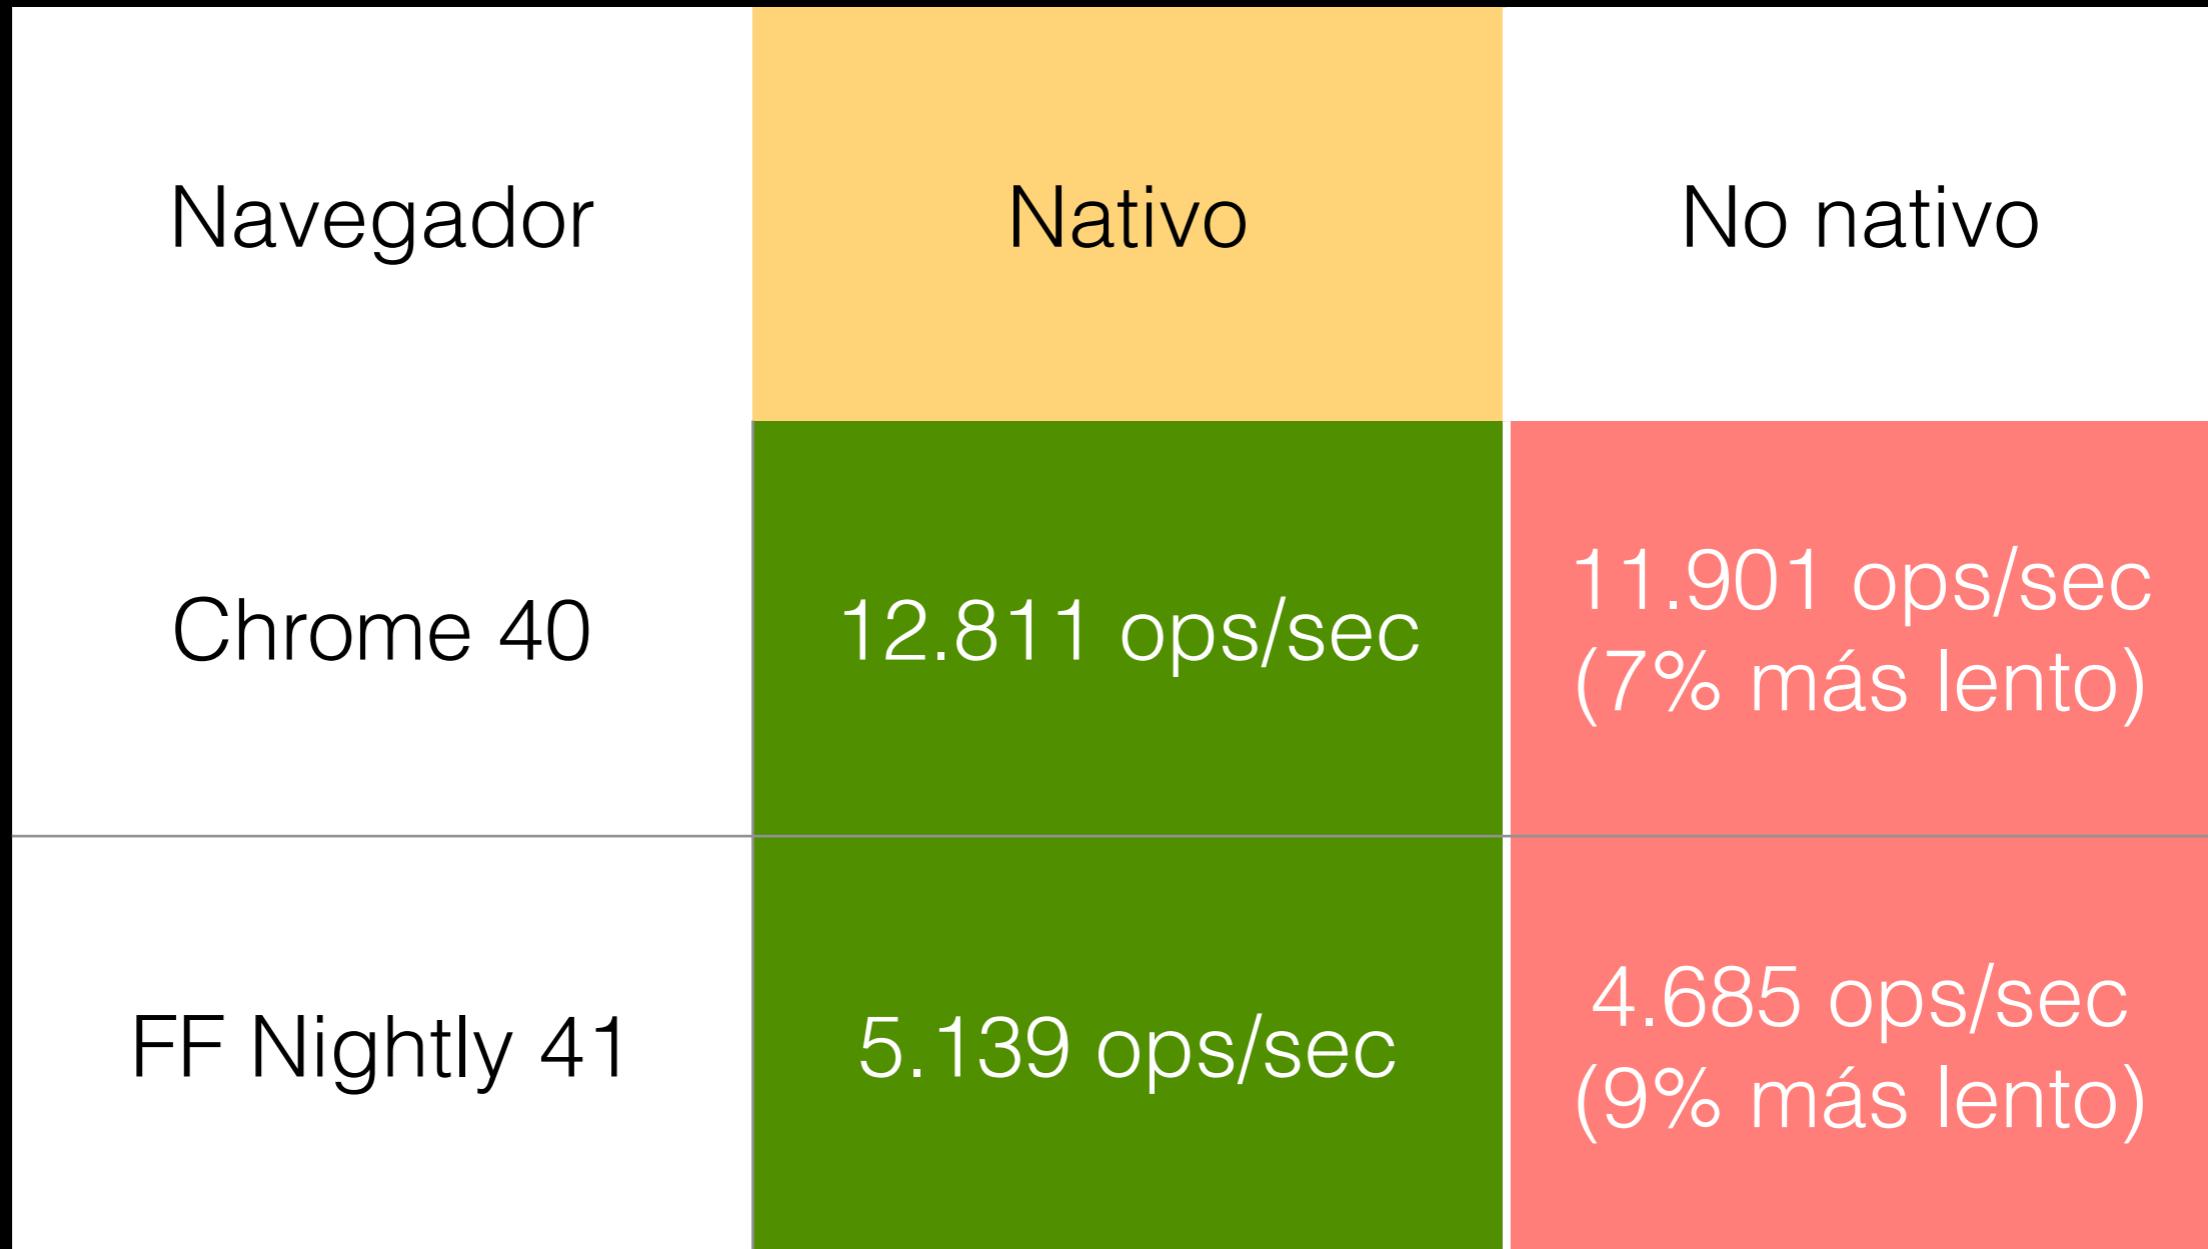
p-j-S

```
> var f = function(a) { return a + 1; };
> worker.postMessage(f);
Uncaught DOMException: Failed to execute 'postMessage'
on 'Worker': An object could not be cloned.
```

Enviar funciones

```
> var f = new Function([ 'a' ], 'return a + 1');  
> f(2)  
3
```

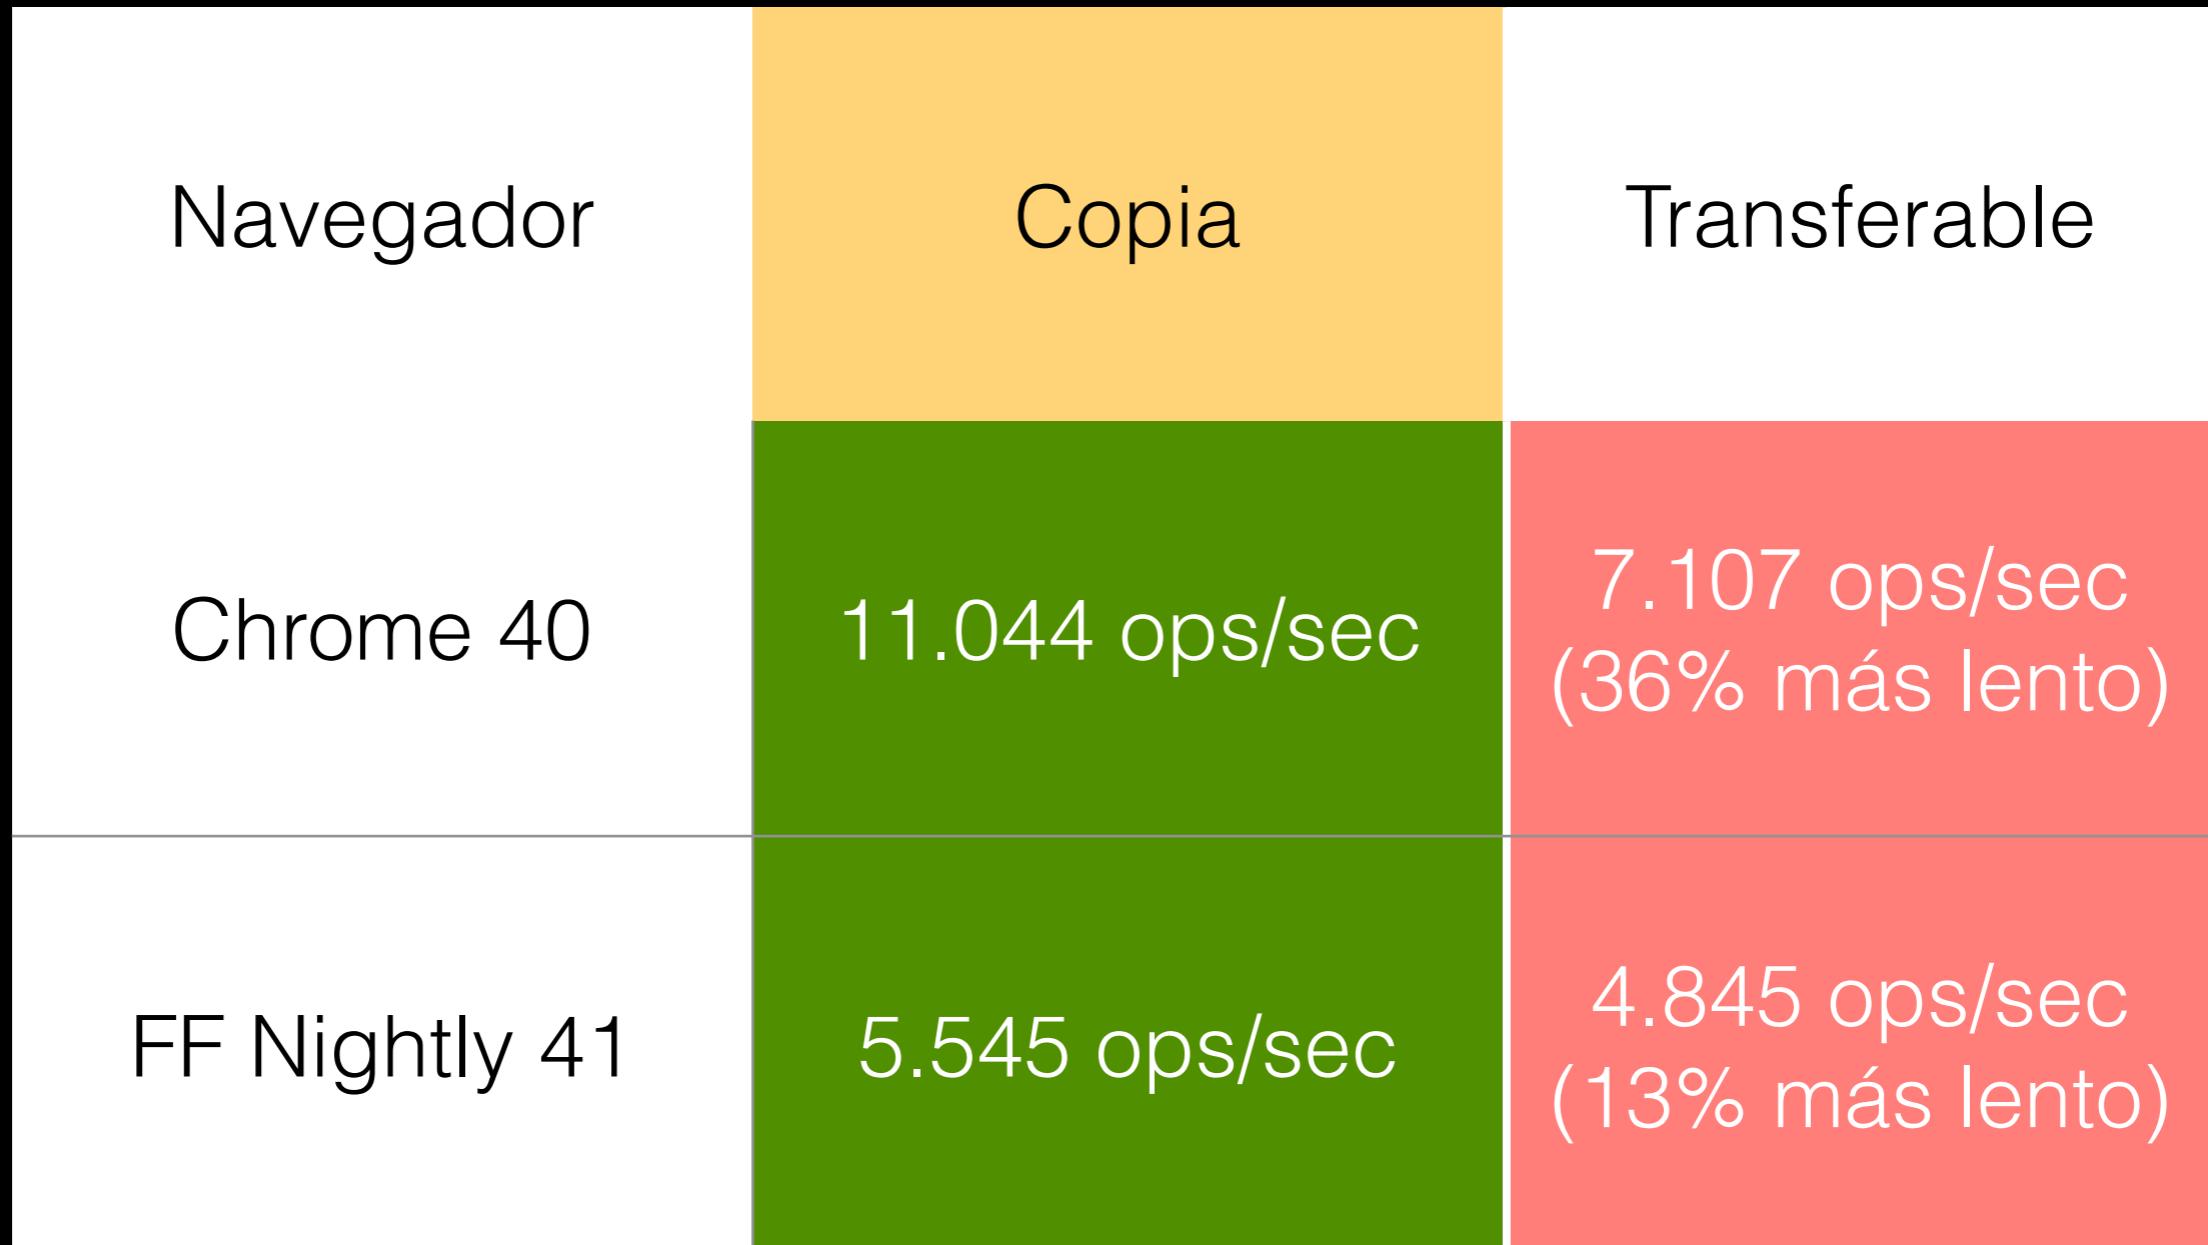
Alternativa



Codificación

Para utilizar con ArrayBuffer

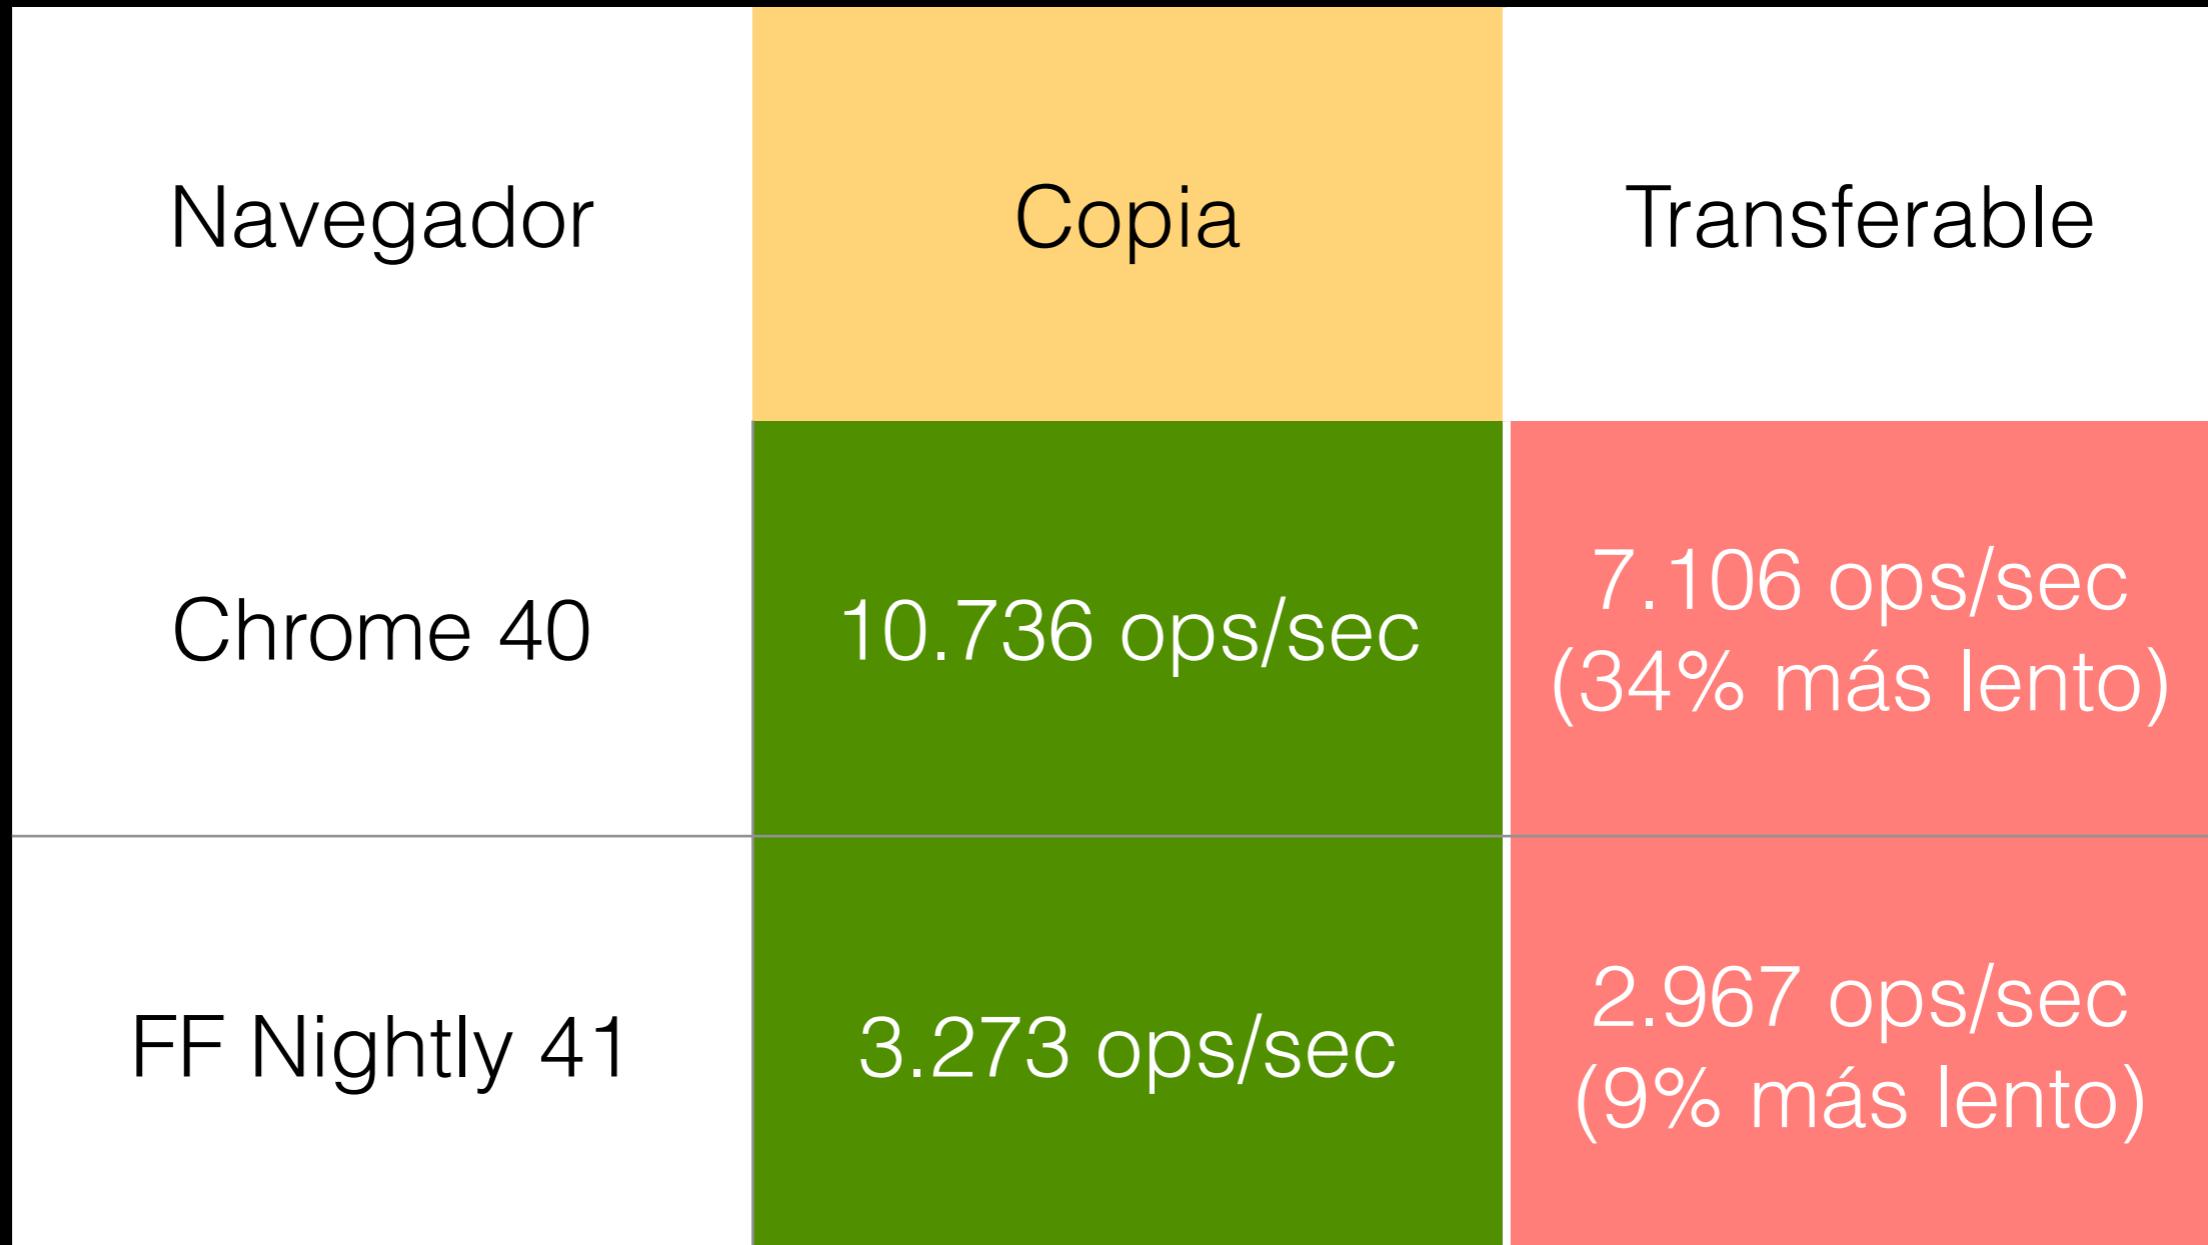
<http://jsperf.com/pjs-serialization-comparison/4>



Pasaje mensaje

Corto, sin memoria compartida

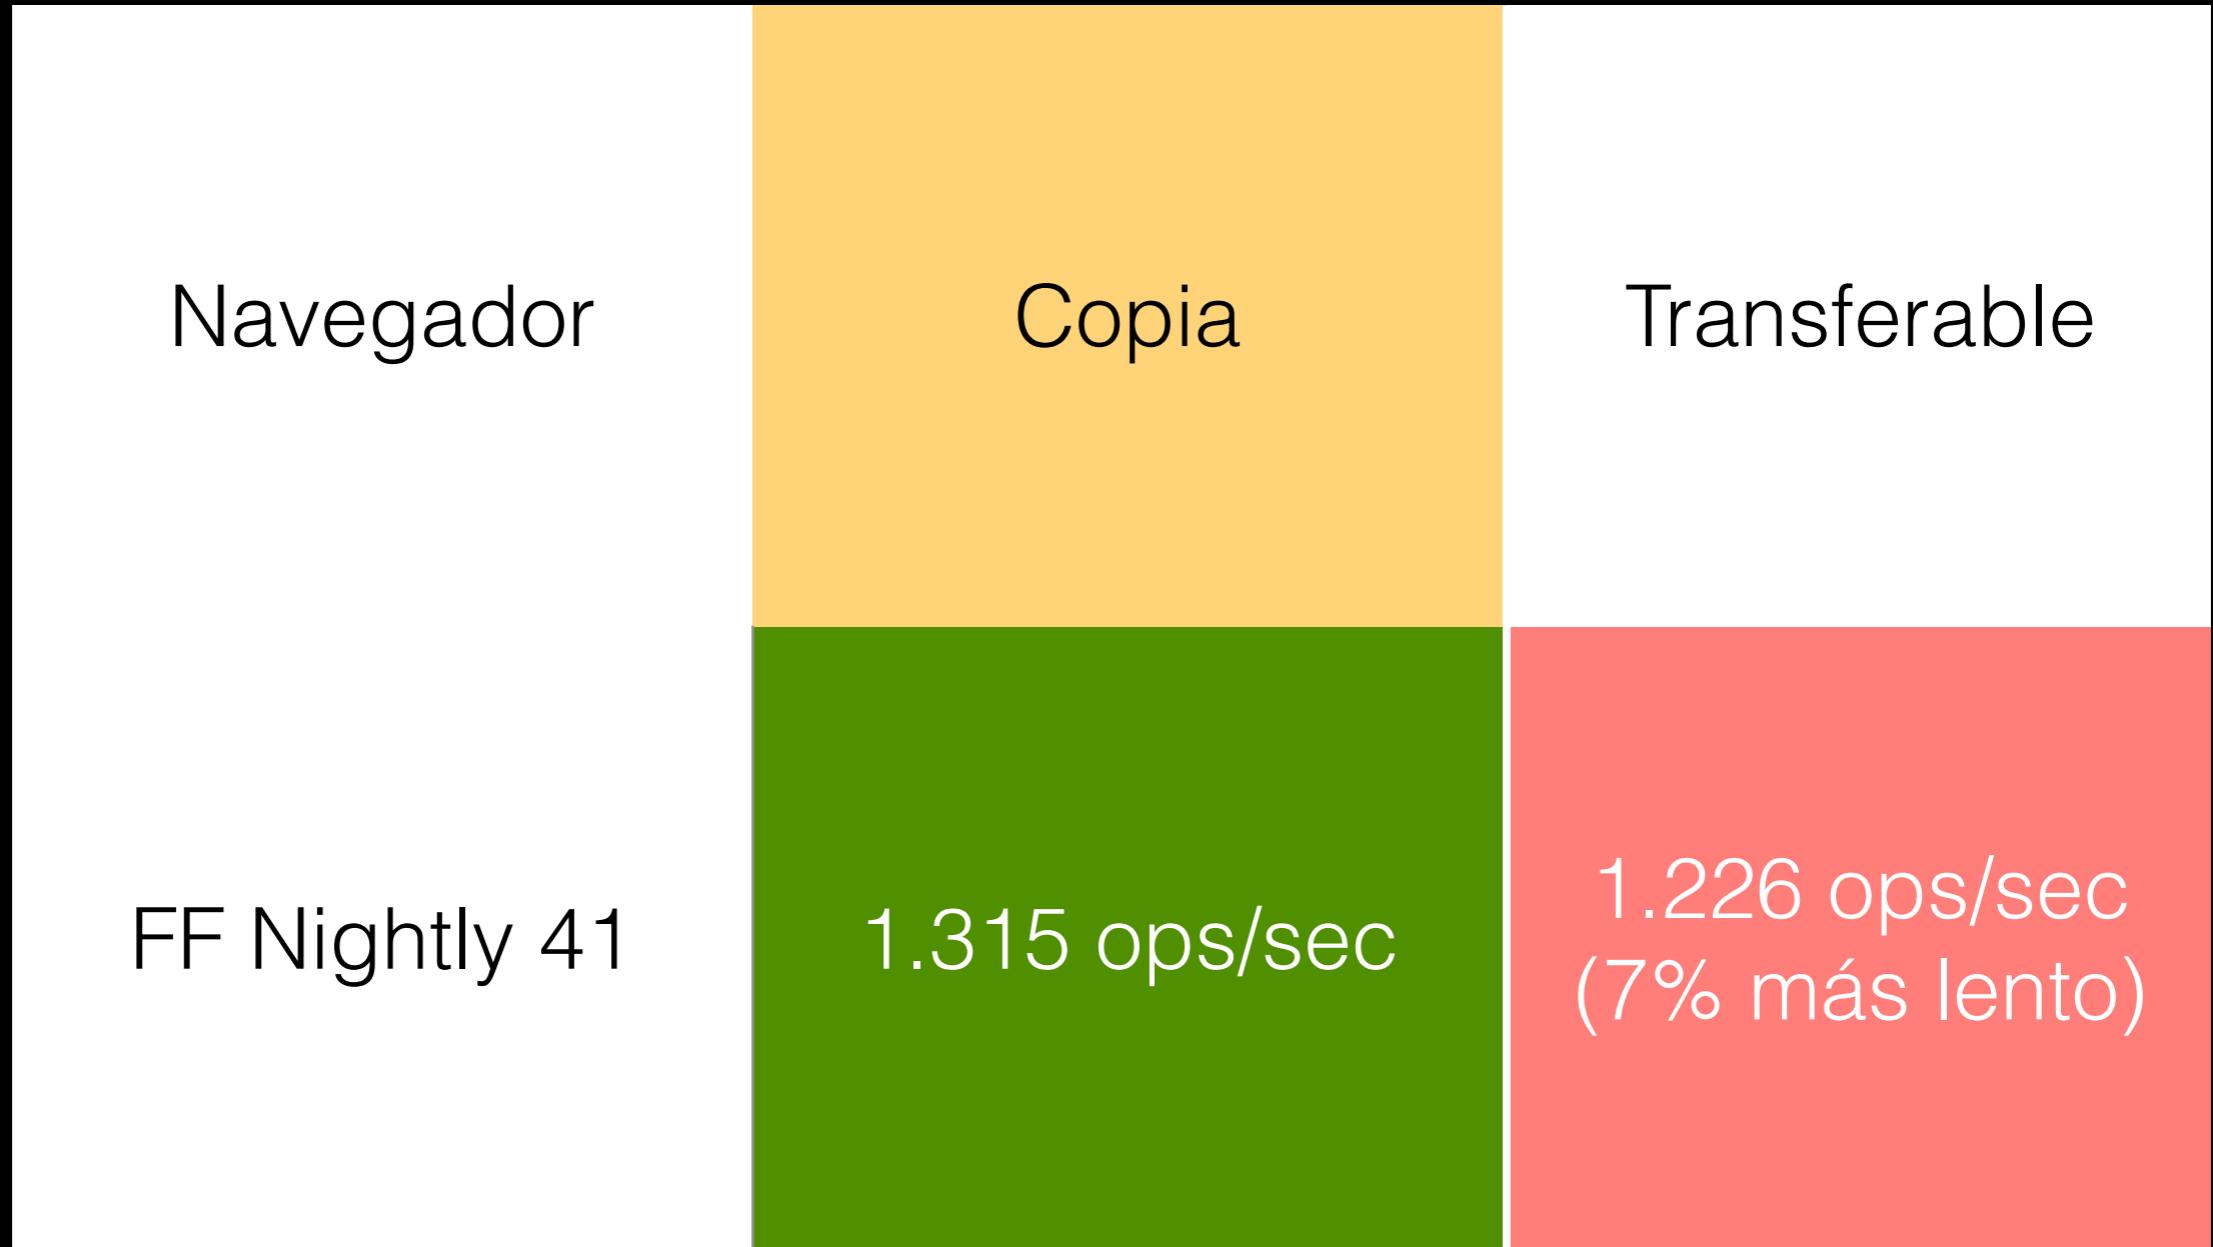
<http://jsperf.com/pjs-serialization/5>



Pasaje mensaje

Largo, sin memoria compartida

<http://jsperf.com/pjs-serialization-long/3>



Pasaje mensaje

Con memoria compartida, a cuatro Web Workers
<http://jsperf.com/shared-function-transfer>

Contexto

Serialización de clausuras

```
var xs = [1,2,3,4];
var fact = 4;

var ys = xs.map(function(x){
    return x * fact;
});

// ys [ 4, 8, 12, 16 ]
```

Clausuras

Closures

Contexto Global

El mismo para todas las cadenas

Contexto local

Específico para un paso de una cadena

```
var xs = new Uint8Array([1,2,3,4]);  
  
// contexto global  
pjs.updateContext({  
  max: 3  
}).then(() => {  
  pjs(xs).filter(  
    (e, ctx) => e <= ctx.max && e >= ctx.min,  
    /* contexto local */ { min: 2 }  
  )  
  .seq().then(ys => {  
    // ys [2 ,3]  
  });  
});
```

Contexto

Ejemplo (1)

```
pjs.updateContext({ A: 0, B: 25, C: 44 })
.then(() => {
  var chain = pjs(xs)
    .map(
      (e, ctx) => { /* code here */ },
      { A: 5, B: 13 }
    ).filter(
      (e, ctx) => { /* code here */ },
      { A: 2, E: 9 }
    );
  ...
}
});
```

Contexto

Ejemplo (2)

Local Context map	
A	5
B	13



Global Context	
A	0
B	25
C	44



Available ctx map	
A	5
B	13
C	44

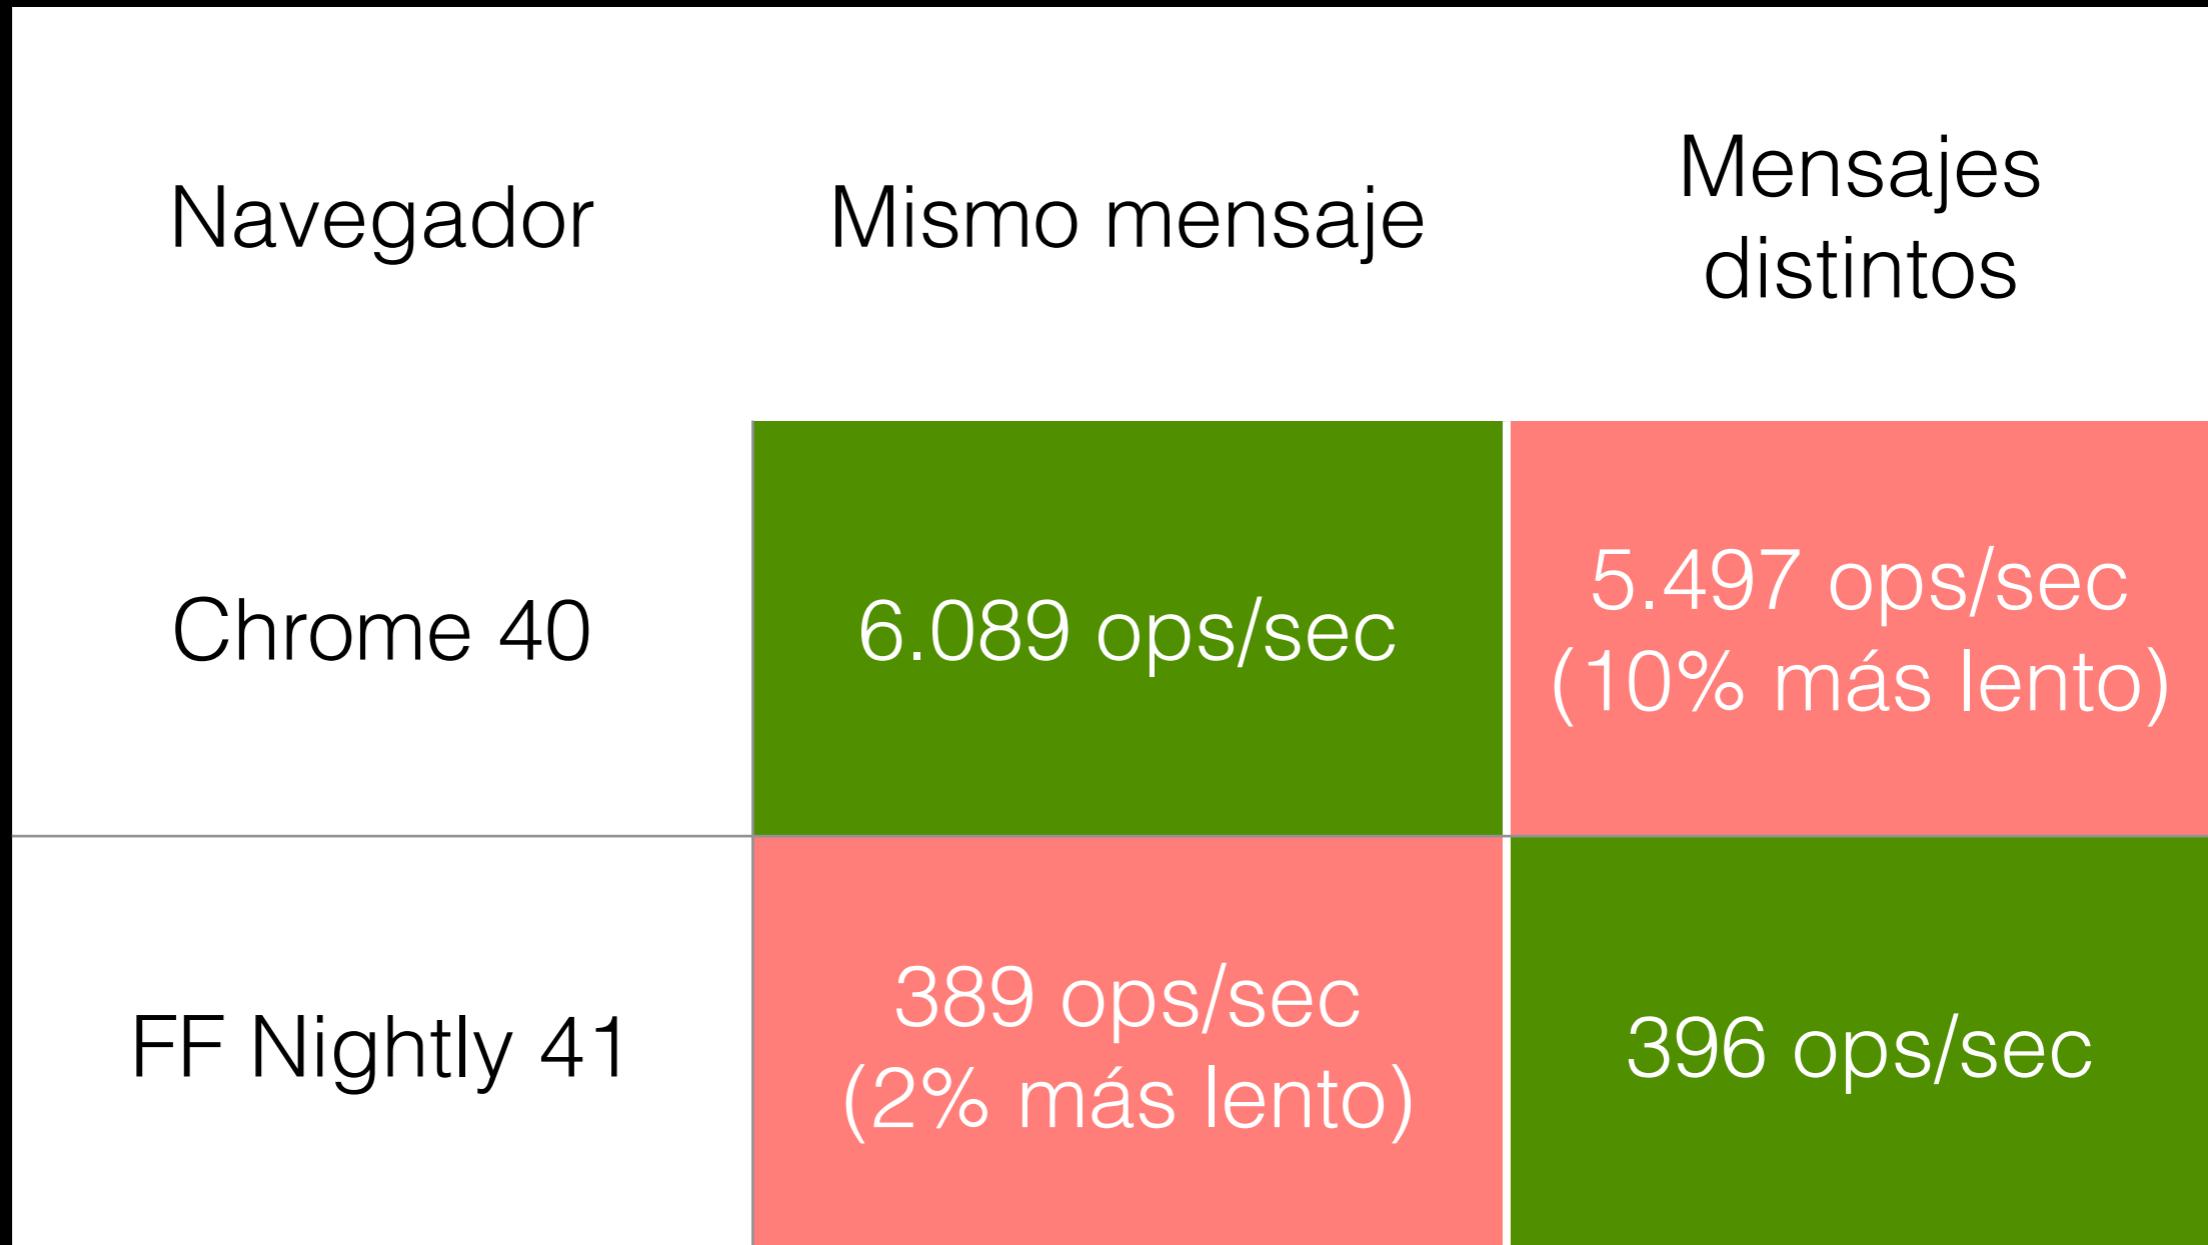
Local Context filter	
A	2
E	9



Global Context	
A	0
B	25
C	44

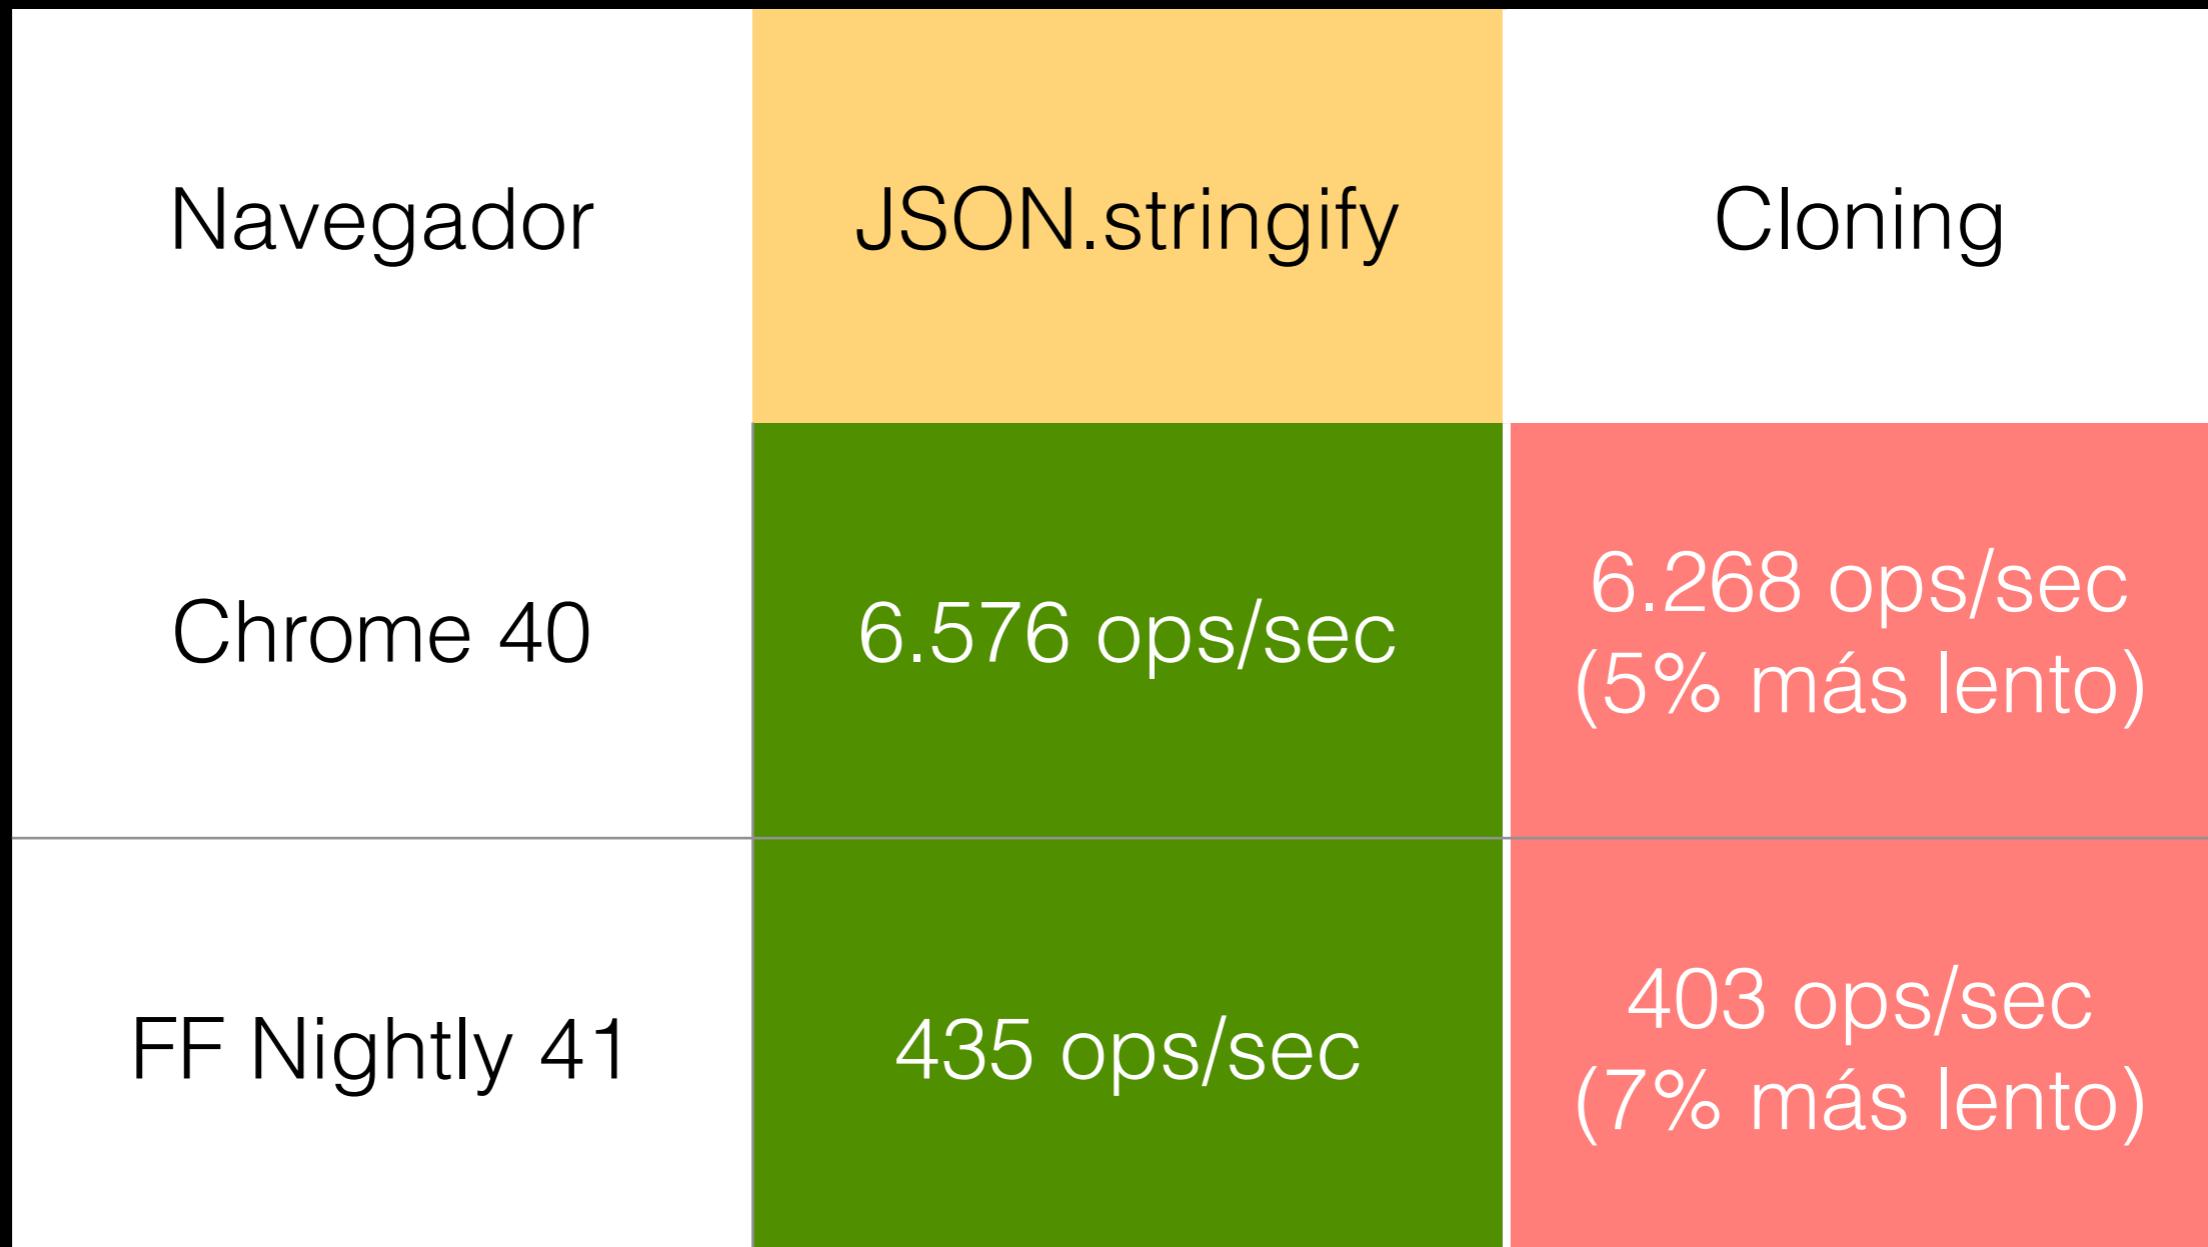
Available ctx filter	
A	2
B	25
C	44
E	9

Union de contextos



Pasaje contexto

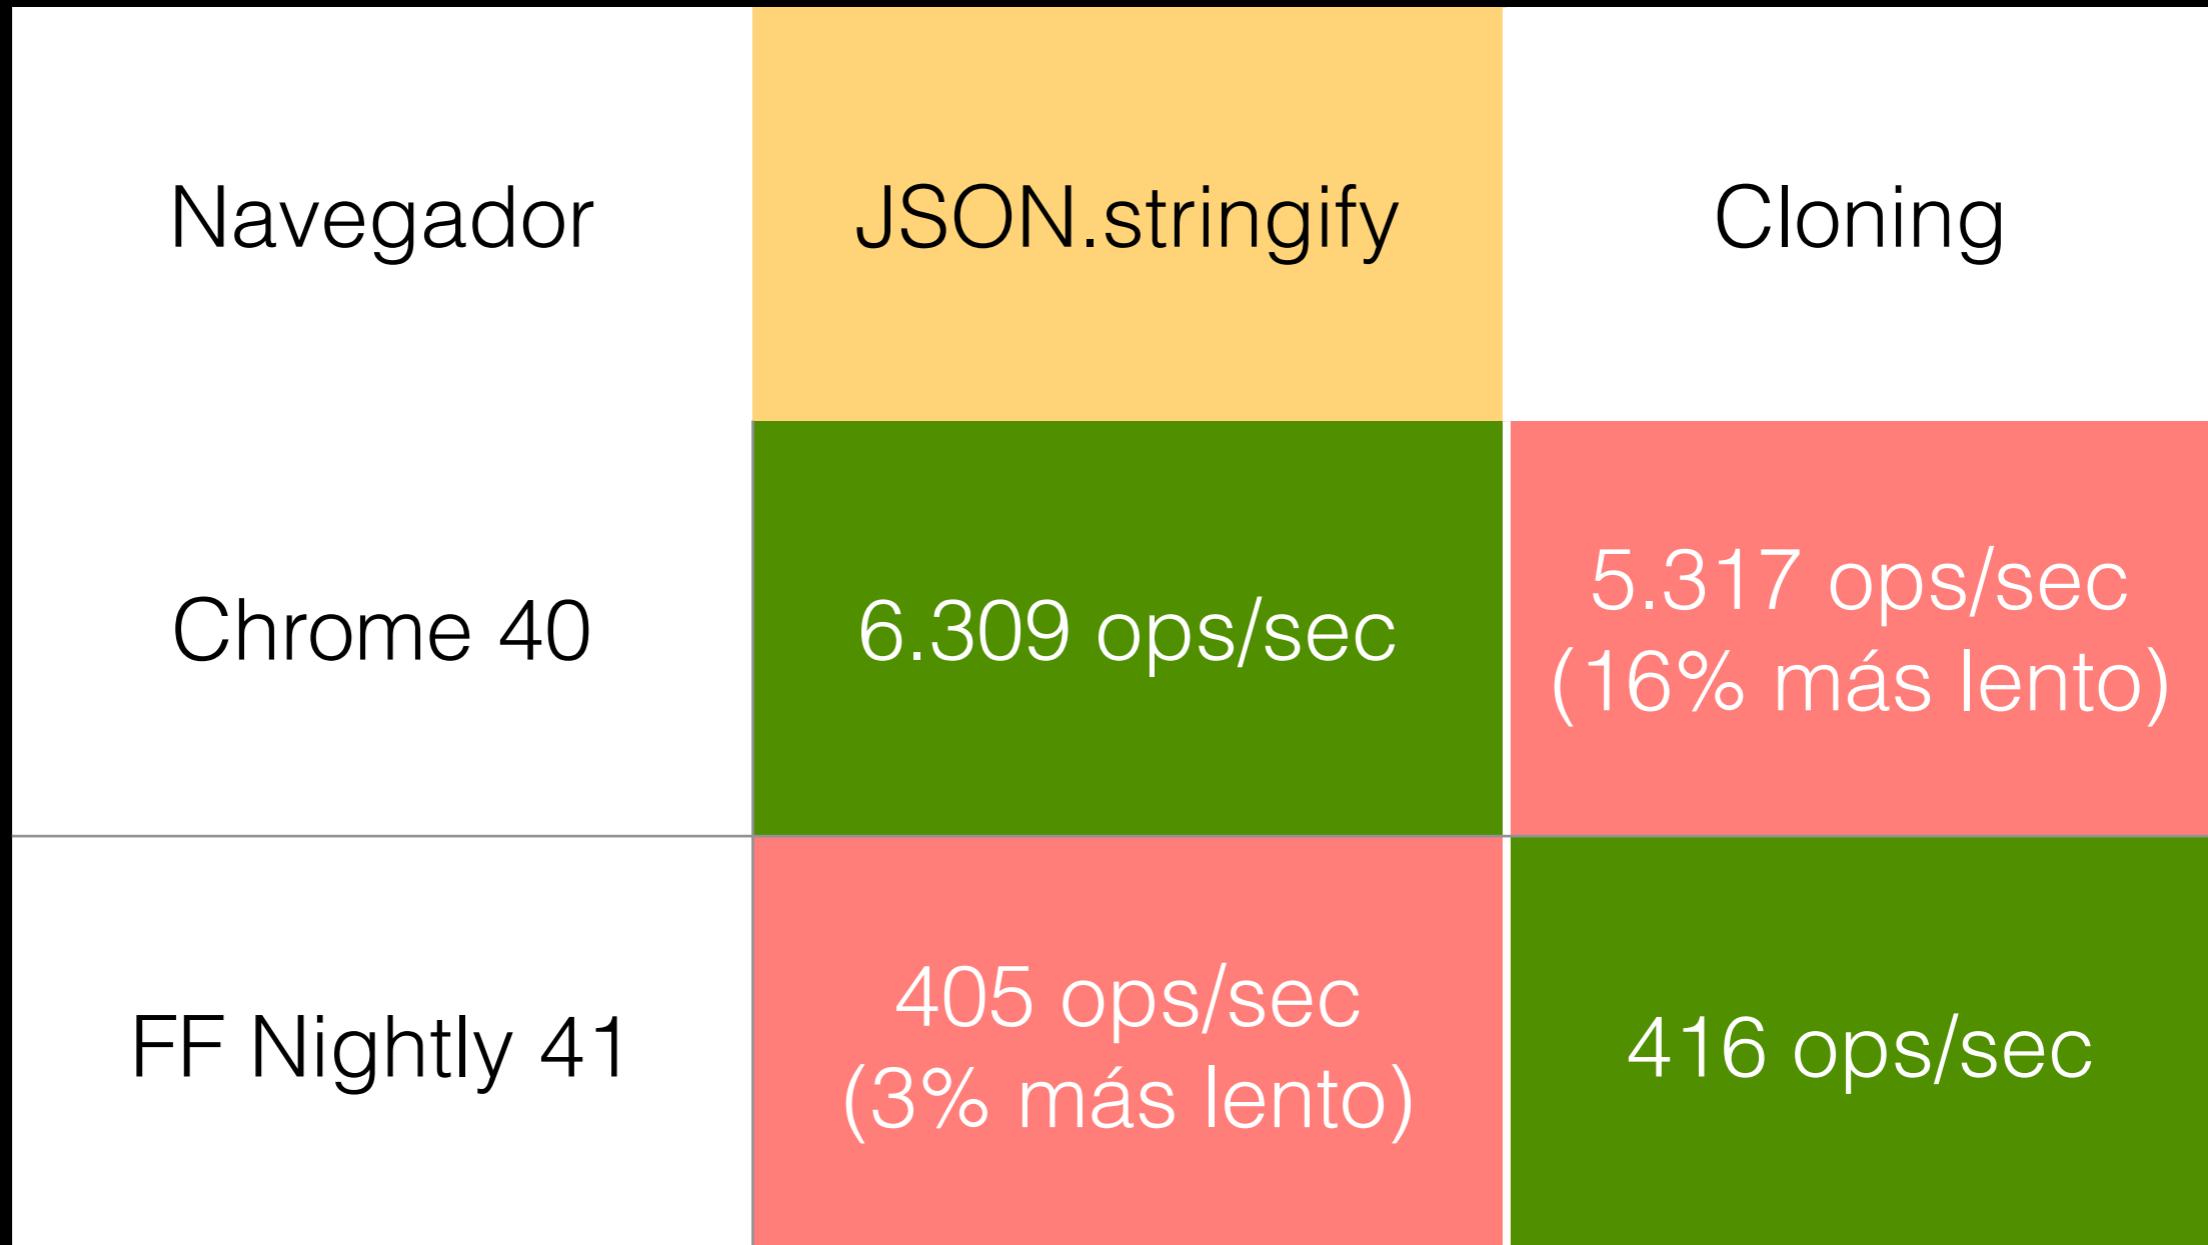
Enviar en mismo mensaje que código o separado
<http://jsperf.com/additional-sendingdata-comparison/2>



Serialización

Contexto pequeño

<http://jsperf.com/additional-clonningdata-comparison/3>

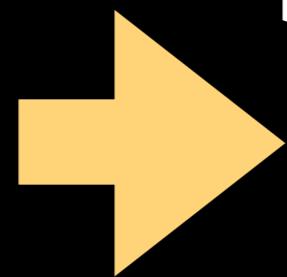


Serialización

Contexto grande

<http://jsperf.com/additional-clonningdata-comparison-long/3>

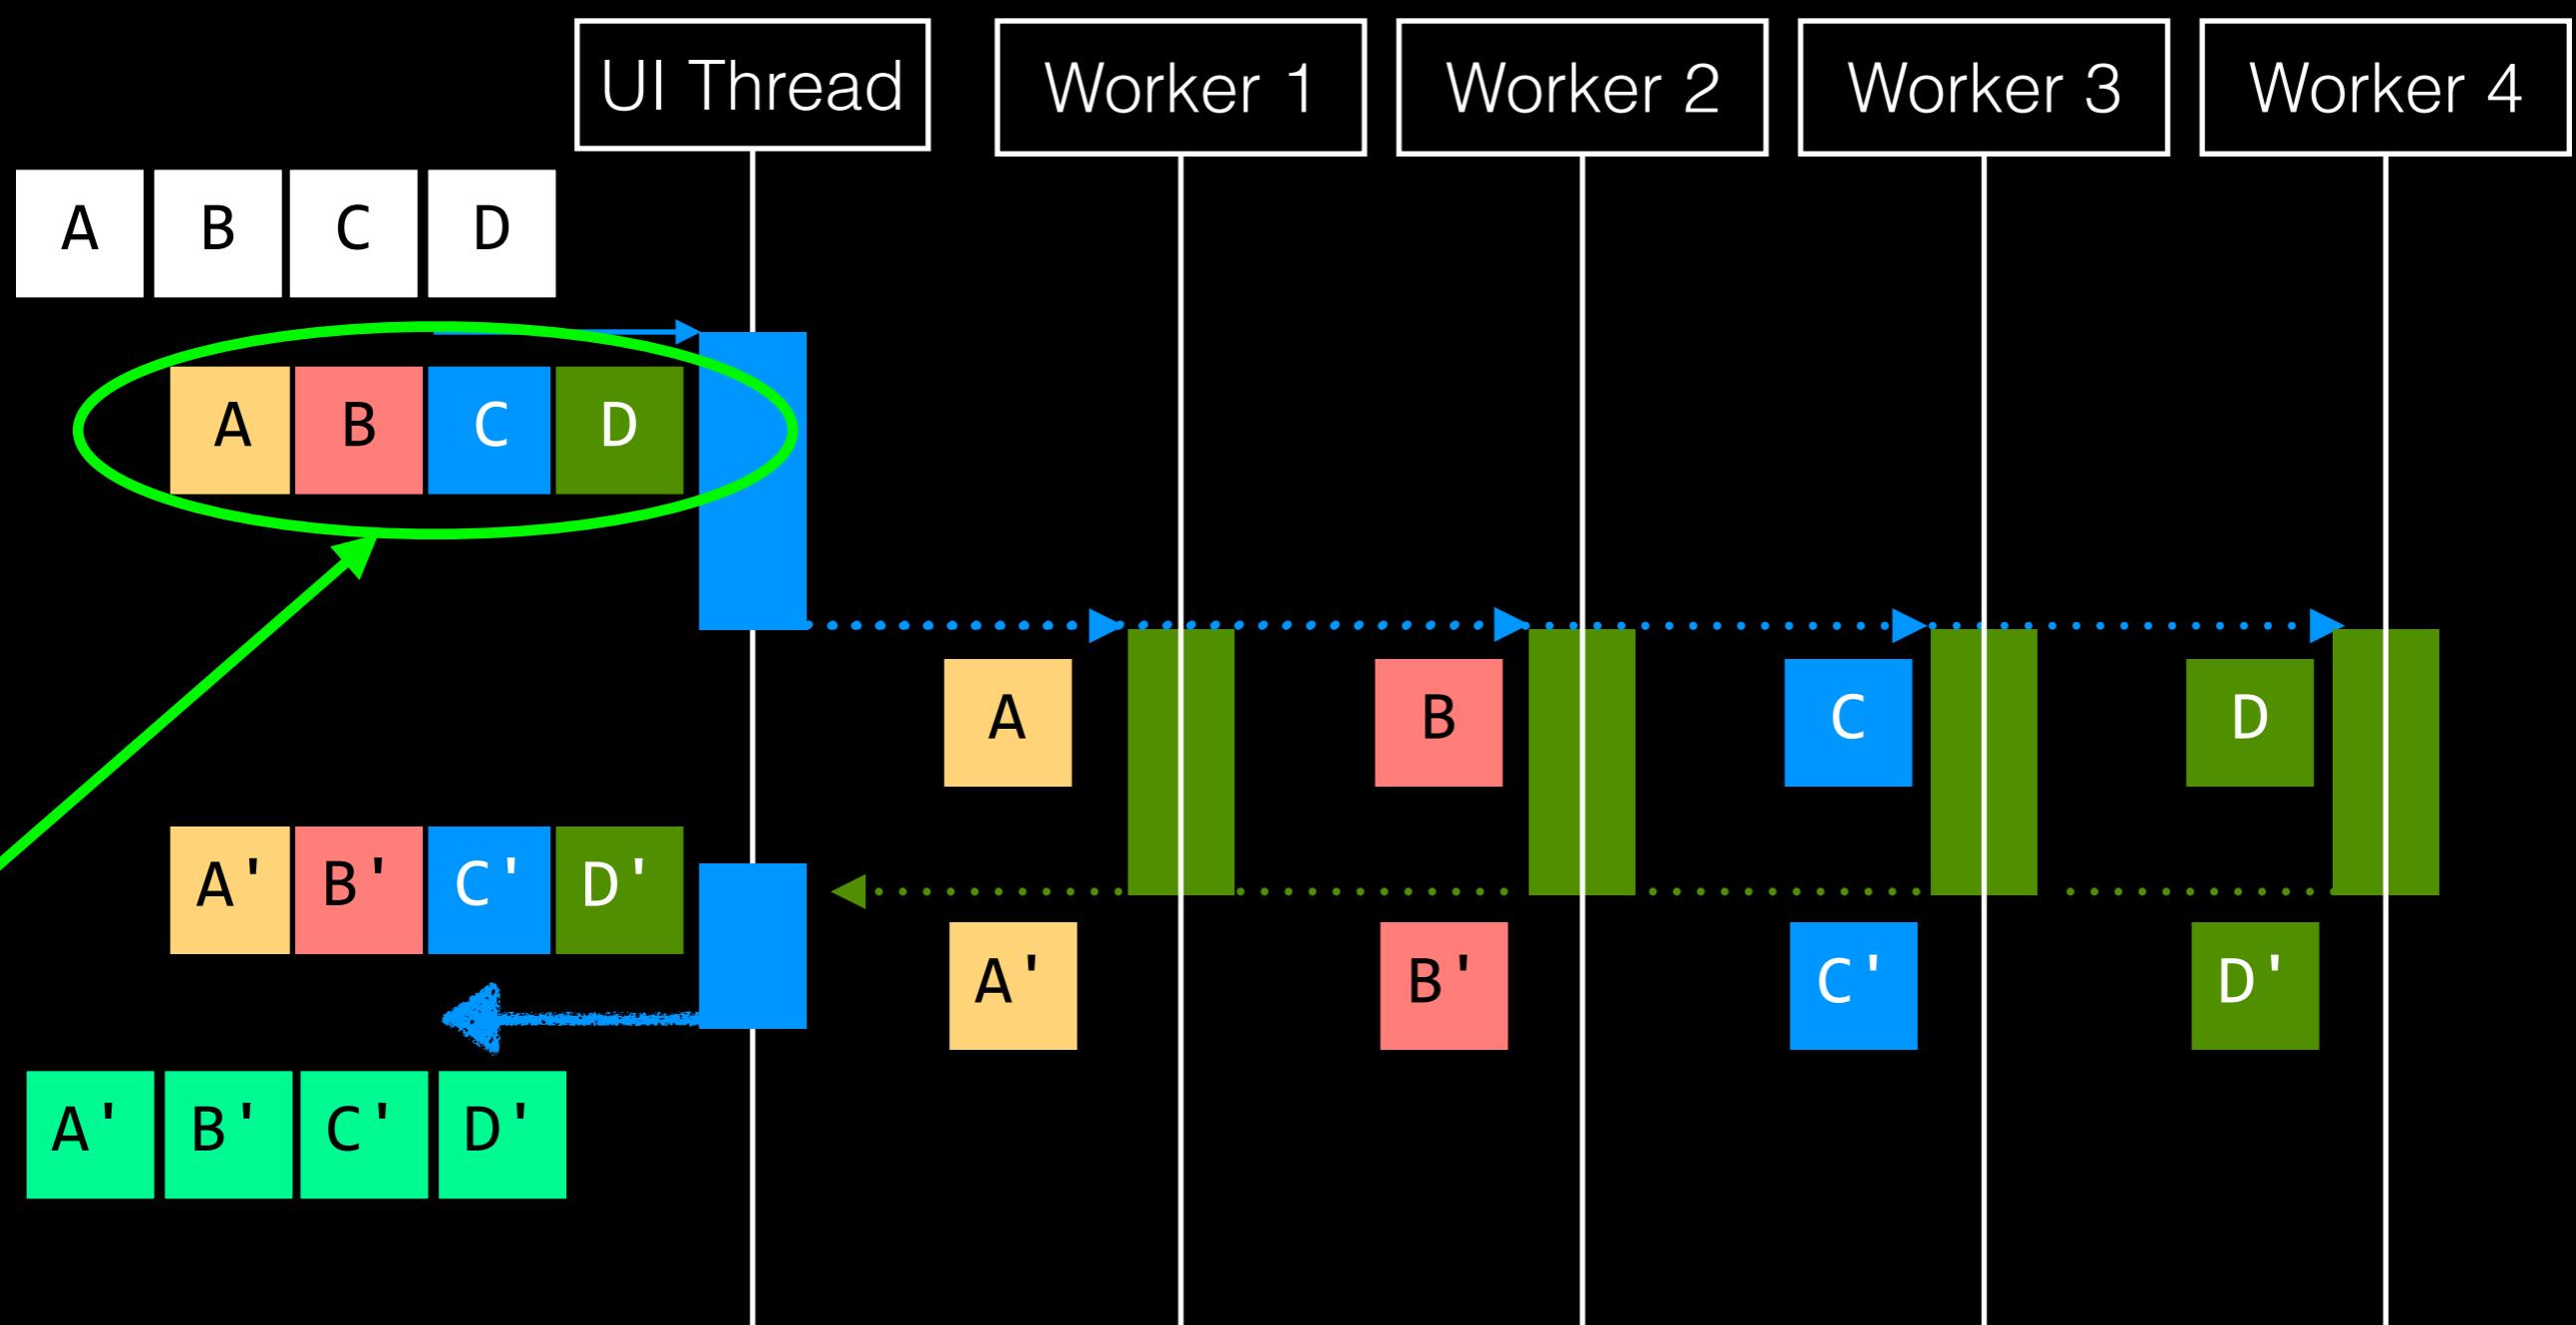
```
{  
  a: 2,  
  b: function(c) {  
    return c + 1;  
  }  
}
```



```
{  
  a: 2,  
  b: {  
    args: [ 'c' ],  
    code: '\n\\treturn c + 1;\n',  
    __isFunction: true  
  }  
}
```

Contexto con funciones

Partición de TypedArray(s)



p-j-S

Navegador	Subarray + Constructor	Subarray + Set	Buffer slice + Constructor
Chrome 40	2868 ops/sec	2.234 ops/sec (22% más lento)	2.247 ops/sec (22% más lento)
FF Nightly 41	1.613 ops/sec (5% más lento)	1.638 ops/sec (3% más lento)	1.692 ops/sec

Partición

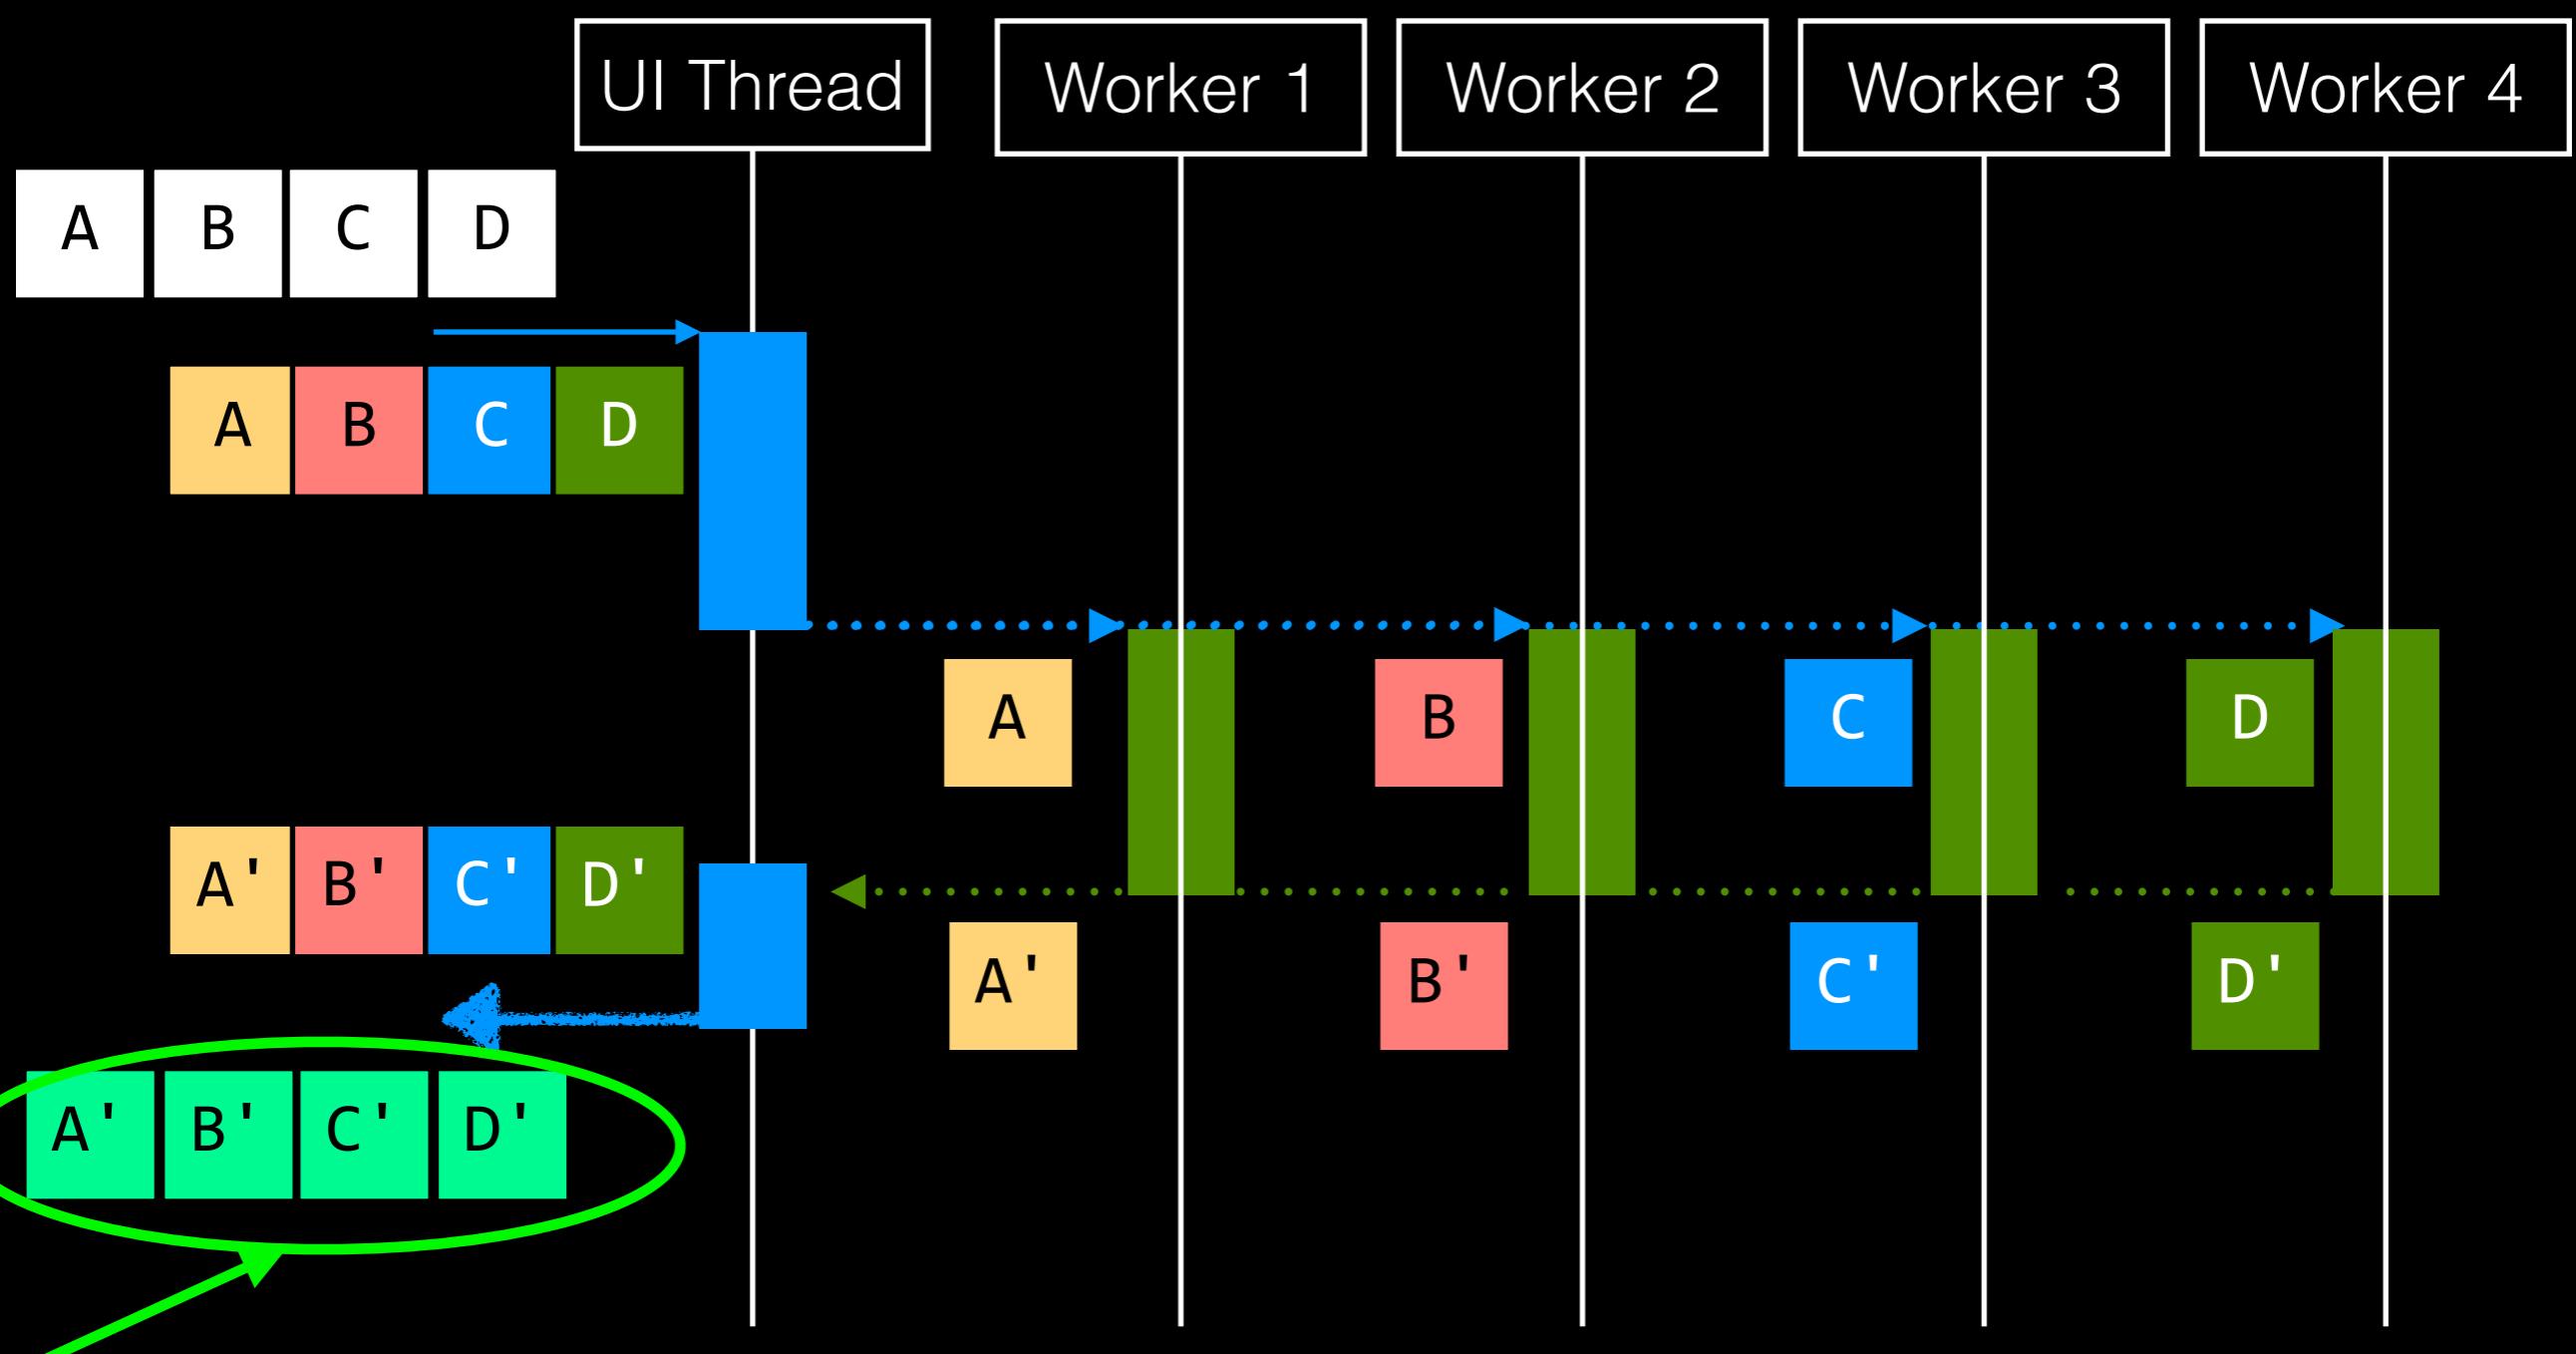
1.000.000 elementos en 4 partes
<http://jsperf.com/arraybuffer-split/3>

```
function typedArraySlice(array, from, to) {  
    return new array.constructor(array.subarray(from, to));  
}
```

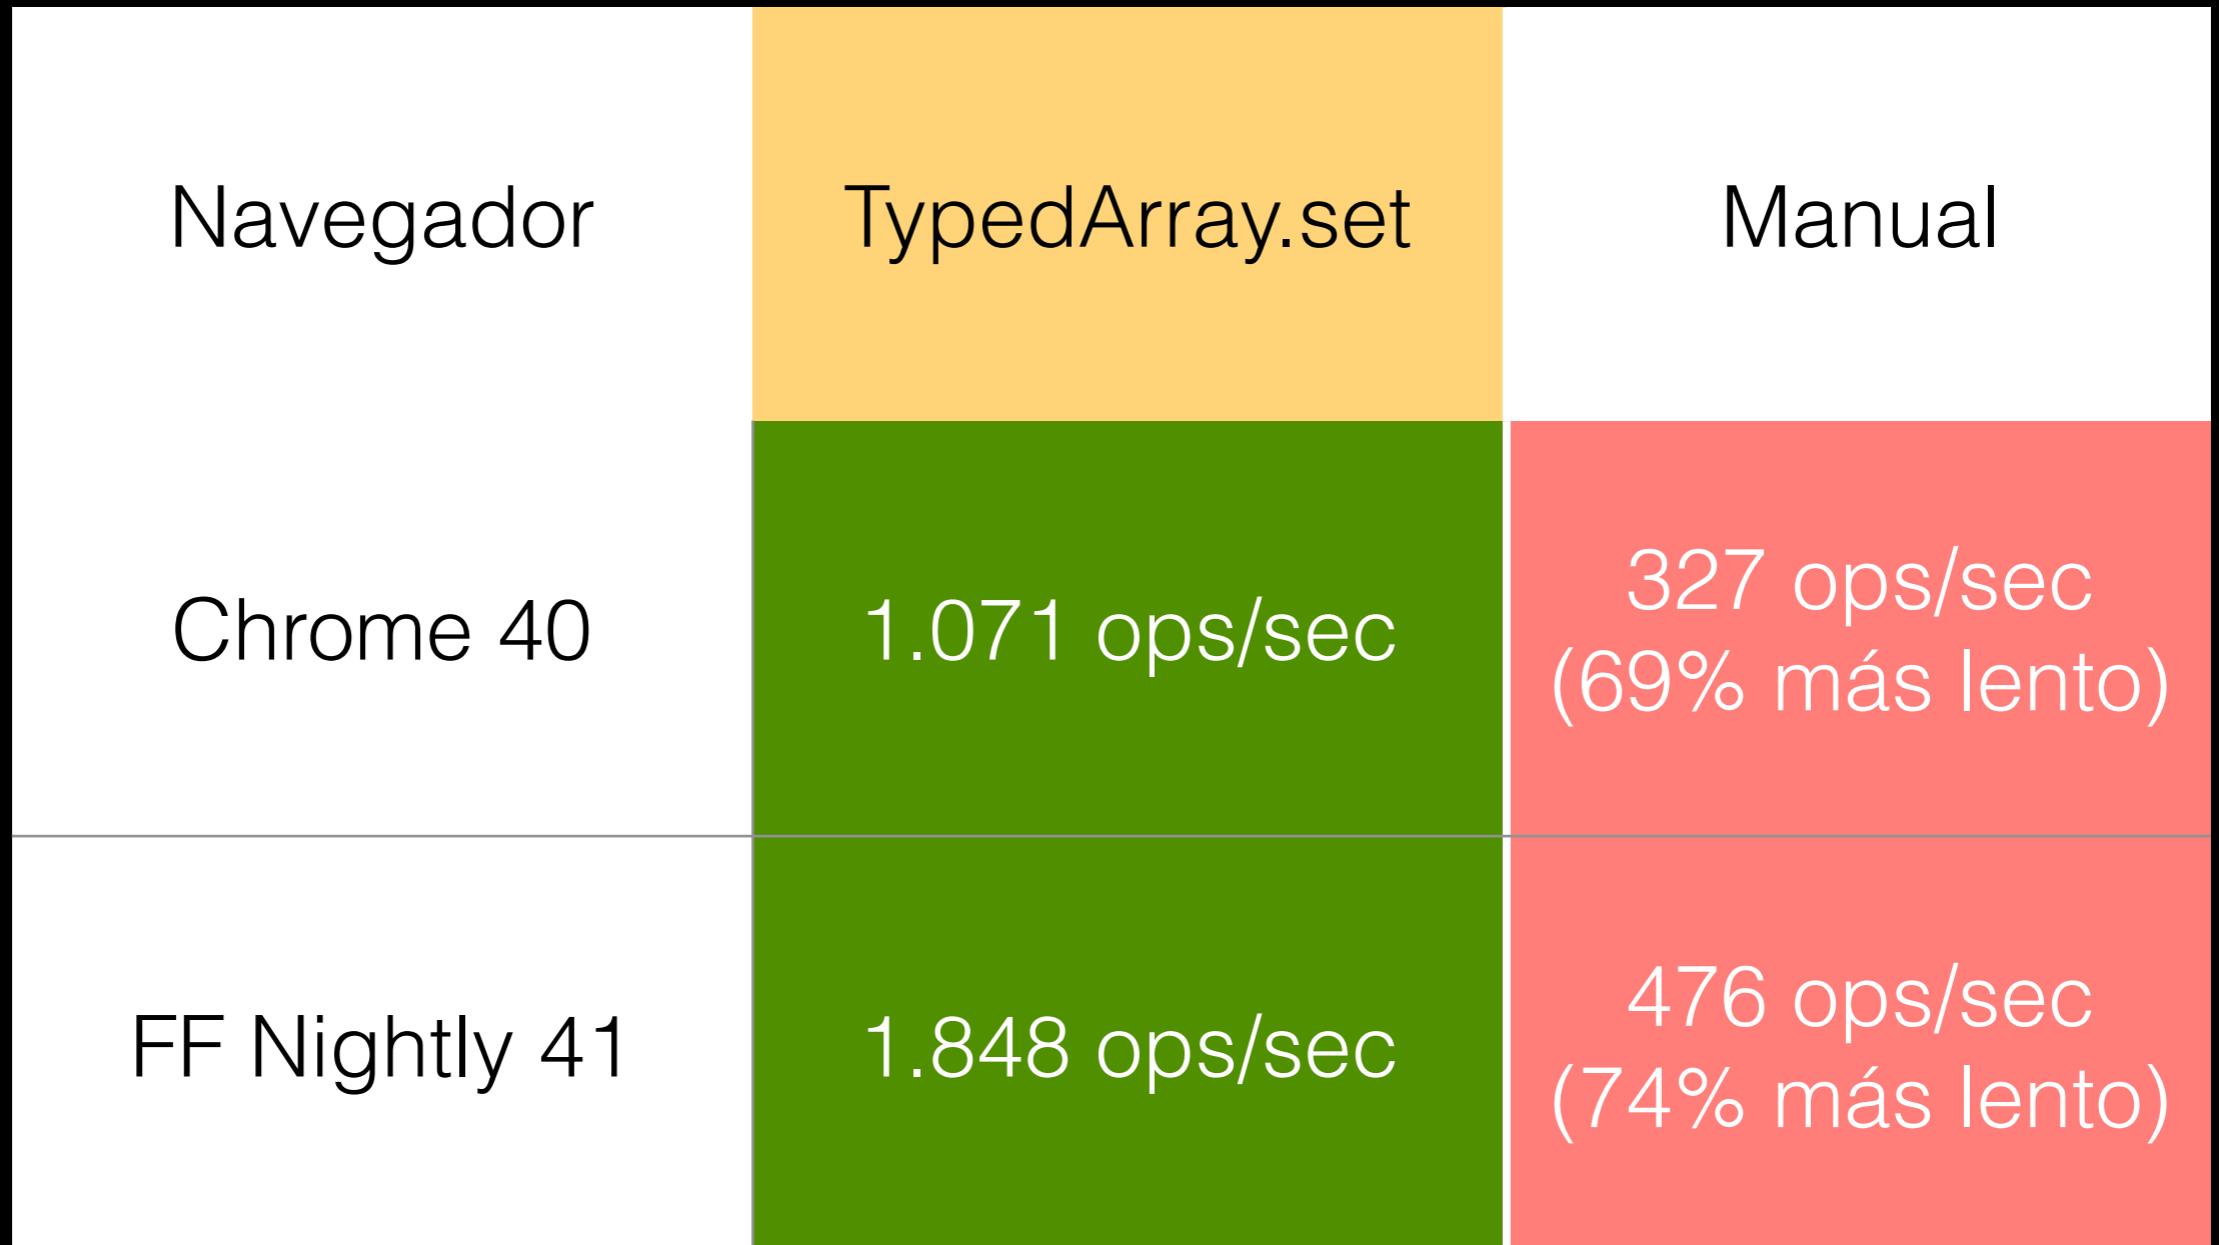
Partición

Ejemplo

Consolidación de resultados

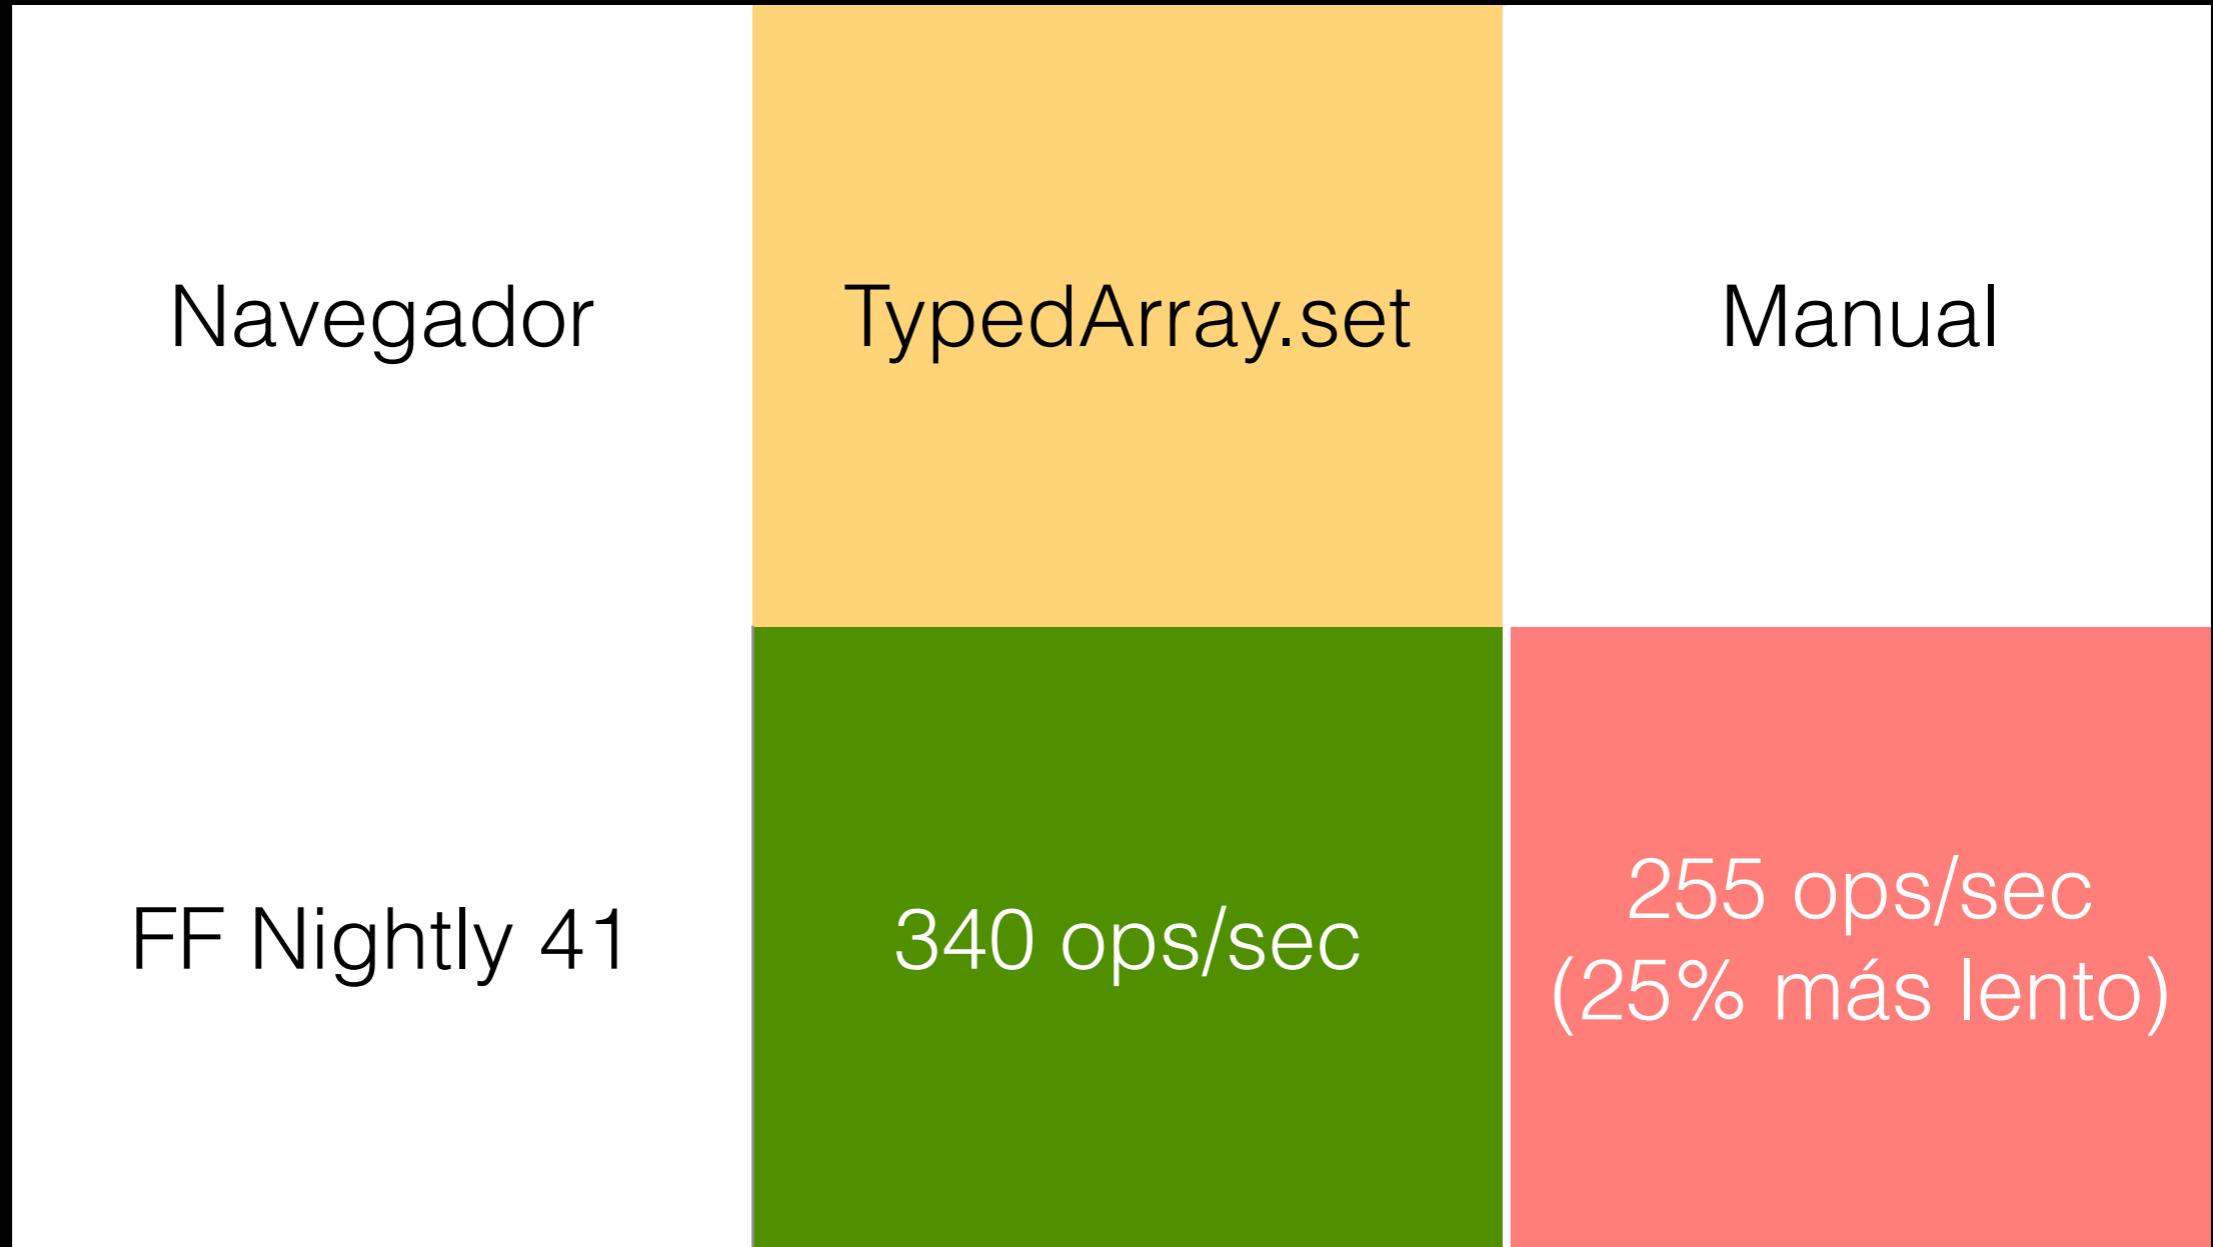


p-j-S



TypedArray

<http://jsperf.com/typedarray-merge/2>

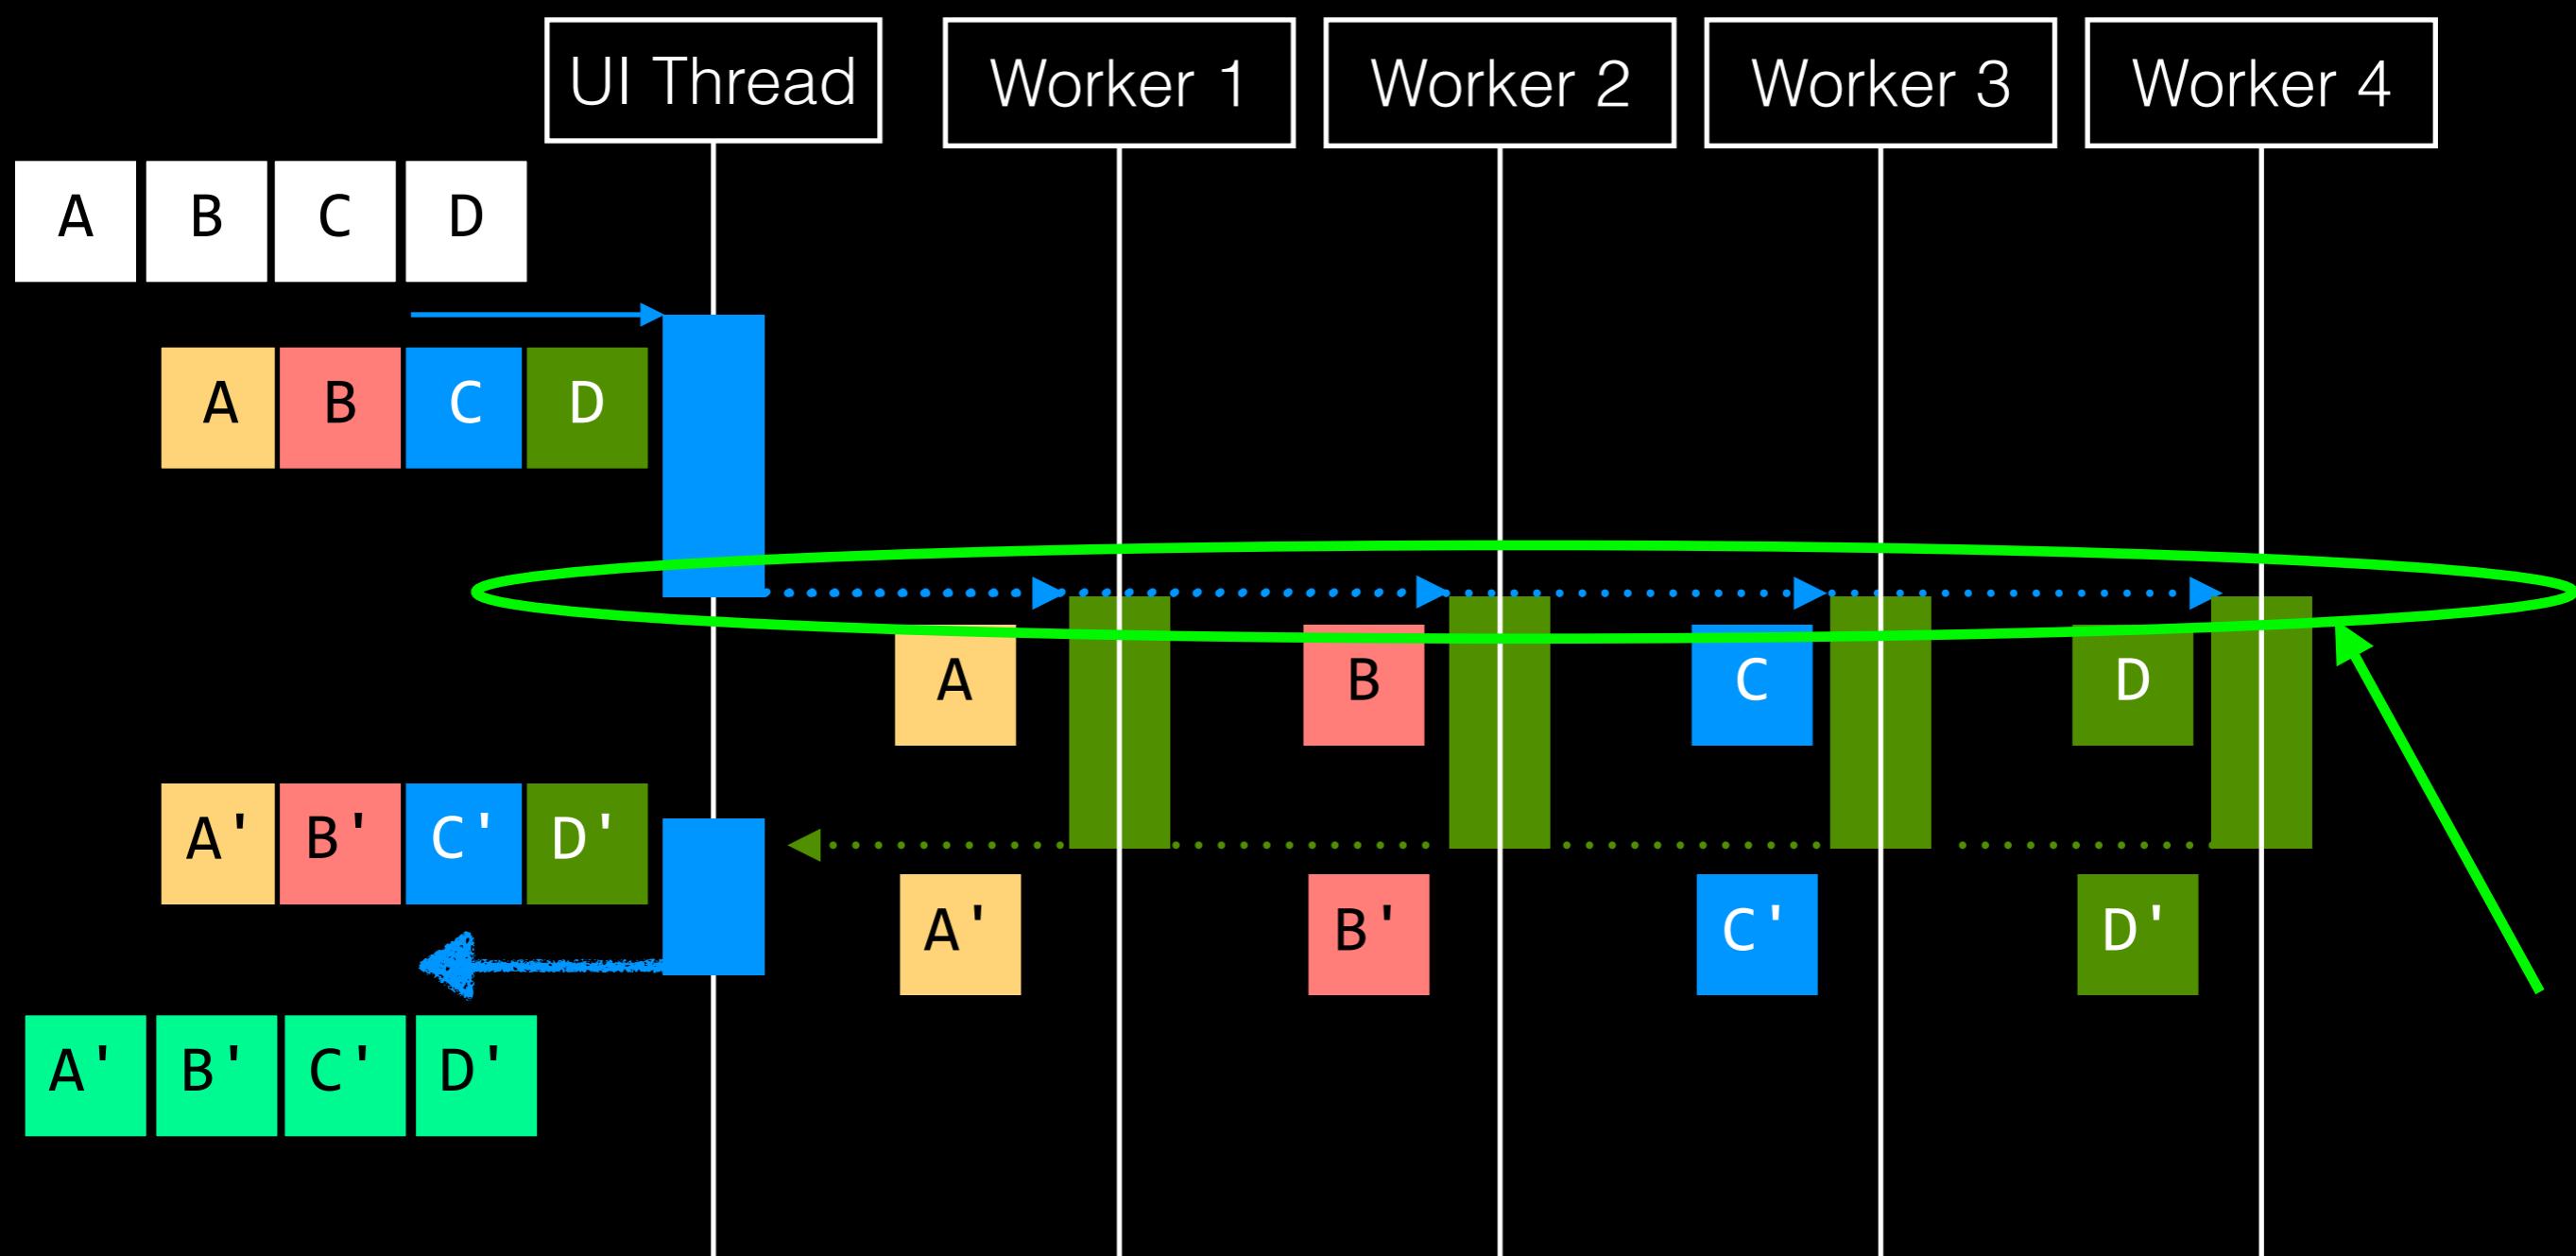


SharedTypedArray

<http://jsperf.com/sharedtypedarray-merge>

Optimizaciones

Tiempo transferencia
funciones

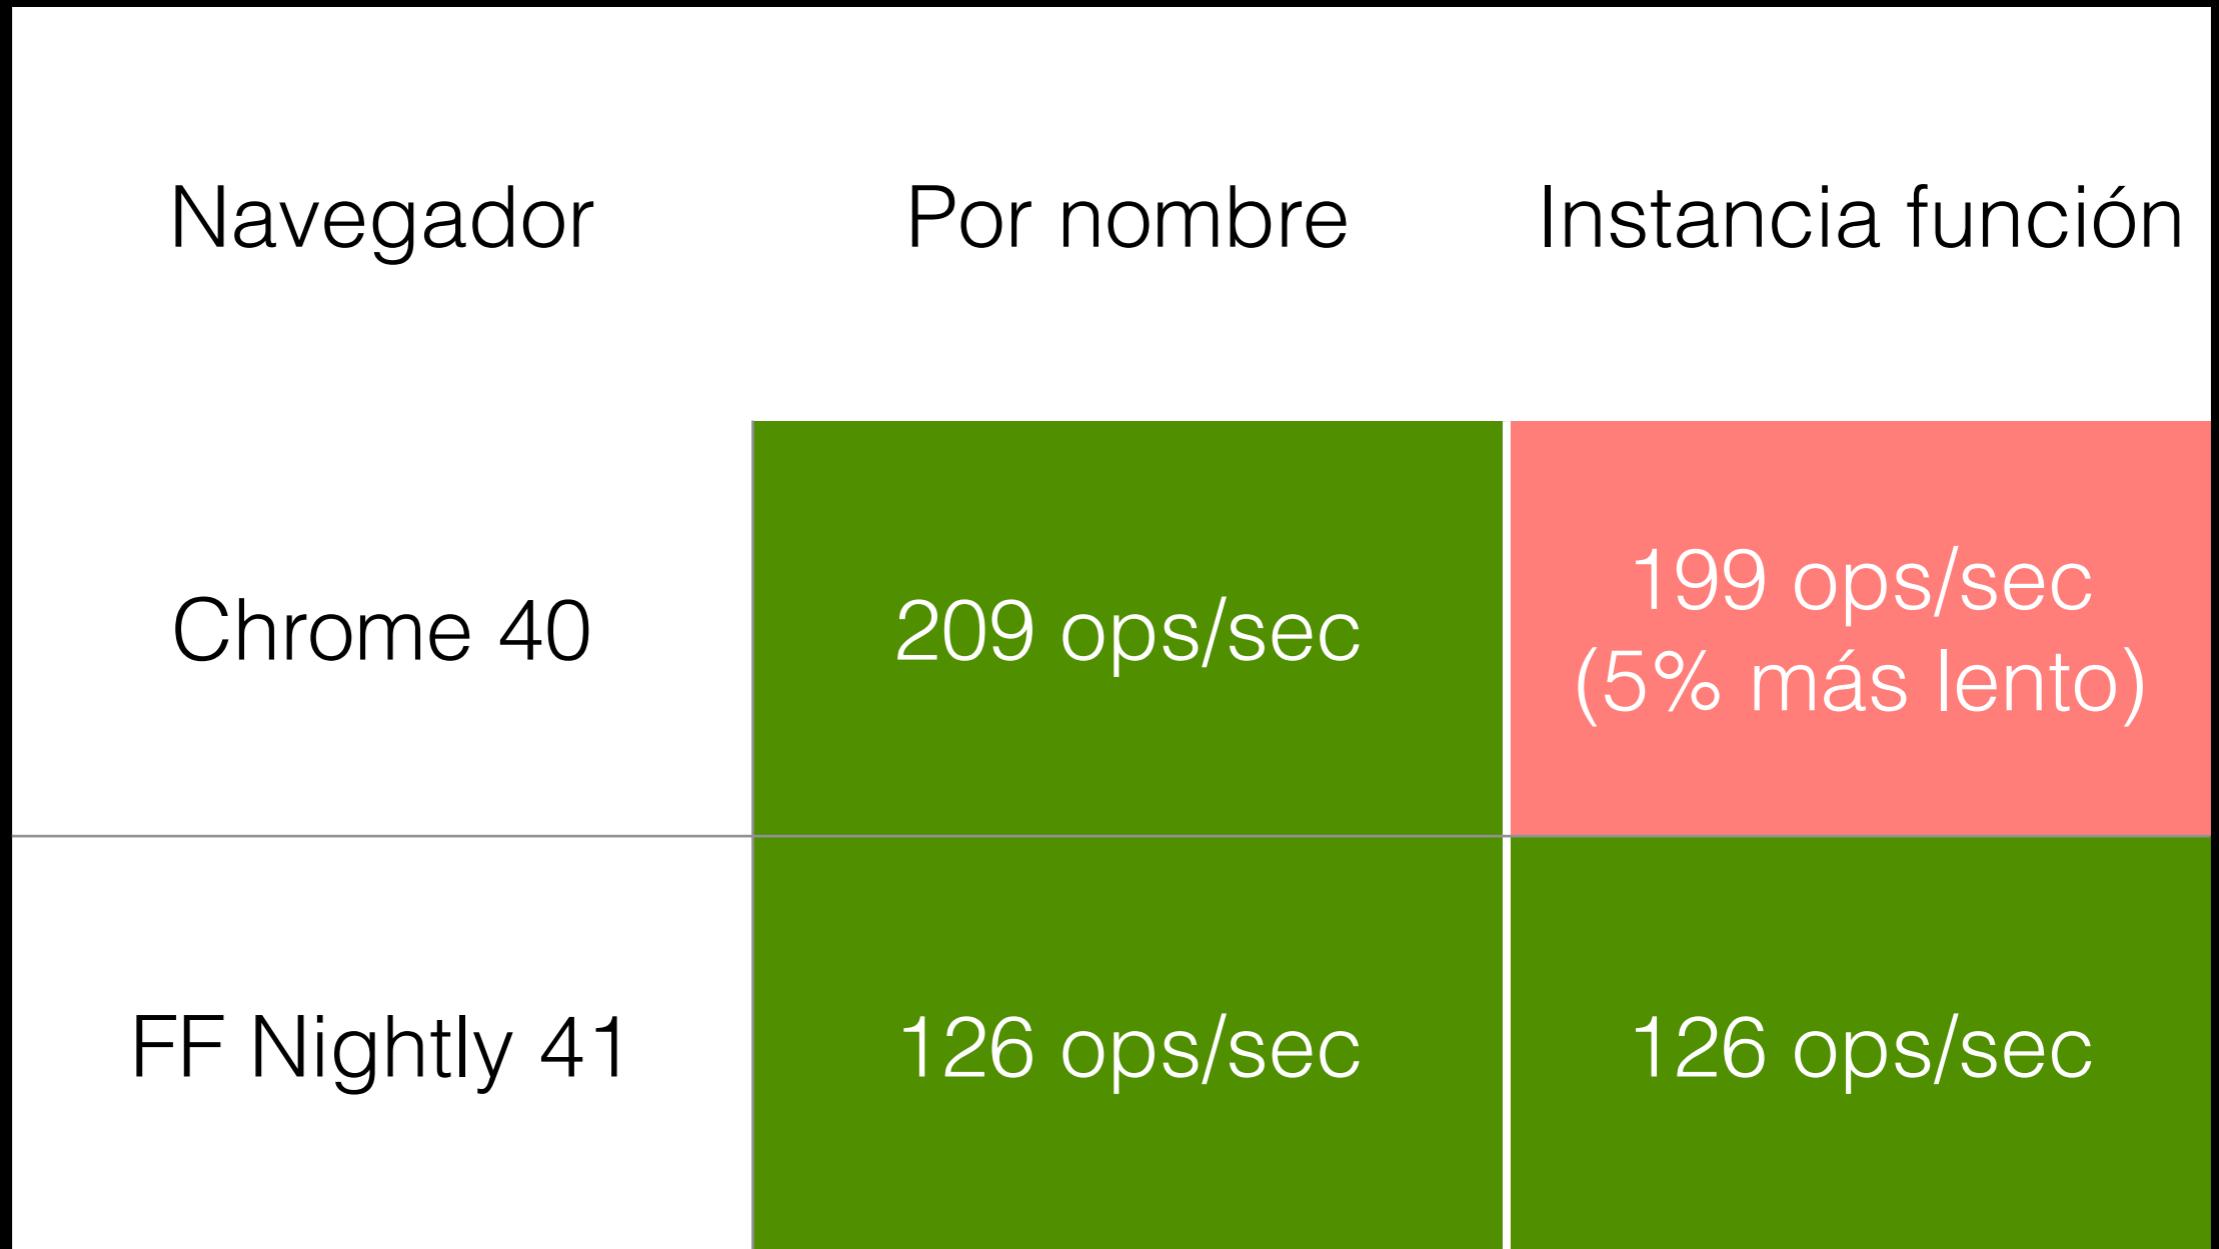


p-j-S

```
var promise = pjs.updateContext({
  add: x => x + 2,
}).then(() => {
  pjs(new Uint8Array([1,2,3,4]))
    .map('add')
    .seq()
    .then(result => {
      // result is [ 3, 4, 5, 6 ]
    });
});
```

Funciones por clave

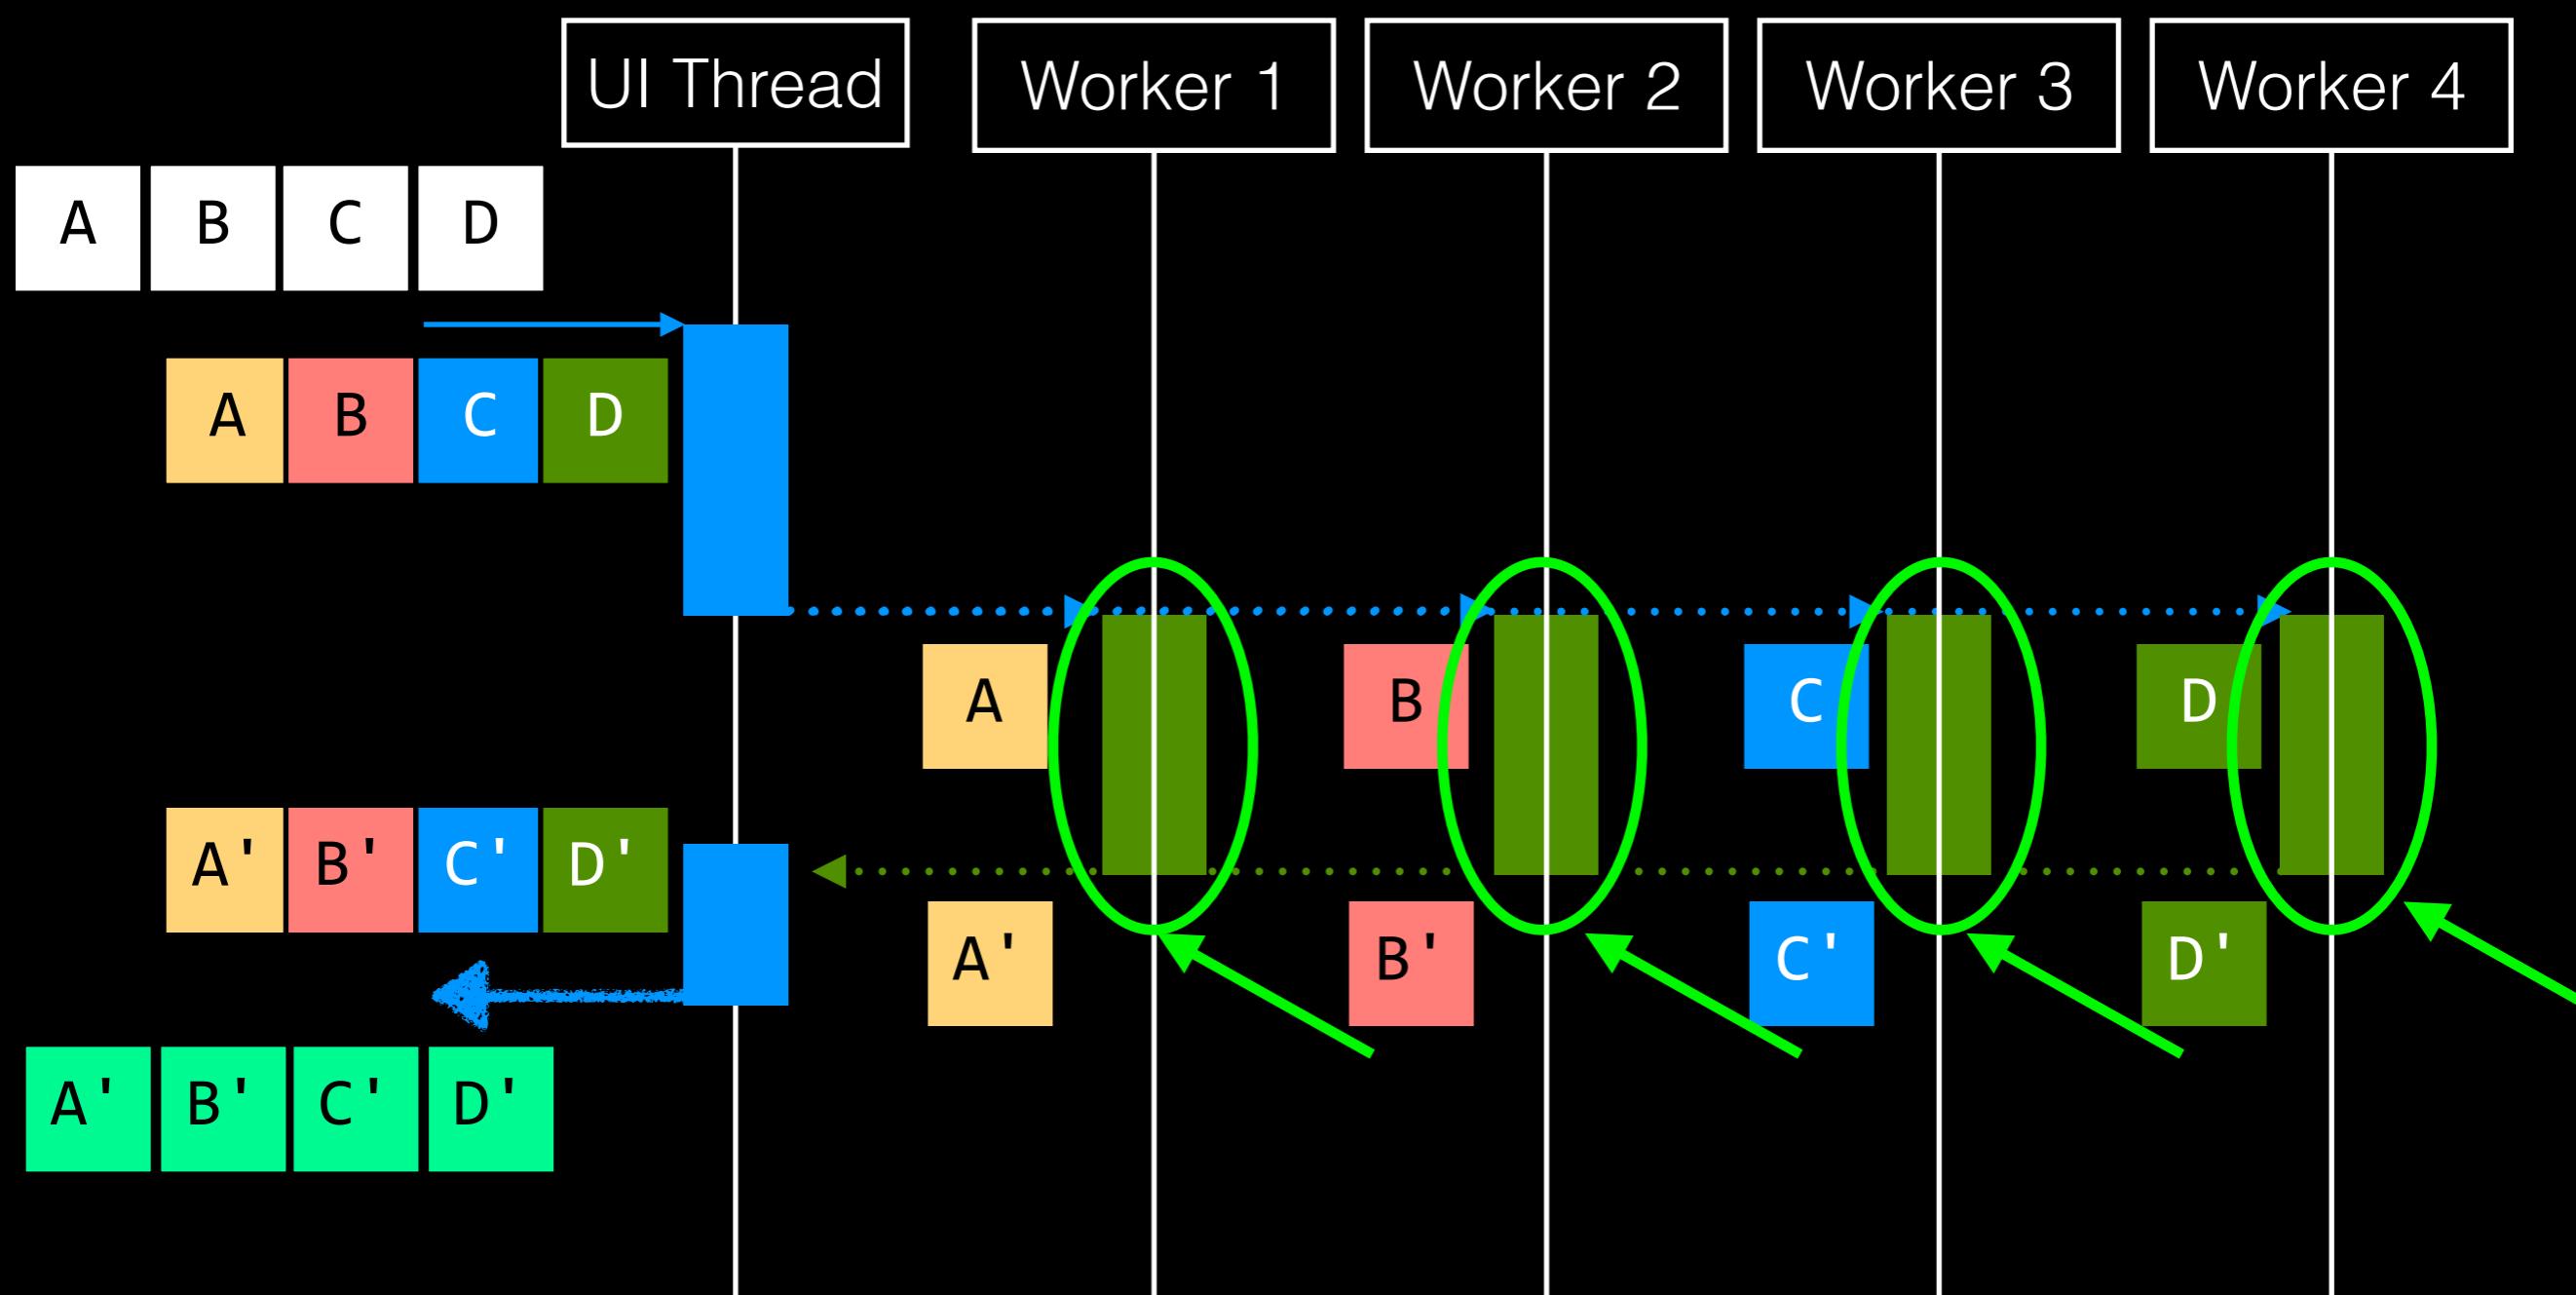
Ejemplo



Funciones por clave

<http://jsperf.com/global-ctx-inlined-func/3>

Inlining



p-j-S

```
function noise() {  
    return Math.random() * 0.5 + 0.5;  
};  
  
function clamp(component) {  
    return Math.max(Math.min(255, component), 0);  
}  
  
function distance(scale, dest, src) {  
    return clamp(scale * dest + (1 - scale) * src);  
};  
  
function map(data, len) {  
    for (var i = 0; i < len; i += 4) {  
        var r = data[i];  
        var g = data[i + 1];  
        var b = data[i + 2];  
  
        data[i] = distance(noise(), (r * 0.393) + (g * 0.769) + (b * 0.189), r);  
        data[i + 1] = distance(noise(), (r * 0.349) + (g * 0.686) + (b * 0.168), g);  
        data[i + 2] = distance(noise(), (r * 0.272) + (g * 0.534) + (b * 0.131), b);  
    }  
};
```

Tono sepia

Secuencial

```
pjs(new Uint32Array(canvasData.data.buffer)).map(pixel => {
  function noise() {
    return Math.random() * 0.5 + 0.5;
 };

  function clamp(component) {
    return Math.max(Math.min(255, component), 0);
  }

  function distance(scale, dest, src) {
    return clamp(scale * dest + (1 - scale) * src);
  };

  var r = pixel & 0xFF;
  var g = (pixel & 0xFF00) >> 8;
  var b = (pixel & 0xFF0000) >> 16;
  var new_r = distance(noise(), (r * 0.393) + (g * 0.769) + (b * 0.189), r);
  var new_g = distance(noise(), (r * 0.349) + (g * 0.686) + (b * 0.168), g);
  var new_b = distance(noise(), (r * 0.272) + (g * 0.534) + (b * 0.131), b);

  return (pixel & 0xFF000000) + (new_b << 16) + (new_g << 8) + (new_r & 0xFF);
}).seq(result => { /* trabajar con el resultado */});
```

Tono sepia

Paralelo

```
pjs(new Uint32Array(canvasData.data.buffer)).map(pixel => {
  function noise() {
    return Math.random() * 0.5 + 0.5;
 };

  function clamp(component) {
    return Math.max(Math.min(255, component), 0);
  }

  function distance(scale, dest, src) {
    return clamp(scale * dest + (1 - scale) * src);
  };

  var r = pixel & 0xFF;
  var g = (pixel & 0xFF00) >> 8;
  var b = (pixel & 0xFF0000) >> 16;
  var new_r = distance(noise(), (r * 0.393) + (g * 0.769) + (b * 0.189), r);
  var new_g = distance(noise(), (r * 0.349) + (g * 0.686) + (b * 0.168), g);
  var new_b = distance(noise(), (r * 0.272) + (g * 0.534) + (b * 0.131), b);

  return (pixel & 0xFF000000) + (new_b << 16) + (new_g << 8) + (new_r & 0xFF);
}).seq(result => { /* trabajar con el resultado */});
```

Tono sepia

Paralelo

```
pjs(buff).map(pixel => {
  var r = pixel & 0xFF;
  var g = (pixel & 0xFF00) >> 8;
  var b = (pixel & 0xFF0000) >> 16;
  var noise_r = Math.random() * 0.5 + 0.5;
  var noise_g = Math.random() * 0.5 + 0.5;
  var noise_b = Math.random() * 0.5 + 0.5;

  var new_r = Math.max(Math.min(255, noise_r * ((r * 0.393) + (g * 0.769) + (b * 0.189))
+ (1 - noise_r) * r), 0);
  var new_g = Math.max(Math.min(255, noise_g * ((r * 0.349) + (g * 0.686) + (b * 0.168))
+ (1 - noise_g) * g), 0);
  var new_b = Math.max(Math.min(255, noise_b * ((r * 0.272) + (g * 0.534) + (b * 0.131))
+ (1 - noise_b) * b), 0);

  return (pixel & 0xFF000000) + (new_b << 16) + (new_g << 8) + (new_r & 0xFF);
}.seq(result => { /* trabajar con resultado */});
```

Tono sepia

Paralelo

Navegador

Manual
inlining

Sin inlining

Chrome 40

27 ops/sec

6 ops/sec
(78% más lento)

FF Nightly 41

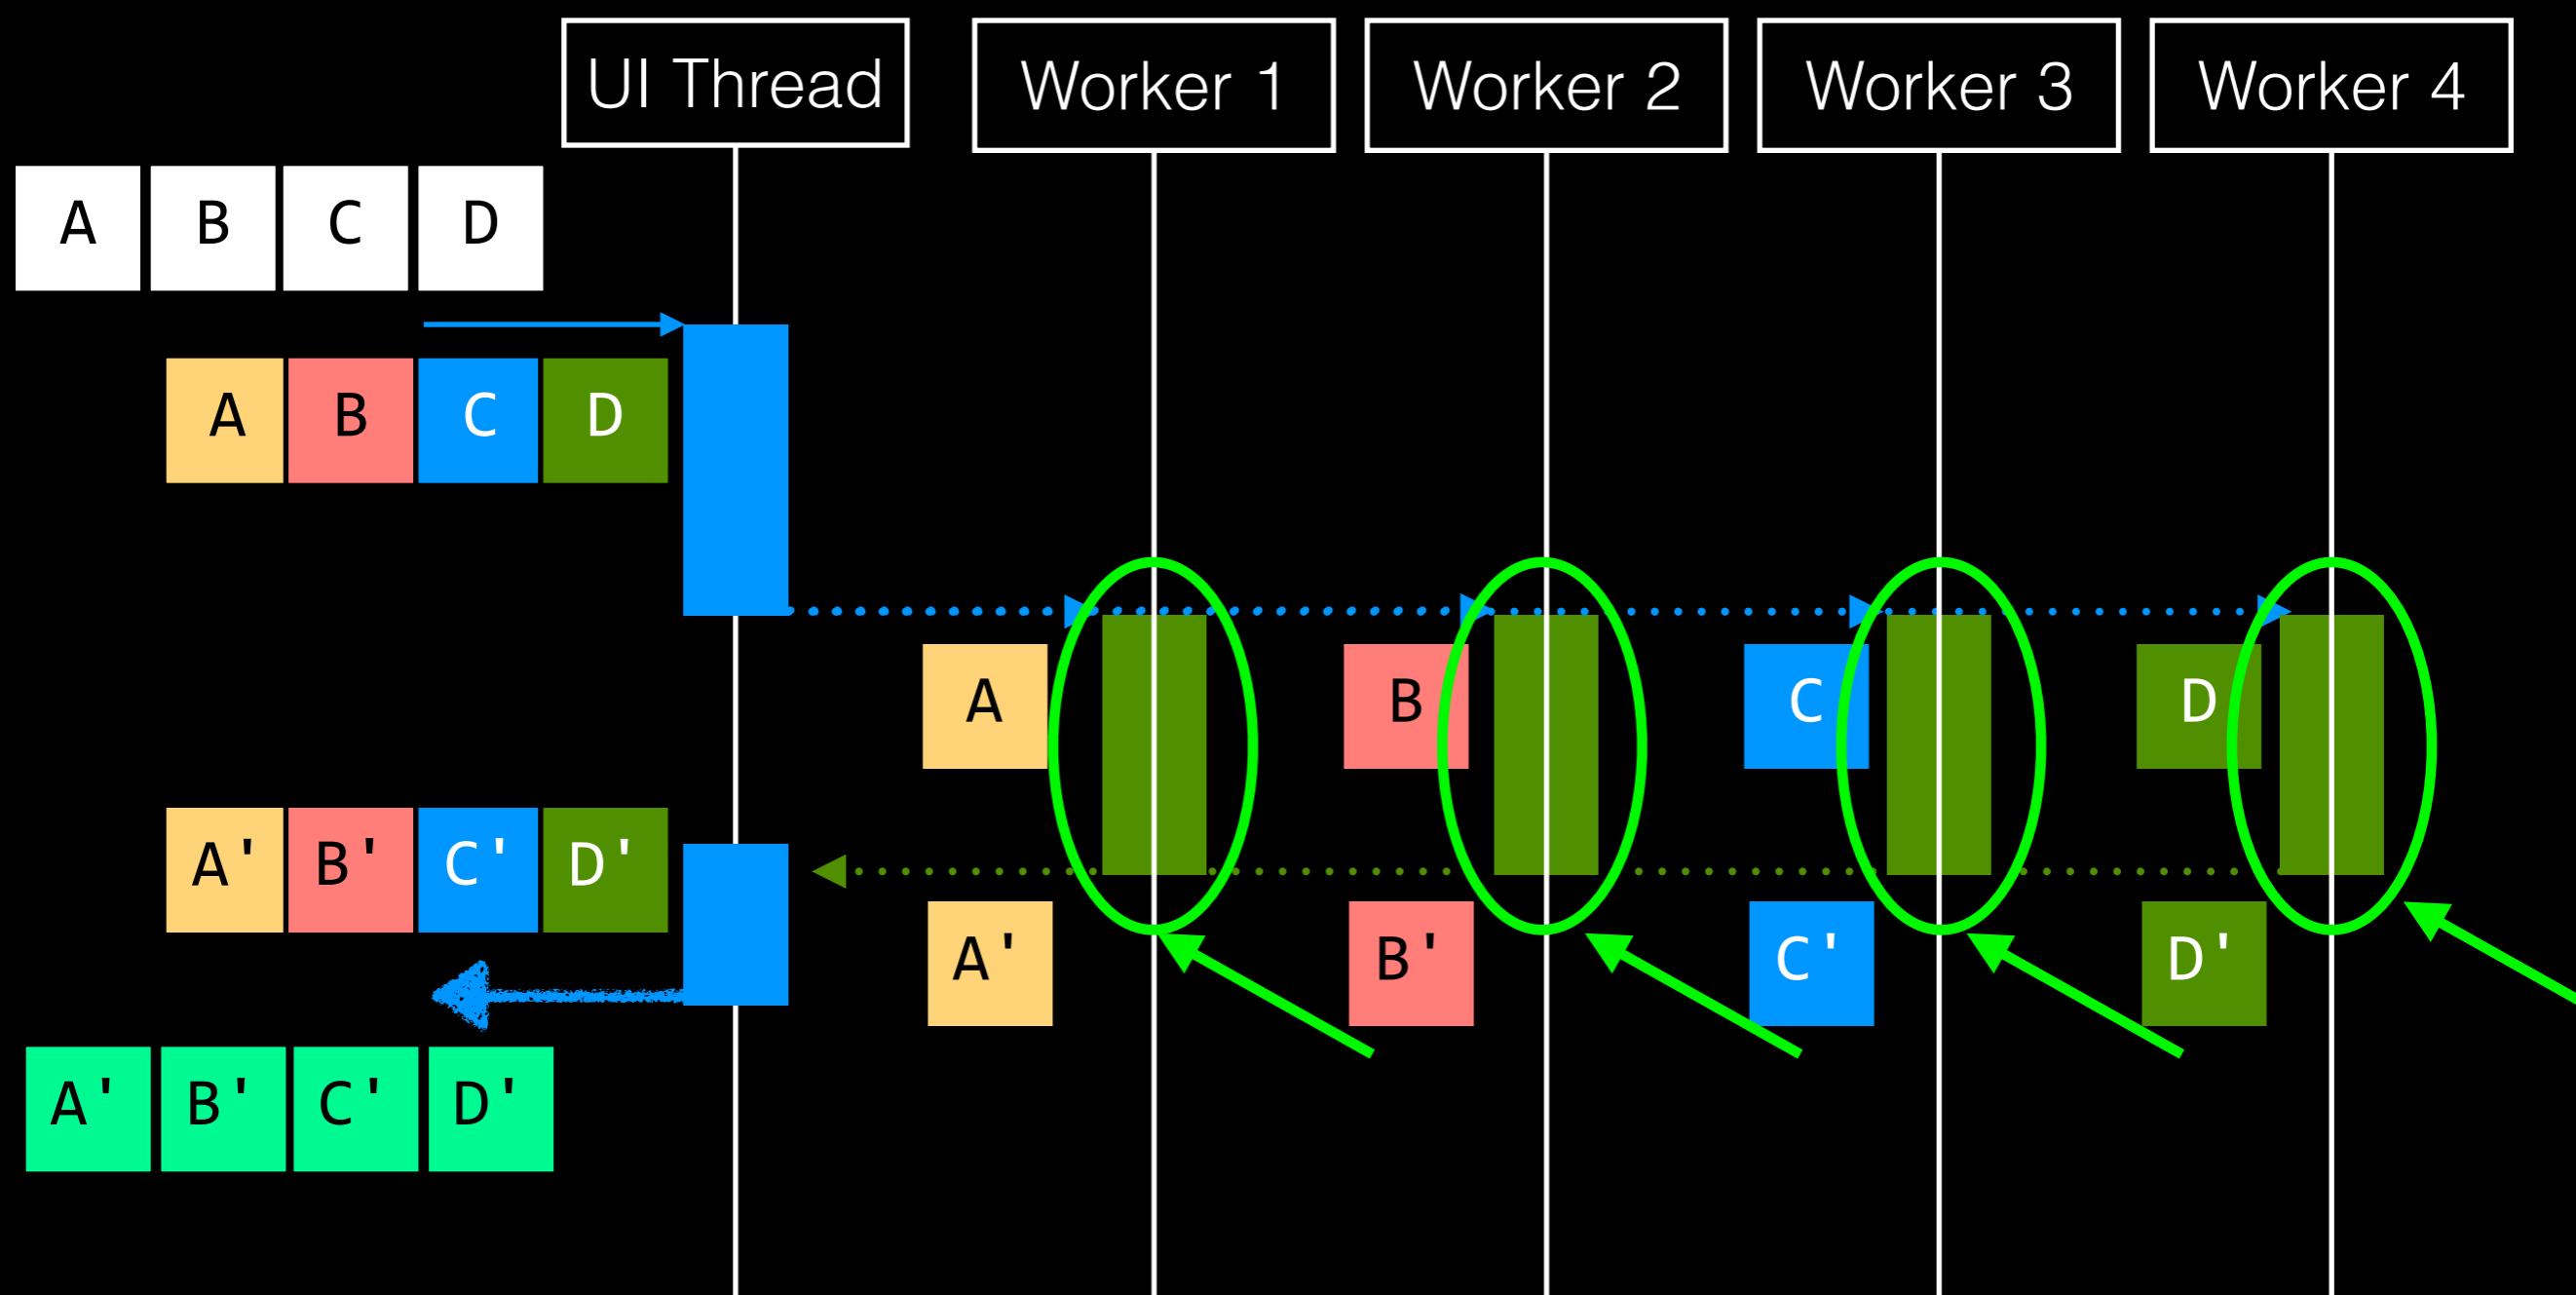
24 ops/sec

15 ops/sec
(38% más lento)

Inlining

<http://jsperf.com/pjs-map-inlining/7>

Caché de funciones



p-j-S

Navegador	Con caché	Sin caché
Chrome 40	26 ops/sec	26 ops/sec
FF Nightly 41	25 ops/sec	21 ops/sec (16% más lento)

Caché de funciones

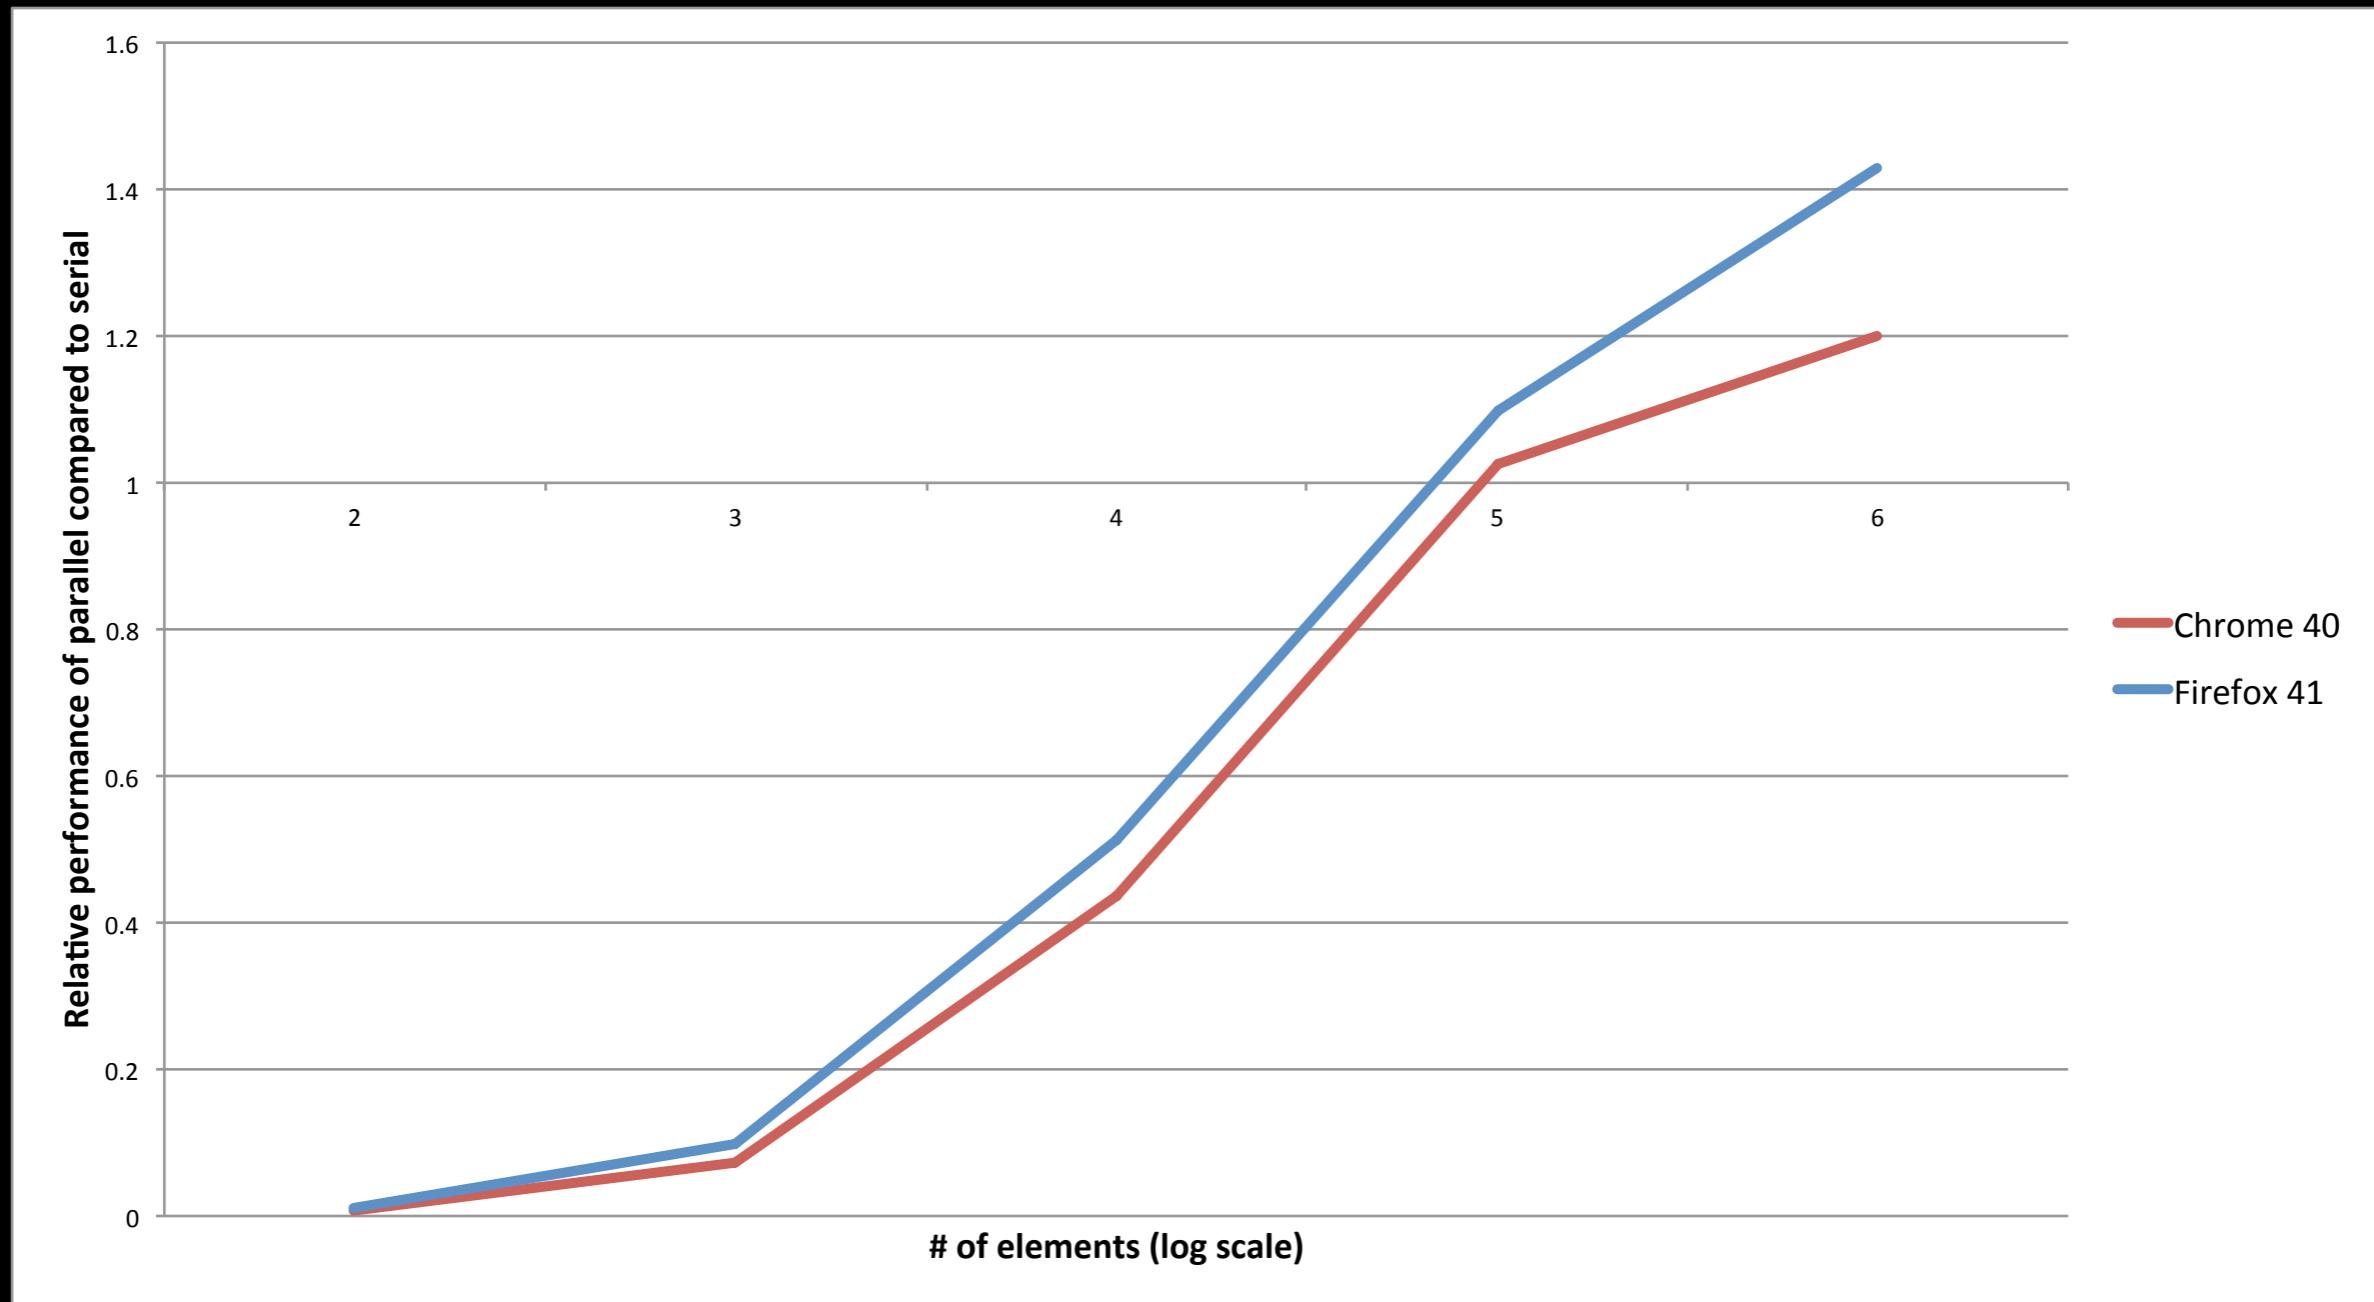
<http://jsperf.com/p-j-s-with-vs-without-function-cache/4>



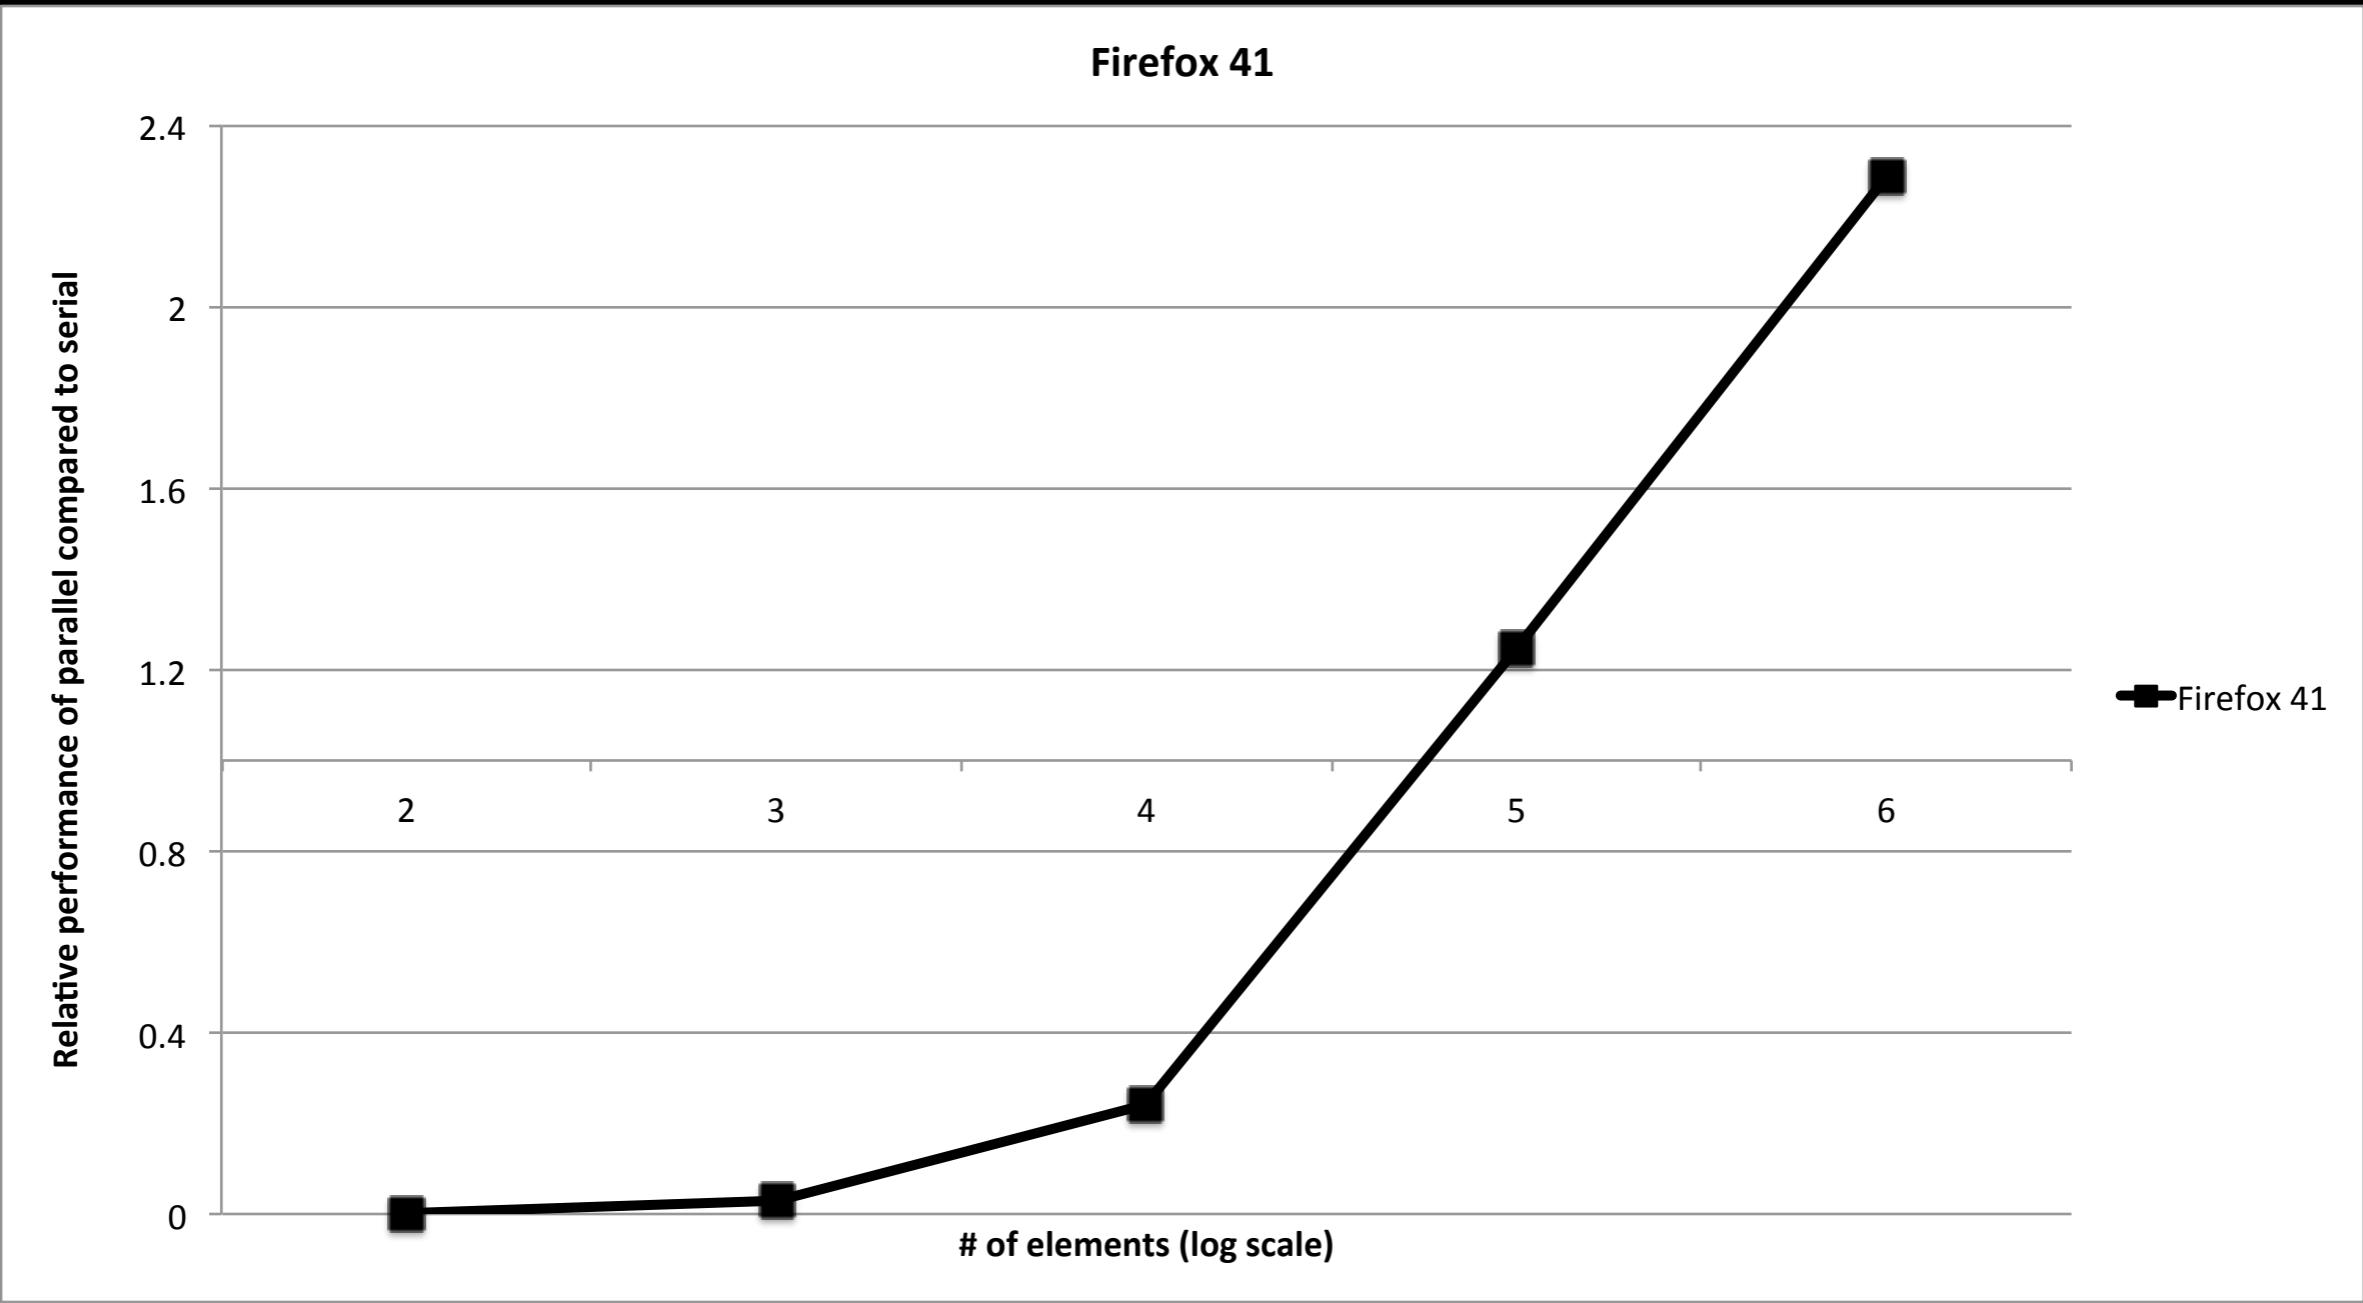
Case de estudio

Transformación de imágenes

Demo



TypedArray



SharedTypedArray

A library to parallelize `map`, `filter` and `reduce` operations on typed arrays through the use of Web Workers.

p-j-s

A JavaScript library that leverages Web Workers to provide parallelism when working with Typed Arrays

 Download ZIP

 Download TAR

 View On GitHub

This project is maintained by pjsteam

Hosted on [GitHub Pages](#) using the Dinky theme

Installing

```
npm i p-j-s
```

Usage

It's as simple as:

```
var pjs = require('p-j-s');

pjs.init({ maxWorkers: 4 }); // initialize the library

pjs(new Uint32Array([1,2,3,4]))
  .filter(function(e){
    return e % 2 === 0;
  })
  .map(function(e){
    return e * 2;
  })
  .seq()
  .then(function (result){
    // result is [4,8] a new Uint32Array

    // if we are not using the library any more cleanup once we are done
    pjs.terminate();
  });
}
```

Sitio web

<http://pjsteam.github.io/pjs/>

Entregables

- Paper propuesto para CARLA 2015
- Reporte técnico con benchmarks y decisiones de diseño
- Repositorio open source con el código de la biblioteca y documentación de uso
- Sitio web para la biblioteca: <http://pjsteam.github.io/pjs/>
- Charla en JSConfUY en la que se presentó la biblioteca

Preguntas?

Gracias!